Generative Models for Classification

Supervised learning in practice

Preprocessing
Explore & prepare data

Data Visualization and Exploration

Data Cleaning

Identify patterns that can be leveraged for learning

- Missing data
- Noisy data
- Erroneous data

Scaling (Standardization)

Feature Extraction

Prepare data for use in scale-dependent algorithms.

Dimensionality reduction eliminates redundant information

Model training

1 Supervised Learning Models: Linear models and KNN
(enough to get started using supervised learning)

5 Other algorithms and concepts:
- **Generative vs discriminative models**
- Parametric vs nonparametric models
- Model ensembles
- Feature/representation learning (neural networks, deep learning)

4 How to control model overfit: regularization strategies for model refinement

Iteratively fine tune the model

Performance evaluation

Make a prediction on validation data

2 Evaluating model performance and comparing models

Classification
Precision, Recall, F

3 How to make decisions using models

Regression
MSE, explained variance, R^2

Classifiers

Covered so far

K-Nearest Neighbors

Perceptron

Logistic Regression

Linear Discriminant Analysis
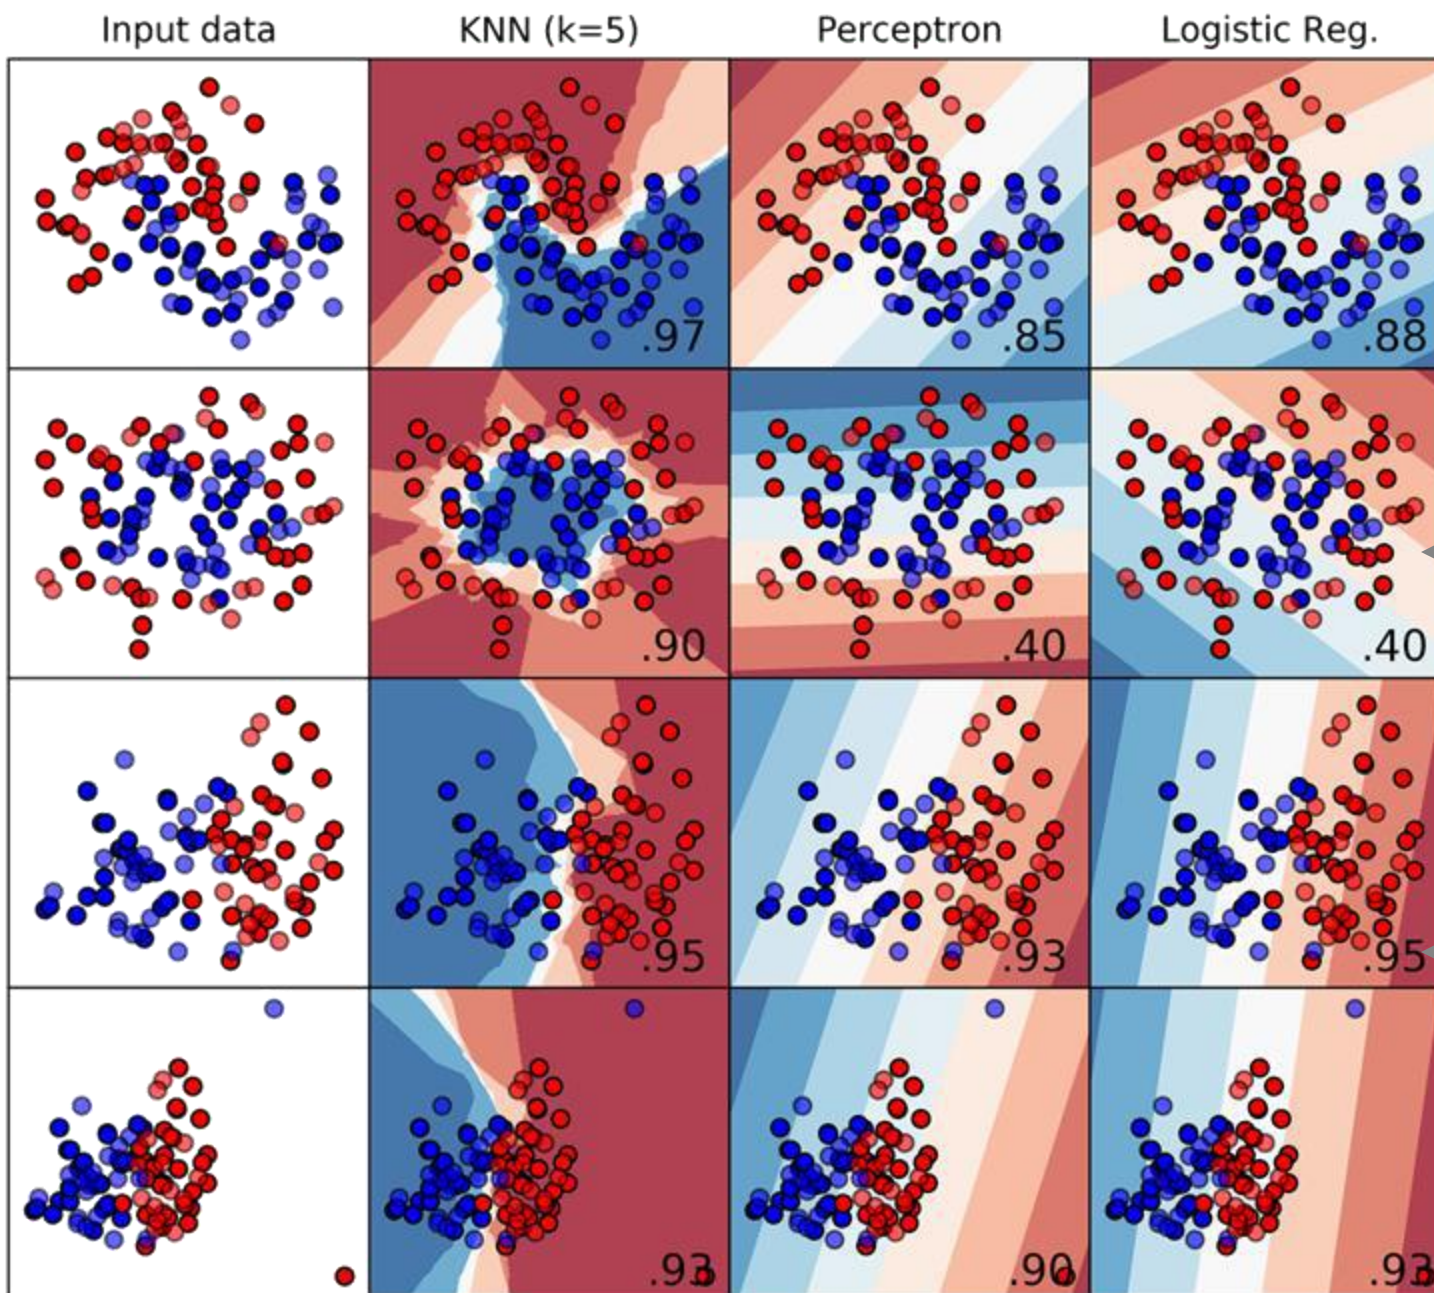
Quadratic Discriminant Analysis

Naïve Bayes

Along the way…
Revisiting Bayes' Rule
Projections from higher dimensions
Multivariate normal distributions

Comparison of classifiers we have seen so far

The color gradient shows the confidence scores

Test data accuracy

Bayes' rule in the context of classification

Bayes' Rule

Likelihood   Prior

$$P(C|X) = \frac{P(X|C)P(C)}{P(X)}$$

Posterior   Evidence

X Features

C Class label

i.e. C $\in \{c_0, c_1\}$ for the binary case

Bayes' Decision Rule:

choose the most probable class given the data

If $P(C_i = c_1|X_i) > P(C_i = c_0|X_i)$ then $\hat{y} = c_1$

otherwise $\hat{y} = c_0$

- If the distributions are correct, this decision rule is **optimal**
- Rarely do we have enough information to use this in practice

Class 1: Light Image

Class 0: Dark Image

Randomly draw a pixel from either of the images:

$$x_i = \boxed{149}$$

Darker pixel values are lower numbers (closer to 0), brighter pixels are higher numbers (closer to 255)

How do we determine which image the sample was most likely to have come from?

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Darker pixel values — Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Prior: $P(Y = y_i)$

0: Dark Image   1: Light Image

Bayes' Rule

$$P(Y = y_i|x) = \frac{P(x|Y = y_i)P(Y = y_i)}{P(x)}$$

Posterior

Likelihood   Prior

Evidence

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence

$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Prior: $P(Y = y_i)$

0: Dark Image   1: Light Image

Bayes' Rule

$$P(Y = y_i | x) = \frac{P(x|Y = y_i)P(Y = y_i)}{P(x)}$$

Posterior — Likelihood — Prior — Evidence

Likelihood

$P(\boldsymbol{x}|Y = \text{Light Image})$

$P(\boldsymbol{x}|Y = \text{Dark Alley})$

Evidence

$P(\boldsymbol{x}) = P(\boldsymbol{x}|y_0)P(y_0) + P(\boldsymbol{x}|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|\boldsymbol{x})$

$P(Y = \text{Light Image}|\boldsymbol{x})$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1 **Class 0: Dark Image** y_0

Prior: $P(Y = y_i)$

0: Dark Image   1: Light Image

Bayes' Rule $\text{Posterior} \quad P(Y = y_i|\boldsymbol{x}) = \dfrac{\overset{\text{Likelihood}}{P(\boldsymbol{x}|Y = y_i)}\overset{\text{Prior}}{P(Y = y_i)}}{\underset{\text{Evidence}}{P(\boldsymbol{x})}}$

Likelihood $P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence
$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|x)$

$P(Y = \text{Light Image}|x)$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1 **Class 0: Dark Image** y_0

Prior: $P(Y = y_i)$

0: Dark Image   1: Light Image

Decision rule:

If $P(Y = \text{Light Image}|x) > P(Y = \text{Dark Image}|x)$ then

else Dark Image

Light Image

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence

$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|x)$

$P(Y = \text{Light Image}|x)$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Green = classified as from Light Image

Grey = classified as from Dark Image

Classifying each of the individual pixels as being either from **Light Image** or **Dark Image** results in classification above

Decision rule:

If $P(Y = \text{Light Image}|x) > P(Y = \text{Dark Image}|x)$ then Light Image

else Dark Image

Likelihood unchanged

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence

$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|x)$

$P(Y = \text{Light Image}|x)$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Prior: $P(Y = y_i)$

Let's assume the sampling of pixels occurred more from the **Dark Image**

Dark Alley   Light Post

Bayes' Rule

$$P(Y = y_i|x) = \frac{P(x|Y = y_i)P(Y = y_i)}{P(x)}$$

Posterior   Likelihood   Prior   Evidence

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence

$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|x)$

$P(Y = \text{Light Image}|x)$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Green = classified as from Light Image

Grey = classified as from Dark Image

Prior: $P(Y = y_i)$

Assuming we the sampling of pixels occurred more from the **Dark Image**

Dark Alley   Light Post

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Evidence

$P(x) = P(x|y_0)P(y_0) + P(x|y_1)P(y_1)$

Posterior

$P(Y = \text{Dark Image}|x)$

$P(Y = \text{Light Image}|x)$

Darker pixel values

Brighter pixel values

Generative models model the **likelihood**

- These can also be used to generate synthetic data
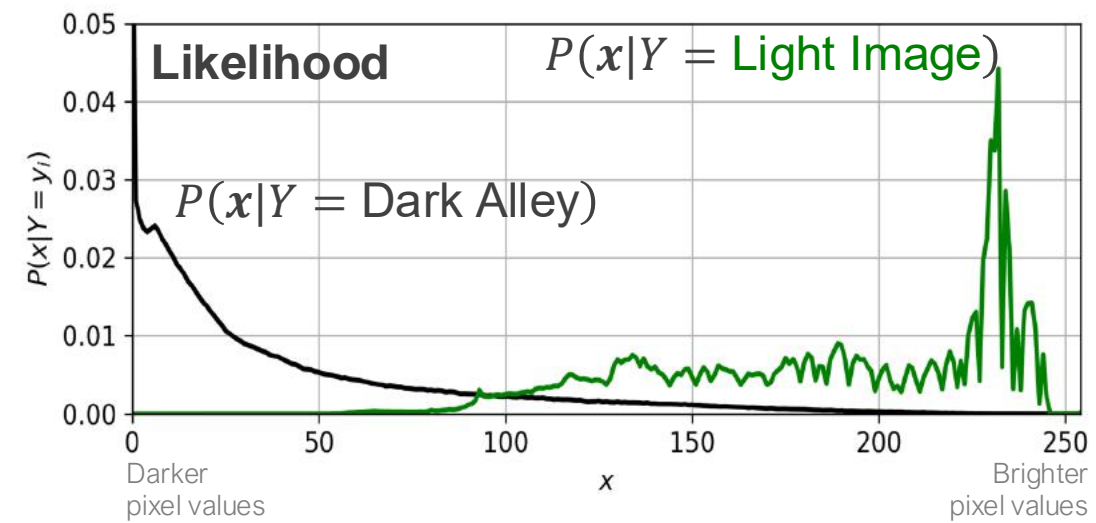- Often good performance when sample size is small

Examples: linear discriminant analysis, naïve Bayes, hidden Markov models, Guassian mixture models, Generative Adversarial Networks

$$P(Y = y_i|x) = \frac{\overset{\text{Likelihood}}{P(x|Y = y_i)}\overset{\text{Prior}}{P(Y = y_i)}}{\underset{\text{Evidence}}{P(x)}}$$

Posterior

Discriminative models model the **posterior**

- Or they just directly estimate labels without a probabilistic interpretation, $f(x) \to y$
- Often better performance for large sample sizes

Examples: logistic regression, support vector machines, neural networks, k nearest neighbors

Likelihood

$P(x|Y = \text{Light Image})$

$P(x|Y = \text{Dark Alley})$

Darker pixel values

Brighter pixel values

Class 1: Light Image y_1

Class 0: Dark Image y_0

Class 0 samples

Class 1 samples

Synthetic "Image" sampled from **Class 0**

Synthetic "Image" sampled from **Class 1**

How to "Generate" Data?

Prior: $P(Y = y_i)$

Dark Alley   Light Post

Bayes' Rule

Posterior $\quad$ Likelihood $\quad$ Prior

$$P(Y = y_i|x) = \frac{P(x|Y = y_i)P(Y = y_i)}{P(x)\ \text{Evidence}}$$

Generative modeling for classification

Assume we have c different classes

For a new sample, classify it as the class with the **largest posterior** $P(y = i | \boldsymbol{x})$

$$i = 1, \dots, c$$

Generative modeling for classification

If we have c different classes, we define a discriminant function, $d_i(\boldsymbol{x})$ for $i = 1, \ldots, c$

If $P(y = i | \boldsymbol{x}) > P(y = j | \boldsymbol{x})$ for $i \neq j$, then we classify feature $\boldsymbol{x}$ to class i

$$P(y = i | \boldsymbol{x}) = \frac{P(\boldsymbol{x} | y = i)P(y = i)}{P(\boldsymbol{x})}$$

$$= \frac{P(\boldsymbol{x} | y = i)P(y = i)}{\sum_{i=1}^{c} P(\boldsymbol{x} | y = i)P(y = i)}$$

Bayes' Rule:

Likelihood   Prior

$$P(Y | X) = \frac{P(\boldsymbol{X} | Y)P(Y)}{P(\boldsymbol{X})}$$

Posterior

Evidence

Denominator is the same for all classes i, so it won't help us tell which class's posterior is higher relative to other classes, so we ignore it going forward

We can define the discriminant function as:

$$d_i(x) = P(\boldsymbol{x} | y = i)P(y = i)$$

If we know the **true likelihood and prior** for our data, this process yields our **Bayes' classifier** (minimum misclassification error classifier)

Generative modeling for classification

$$d_i(x) = P(\boldsymbol{x}|y = i)P(y = i)$$

We **rarely** know our **true likelihood** for our data so we need to assume a form for the distributions and approximate

1 Assume a form for $P(\boldsymbol{x}|y = i)$

Gaussian ⟶ Linear and Quadratic Discriminant Analysis
Gaussian mixture models
Nonparametric density estimates

If we assume **independent features** → **Naïve Bayes**

(remember $d_i(x)$ proportional to the posterior)

**2** Assign the class, i, for which $d_i(x)$ is largest

Applies to both binary and multiclass problems

Discriminant Function: 2 classes

Decision rule:

Class 1 if: $\dfrac{d_1(\boldsymbol{x})}{d_0(\boldsymbol{x})} > \lambda_{thresh}$

Otherwise, class 0

Duda, Hart, and Stork, Pattern Classification, 2006

Discriminant Function: 2 classes

We build a classifier that assigns the class with the higher posterior probability:

If $\quad \dfrac{d_1(\boldsymbol{x})}{d_0(\boldsymbol{x})} = \dfrac{P(\boldsymbol{x}|y=1)P(y=1)}{P(\boldsymbol{x}|y=0)P(y=0)} > 1 \qquad$ Assign class 1, else class 0

Assumes these likelihoods are normal

$$N(\mu_1, \Sigma_1)$$

We can write our decision rule as a function of likelihoods (**likelihood ratio**)

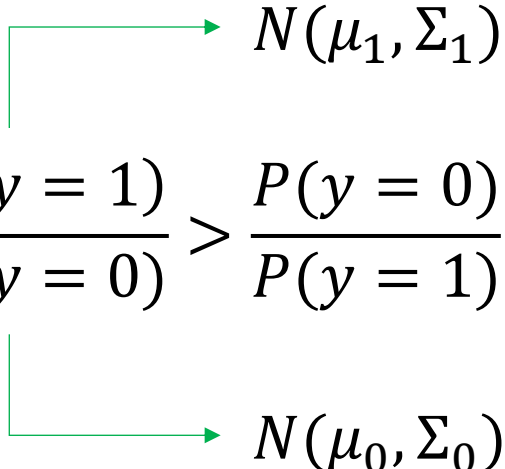$$\frac{P(\boldsymbol{x}|y=1)}{P(\boldsymbol{x}|y=0)} > \frac{P(y=0)}{P(y=1)}$$

Estimate the class-conditional mean and covariance matrix from the data

$$N(\mu_0, \Sigma_0)$$

Discriminant Function: 2 classes

We build a classifier that assigns the class with the higher posterior probability:

$$N(\mu_1, \Sigma_1)$$

Likelihood ratio:
$$\frac{P(\boldsymbol{x}|y=1)}{P(\boldsymbol{x}|y=0)} > \frac{P(y=0)}{P(y=1)}$$

$$N(\mu_0, \Sigma_0)$$

If we assume the class conditional distributions are **Gaussian**, this represents

Quadratic Discriminant Analysis

If we further assume the **covariance matrices for each class are the same**, $\Sigma_0 = \Sigma_1$, this represents

Linear Discriminant Analysis

Simple example with one feature

$P(\boldsymbol{x}|y=0)$ $P(\boldsymbol{x}|y=1)$

$$\frac{P(\boldsymbol{x}|y=1)}{P(\boldsymbol{x}|y=0)} > \frac{P(y=0)}{P(y=1)}$$

LDA Decision Rule

Bayes Decision Rule

Number of Examples
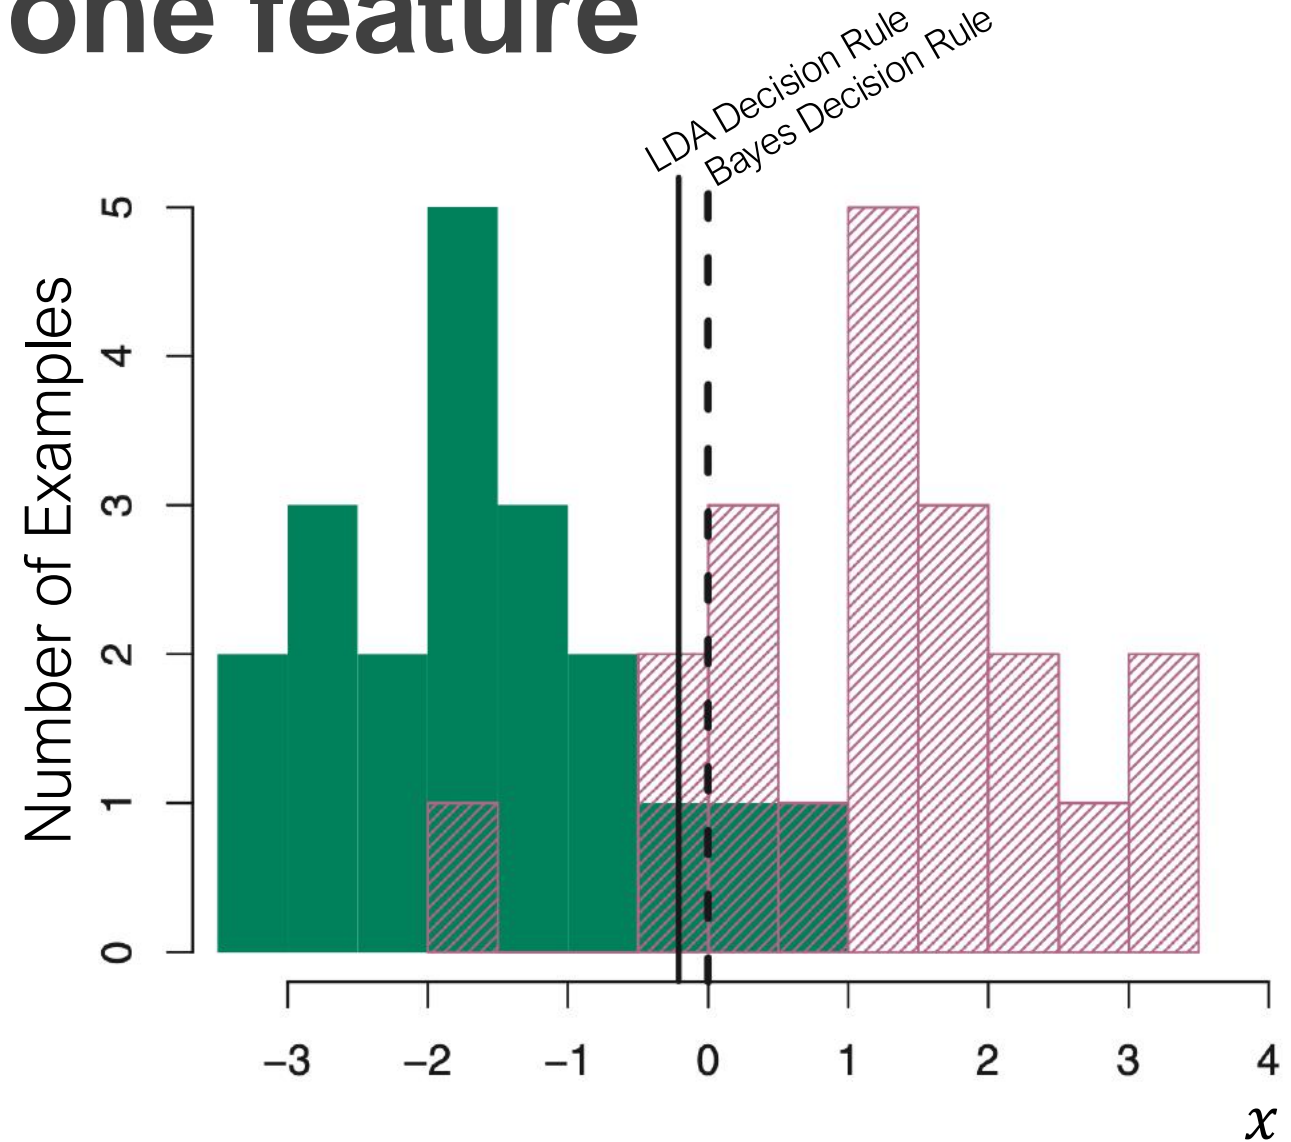
Figures from James et al. - Introduction to Statistical Learning

Covariance matrix

$$\mathbf{X}_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{iD} \end{bmatrix}$$

Vector of observation i

x_{ij}

Observation index   Predictor index

Predictors ⟶

$$\mathbf{X} = \text{Observations} \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1D} \\ x_{21} & x_{22} & \cdots & x_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \cdots & x_{ND} \end{bmatrix}$$

$$\boldsymbol{\Sigma} = \frac{1}{N} \sum_{i=1}^{N} (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \rightarrow [D \times D]$$

$$[D \times 1][1 \times D]$$

Mean of each predictor

If $\bar{x}_j = 0$ for all j

This will be the case IF the data are standardized

$$\boldsymbol{\Sigma} = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} & \cdots & \Sigma_{1D} \\ \Sigma_{21} & \Sigma_{22} & \cdots & \Sigma_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ \Sigma_{D1} & \Sigma_{D2} & \cdots & \Sigma_{DD} \end{bmatrix}$$

$$\Sigma_{jk} = \frac{1}{N} \sum_{i=1}^{N} (x_{ij} - \bar{x}_j)(x_{ik} - \bar{x}_k)$$

$$= \text{cov}(X_j, X_k)$$

$$= E[(X_j - \mu_j)(X_k - \mu_k)]$$

$$\Sigma_{jk} = \frac{1}{N} \sum_{i=1}^{N} x_{ij} x_{ik}$$

$$= \frac{1}{N} \mathbf{x}_j^T \mathbf{x}_k$$

$$= E[X_j X_k]$$

$$\boldsymbol{\Sigma} = \begin{bmatrix} \sigma_1^2 & \Sigma_{12} & \cdots & \Sigma_{1D} \\ \Sigma_{21} & \sigma_2^2 & \cdots & \Sigma_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ \Sigma_{D1} & \Sigma_{D2} & \cdots & \sigma_D^2 \end{bmatrix}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^{N} (x_{ij} - \bar{x}_j)^2$$

$$= E[(X_j - \mu_j)^2]$$

$$\boldsymbol{\Sigma} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$

Covariance and Correlation

Relationship between covariance and correlation

$$\text{corr}(X,Y) = \frac{\text{cov}(X,Y)}{\sqrt{\text{var}(X)\text{var}(Y)}}$$

When $\text{var}(X) = \text{var}(Y) = 1$, then:

$$\text{corr}(X,Y) = \text{cov}(X,Y)$$

If each of the features have been standardized, this means this matrix is:

$$\mathbf{\Sigma} = \begin{bmatrix} 1 & \rho_{12} & \cdots & \rho_{1D} \\ \rho_{21} & 1 & \cdots & \rho_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ \rho_{D1} & \rho_{D2} & \cdots & 1 \end{bmatrix}$$

Covariance Matrix Examples

$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 1 & -0.9 \\ -0.9 & 1 \end{bmatrix}$$

Key model differences

Quadratic Discriminant Analysis (QDA) $\boldsymbol{\Sigma}_0 \neq \boldsymbol{\Sigma}_1$ Each class (e.g., 0 or 1) may have a **unique** covariance matrix

Linear Discriminant Analysis (LDA) $\boldsymbol{\Sigma}_0 = \boldsymbol{\Sigma}_1$ Every class (e.g., 0 or 1) has an **identical** covariance matrix

Naïve Bayes with Gaussian Likelihoods
$$\boldsymbol{\Sigma}_i = \begin{bmatrix} \sigma_1^2 & 0 & \cdots & 0 \\ 0 & \sigma_2^2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_D^2 \end{bmatrix}$$
Every class (e.g., $i = 0$ or 1) has a **diagonal** covariance matrix

Input data	KNN (k=5)	Perceptron	Logistic Reg.	LDA	QDA
	.97	.85	.88	.88	.85
	.90	.40	.40	.40	.72
	.95	.93	.95	.95	.93
	.93	.90	.93	.90	.90
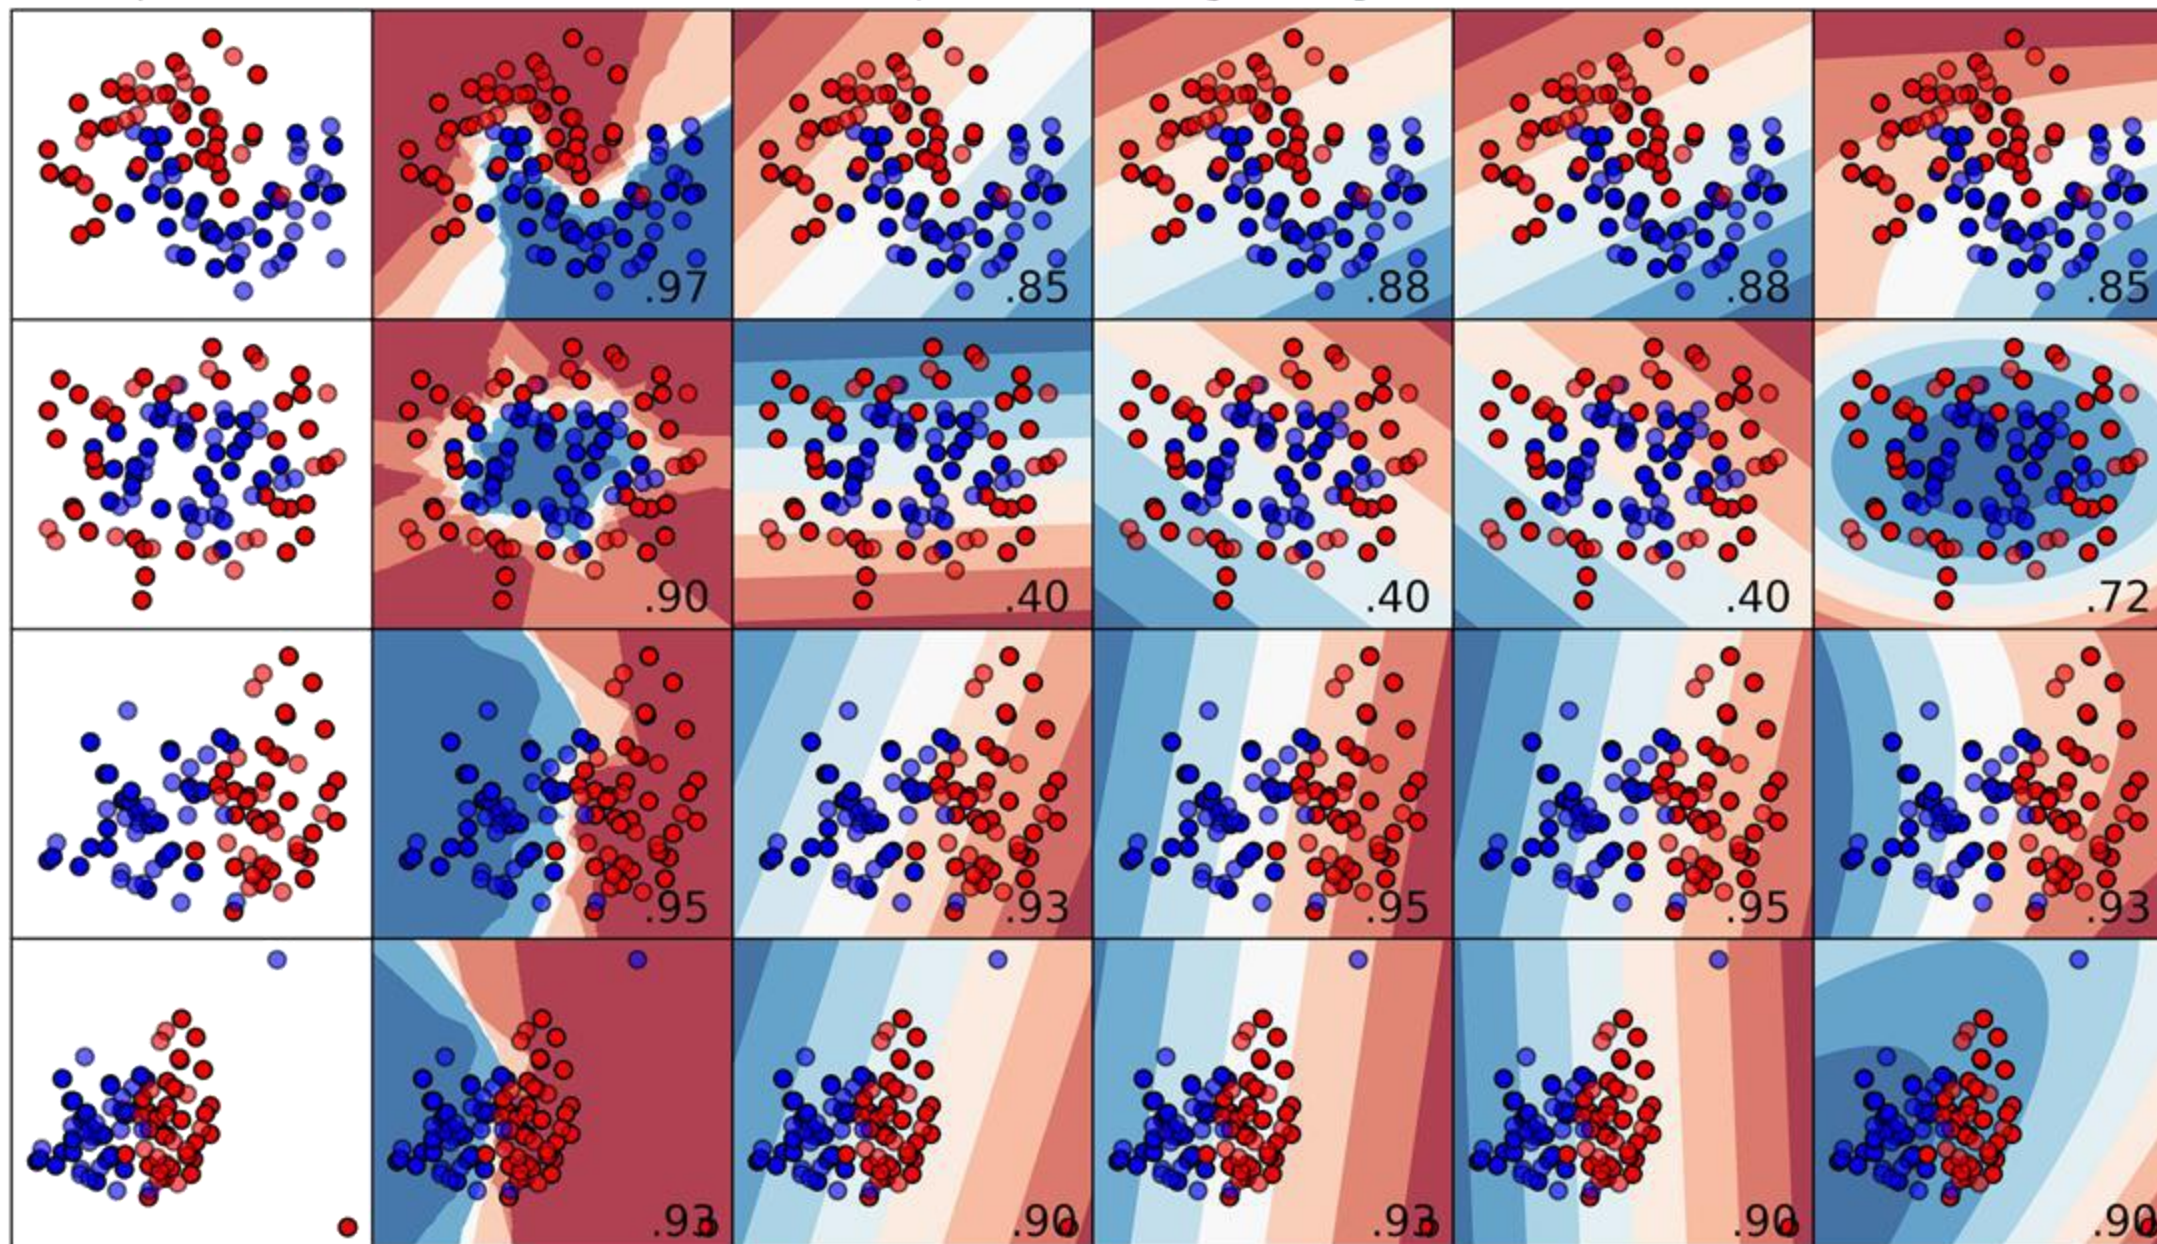
Joint vs Marginal Densities

The marginal densities don't factor in relationships between features

What if the joint density is too hard to estimate?
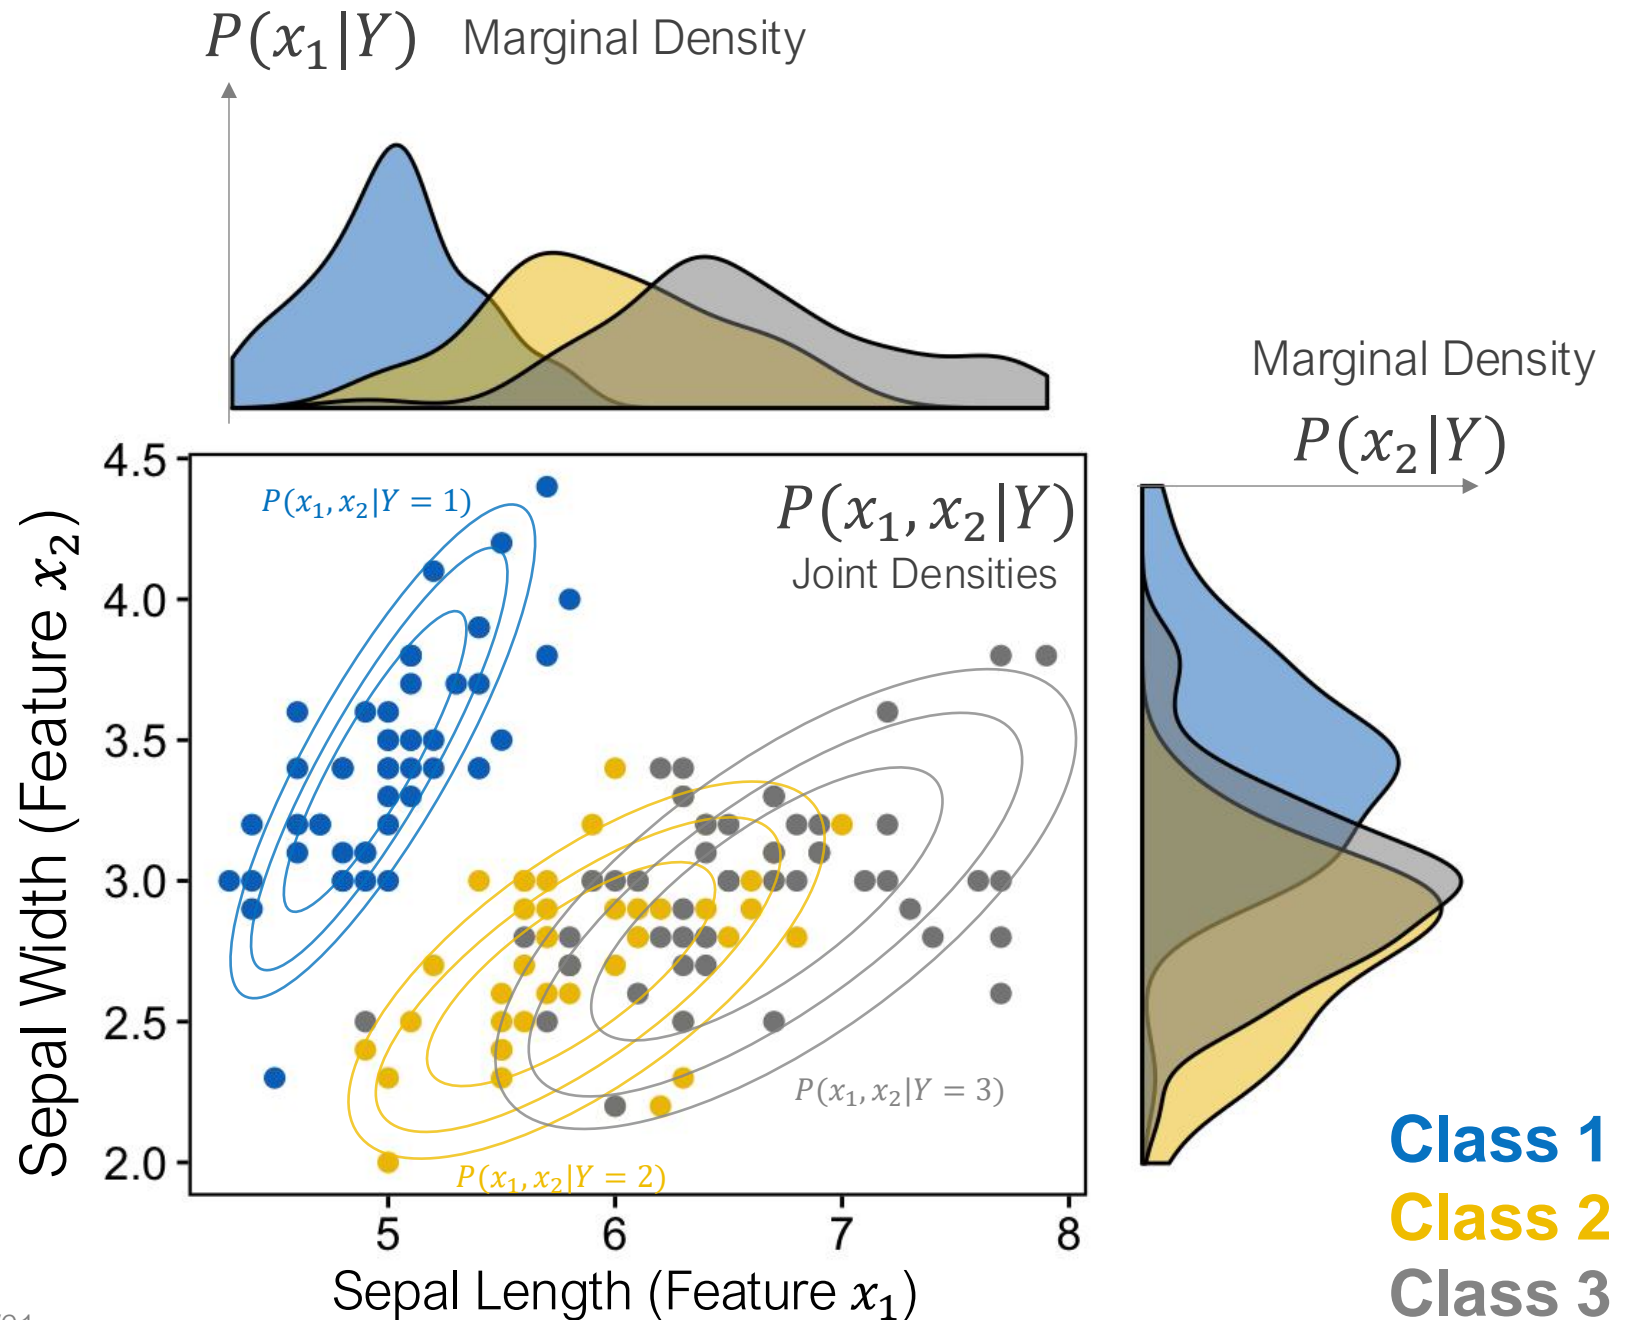
$P(x_1|Y)$ Marginal Density

$P(x_1, x_2|Y)$
Joint Densities

$P(x_1, x_2|Y=1)$

$P(x_1, x_2|Y=2)$

$P(x_1, x_2|Y=3)$

Sepal Width (Feature x_2)

Sepal Length (Feature x_1)

Marginal Density

$P(x_2|Y)$

Class 1
Class 2
Class 3

Naïve Bayes

Start with our original expression for our discriminant function
(proportional to posterior distribution)

$$d_i(\boldsymbol{x}) = P(\boldsymbol{x}|y = i)P(y = i)$$

Write out the full expression with all the terms in $\boldsymbol{x}$
(assume p predictors/features)

$$d_i(x_1, x_2, \dots, x_p) = P(x_1, x_2, \dots, x_p|y = i)P(y = i)$$

Assumption: Given the class, the features are independent

$$d_i(x_1, x_2, \dots, x_p) = P(y = i)\prod_{j=1}^{p} P(x_j|y = i)$$

Predict the class with the highest discriminant function
(i.e. posterior probability)

Naïve Bayes

We assign the class that has the largest discriminant (i.e. posterior probability)

$$d_i(x_1, x_2, \ldots, x_p) = P(y = i) \prod_{j=1}^{p} P(x_j | y = i)$$

This implies we estimate the density of each feature **separately**

This independence assumption is a strong assumption that is rarely valid

Considerably simplifies computation and data needs

Is flexible to allow for different distributional forms (i.e. Gaussian) or nonparametric techniques for the likelihood $P(x_j | y = i)$

Naïve Bayes: Gaussian example

We assign the class that has the largest discriminant (i.e. posterior probability)

$$d_i(x_1, x_2, \ldots, x_p) = P(y = i) \prod_{j=1}^{p} P(x_j | y = i)$$

This implies we estimate the density of each feature **separately**

If $P(x_j | y = i)$ is $N(\mu_{ji}, \sigma_{ji}^2)$, so for each class we estimate one mean and variance for each of the p features and for each class. We multiply **univariate** distributions together

$$d_i(x_1, x_2, \ldots, x_p) = P(y = i) \prod_{j=1}^{p} N(\mu_{ji}, \sigma_{ji}^2)$$

Naïve Bayes: Parameters
p predictors, c classes

Gaussian Naïve Bayes ($p = 2$)

For each predictor, x_i, and class, y_j:
$$\left(\mu_{ij}, \sigma_{ij}^2\right)$$

Total parameters = $2cp$

$P(x_i|Y = 0)$

$P(x_1|Y = 0)$

$P(x_2|Y = 0)$

$P(x_i|Y = 1)$

$P(x_1|Y = 1)$

$P(x_2|Y = 1)$

x_i

x_i

Without the Naïve Bayes independence assumption, each class would be a multivariate Gaussian with $(\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$

$$\boldsymbol{\mu}_j = \begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_p \end{bmatrix}, \boldsymbol{\Sigma}_j = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12}^2 & \cdots & \sigma_{1p}^2 \\ \sigma_{21}^2 & \sigma_{22}^2 & \cdots & \sigma_{2p}^2 \\ \vdots & \vdots & & \vdots \\ \sigma_{p1}^2 & \sigma_{p2}^2 & \cdots & \sigma_{pp}^2 \end{bmatrix}$$

Total parameters = $c\left(p + \frac{p^2 - p}{2} + p\right)$

Multivariate Gaussian ($p = 2$)

$P(x_1, x_2|Y = 0)$

$P(x_1, x_2|Y = 1)$

x_2

x_1

x_2

x_1

Image from ISL

Input data	KNN (k=5)	Perceptron	Logistic Reg.	LDA	QDA	Naive Bayes
	.97	.85	.88	.88	.85	.88
	.90	.40	.40	.40	.72	.70
	.95	.93	.95	.95	.93	.95
	.93	.90	.93	.90	.90	.90
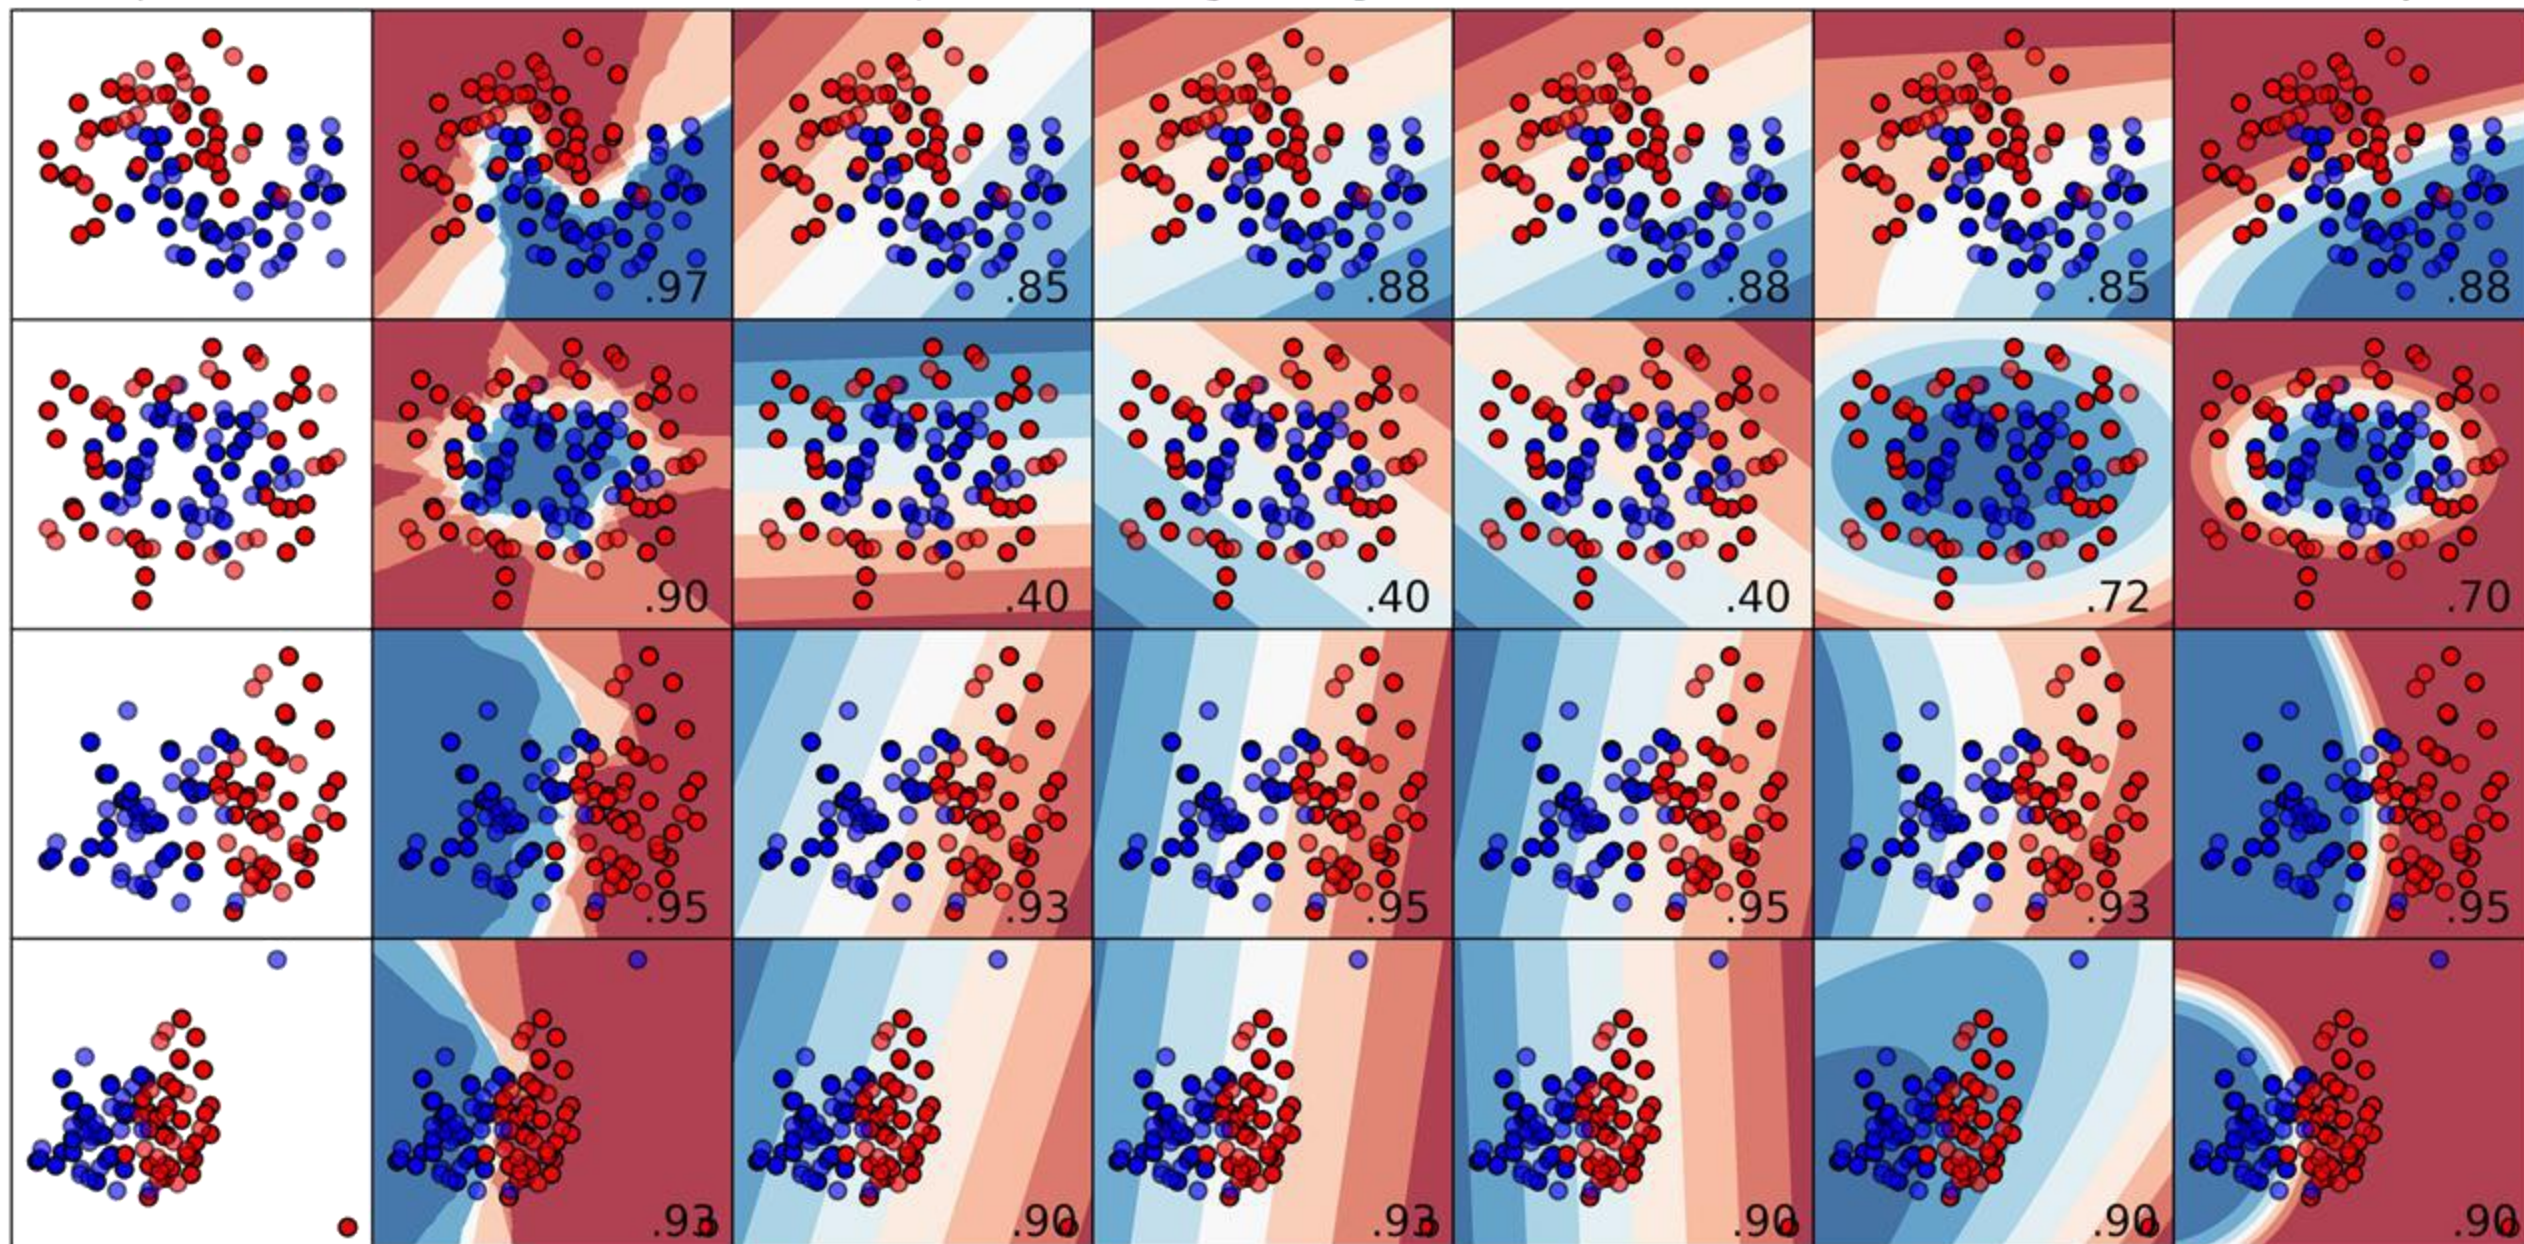
Classifiers

Covered so far

K-Nearest Neighbors

Perceptron

Logistic Regression

Linear Discriminant Analysis

Quadratic Discriminant Analysis

Naïve Bayes

Have closed-form solutions

Apply to multiclass problems

No hyperparameters

Fast to train

Requires small amounts of training data

Only choice is the form of P(X|Y) (otherwise no parameter choices)

Fast to train

Examples of deep generative models…

Variational Auto Encoders
Normalizing Flows
Generative Adversarial Networks (GANs)
Diffusion Models

Face Synthesis

Image Synthesis (link)

Karras et al. 2018, NVIDIA: Progressive growing of GANS for improved quality, stability, and variation

These images are all synthetic

Synthetic Generation

Karras, T., Laine, S. and Aila, T., 2019. A style-based generator architecture for generative adversarial networks. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 4401-4410).

Synthetic Images

Style Mixing:

From source B:
Pose, general hair style, face shape, eyeglasses

Hair style, eyes open/closed

Color scheme and microstructure

Supervised learning in practice

Preprocessing
Explore & prepare data

Data Visualization and Exploration

Data Cleaning

Identify patterns that can be leveraged for learning

• Missing data
• Noisy data
• Erroneous data

Scaling (Standardization)

Feature Extraction

Prepare data for use in scale-dependent algorithms.

Dimensionality reduction eliminates redundant information

Model training

1
Supervised Learning Models: Linear models and KNN
(enough to get started using supervised learning)

5
Other algorithms and concepts:
• **Generative vs discriminative models**
• Parametric vs nonparametric models
• Model ensembles
• Feature/representation learning (neural networks, deep learning)

4
How to control model overfit: regularization strategies for model refinement

Iteratively fine tune the model

Performance evaluation

Make a prediction on validation data

2
Evaluating model performance and comparing models

3
How to make decisions using models

Classification

Regression
MSE, explained variance, R^2