


Neural Networks II

What is a neural network and **how does it work?**

How do we **choose model weights?**
(i.e. how do we fit our model to data)

What are the challenges of using neural networks?

Forward propagation to create prediction and calculate training error / cost



Backpropagation lets us **compute the gradient** with respect to each of the parameters so we can tune them with **gradient descent**


$$E = \frac{1}{2} \sum_k (\hat{y}_k - y_k)^2$$

(gradient descent)

$$w_{ij} \leftarrow w_{ij} - \eta \frac{\partial E}{\partial w_{ij}}$$

Backpropagation is simply
the recursive application of
the chain rule

Example #1



$$y = 2a \quad \frac{\partial y}{\partial a} = 2$$


$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial a} \frac{\partial a}{\partial x} = (2)(3) = 6$$

$$a = 3x \quad \frac{\partial a}{\partial x} = 3$$

Chain Rule

Along a path we apply the chain rule

Example #2



$$y = 2a_1 + a_2$$

$$a_1 = 3x$$

$$a_2 = -x$$

$$\frac{\partial y}{\partial x} = \frac{\partial}{\partial x} (2a_1 + 1a_2)$$

Sum Rule

Across paths we apply the sum rule

$$= (2) \frac{\partial a_1}{\partial x} + (1) \frac{\partial a_2}{\partial x}$$

Chain Rule

$$= \frac{\partial y}{\partial a_1} \frac{\partial a_1}{\partial x} + \frac{\partial y}{\partial a_2} \frac{\partial a_2}{\partial x}$$

$$= (2)(3) + (1)(-1)$$

$$= 5$$

Example #3


$$y = \sum_{j=1}^N w_j^{(2)} a_j$$

$$\frac{\partial y}{\partial a_i} = w_i^{(2)}$$

$$a_i = w_i^{(1)} x$$

$$\frac{\partial a_i}{\partial x} = w_i^{(1)}$$

$$\frac{\partial y}{\partial x} = \sum_{j=1}^N \frac{\partial y}{\partial a_j} \frac{\partial a_j}{\partial x} = \sum_{j=1}^N w_j^{(2)} w_j^{(1)}$$



Example #3

$$y = \sum_{j=1}^N w_j^{(2)} a_j$$

$$\frac{\partial y}{\partial a_i} = w_i^{(2)}$$


$$a_i = w_i^{(1)} x$$

$$\frac{\partial a_i}{\partial x} = w_i^{(1)}$$

Derivatives with respect to the weights:

$$\frac{\partial y}{\partial w_i^{(2)}} = \frac{\partial}{\partial w_i^{(2)}} \left(\sum_{j=1}^N w_j^{(2)} a_j \right) = a_i$$

$$\frac{\partial y}{\partial w_i^{(1)}} = \frac{\partial y}{\partial a_i} \frac{\partial a_i}{\partial w_i^{(1)}} = w_i^{(2)} x$$



Backpropagation intuitively

Consider a derivative of a complicated function that can be represented as a long chain rule application

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial w} \underbrace{\frac{\partial w}{\partial x}}_{\frac{\partial f}{\partial x}} \frac{\partial x}{\partial y} \frac{\partial y}{\partial z}$$

$$\frac{\partial f}{\partial x} \quad \text{Chain rule equality}$$


This process of using the next step in the chain rule is backpropagation

$$\begin{aligned} \frac{\partial f}{\partial z} &= \frac{\partial f}{\partial w} \underbrace{\frac{\partial w}{\partial x}}_{\frac{\partial f}{\partial x}} \frac{\partial x}{\partial y} \frac{\partial y}{\partial z} &= \frac{\partial f}{\partial x} \underbrace{\frac{\partial x}{\partial y} \frac{\partial y}{\partial z}}_{\frac{\partial f}{\partial z}} &= \frac{\partial f}{\partial y} \underbrace{\frac{\partial y}{\partial z}}_{\frac{\partial f}{\partial z}} \\ && &= \frac{\partial f}{\partial z} \end{aligned}$$

Simple example

$$f(a, b) = a + b$$


This graph to the right represents this function



Local derivatives (one for each edge input into a node):

$$\frac{\partial y}{\partial a} = 1, \quad \frac{\partial y}{\partial b} = 1$$

$$f(\mathbf{x}, \mathbf{w}) = \frac{1}{w_0 x_0 + w_1}$$




- 1 Get inputs
- 2 Forward Pass
(forward propagate values)
- 3 Backward Pass
(back propagate gradients)

$$\begin{aligned} \frac{\partial f}{\partial y_2} &= \frac{\partial f}{\partial y_3} \frac{\partial y_3}{\partial y_2} \\ &= (1) \left(-\frac{1}{y_2^2} \right) \\ &= (1) \left(-\frac{1}{(-2)^2} \right) \\ &= -0.25 \end{aligned}$$

| | | | | |
|--------------------------------------|--|---|--|---------------------------------------|
| Local derivatives: (one per edge) | $\frac{\partial y_1}{\partial w_0} = x_0, \quad \frac{\partial y_1}{\partial x_0} = w_0$ | $\frac{\partial y_2}{\partial y_1} = \frac{\partial y_2}{\partial w_1} = 1$ | $\frac{\partial y_3}{\partial y_2} = -\frac{1}{y_2^2}$ | $\frac{\partial f}{\partial y_3} = 1$ |
|--------------------------------------|--|---|--|---------------------------------------|

Let's try an example closer to a real neural network



$$E = \frac{1}{2}(\hat{y} - y)^2$$

Forward propagation

In this particular case, this could be written as a single function of z_0


$$\hat{y} = z_2 = \sigma(w_2 \sigma(w_1 z_0))$$

We can calculate the error:

$$E = \frac{1}{2}(\hat{y} - y)^2$$

We want to estimate the gradient with respect to each parameter

$\frac{\partial E}{\partial w_i}$ **Backpropagation:** an efficient way of calculating these values



Forward propagation

$$E = \frac{1}{2}(\hat{y} - y)^2$$

$$\frac{\partial E}{\partial \hat{y}} = \hat{y} - y$$

$$\hat{y} = \sigma(a_2)$$

$$\frac{\partial \hat{y}}{\partial a_2} = \sigma'(a_2)$$

$$a_2 = w_2 z_1$$

$$\frac{\partial a_2}{\partial z_1} = w_2$$

$$z_1 = \sigma(a_1)$$

$$\frac{\partial z_1}{\partial a_1} = \sigma'(a_1)$$

$$a_1 = w_1 z_0$$

$$\frac{\partial a_1}{\partial w_1} = z_0 = x$$


Backpropagation

Let's calculate $\frac{\partial E}{\partial w_1}$

$$\frac{\partial E}{\partial w_1} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_2} \frac{\partial a_2}{\partial z_1} \frac{\partial z_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

$$\frac{\partial E}{\partial w_1} = (\hat{y} - y) \sigma'(a_2) w_2 \sigma'(a_1) z_0$$

We know all these quantities from forward propagation



Forward propagation (left to right):

$$\begin{aligned} E &= \frac{1}{2}(\hat{y} - y)^2 & \frac{\partial E}{\partial \hat{y}} &= \hat{y} - y \\ \hat{y} &= \sigma(a_2) & \frac{\partial \hat{y}}{\partial a_2} &= \sigma'(a_2) \\ a_2 &= w_2 z_1 & \frac{\partial a_2}{\partial z_1} &= w_2 \\ z_1 &= \sigma(a_1) & \frac{\partial z_1}{\partial a_1} &= \sigma'(a_1) \\ a_1 &= w_1 z_0 & \frac{\partial a_1}{\partial w_1} &= z_0 = x \end{aligned}$$

Backpropagation (right to left):


$$\frac{\partial E}{\partial w_1} = \underbrace{\frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_2} \frac{\partial a_2}{\partial z_1} \frac{\partial z_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}}_{\frac{\partial E}{\partial a_1}}$$

These derivatives with respect to the activations, a_i , allow us to quickly calculate each of our parameter derivatives:

$$\frac{\partial E}{\partial w_i} = \frac{\partial E}{\partial a_i} \frac{\partial a_i}{\partial w_i} = \frac{\partial E}{\partial a_i} z_{i-1}$$

δ_i (common shorthand)

And here's an actual mini neural network...



Activations in the other layer are the function $g(x)$, with derivative $g'(x)$

Linear activation for regression in the output layer

Local Relationship(s)

$$a_i = w_{i0}x_0 + w_{i1}x_1 + w_{i2}x_2$$

$$z_i = g(a_i)$$

$$\hat{y} = w_0z_0 + w_1z_1 + w_2z_2$$

$$C(y, \hat{y}) = E = \frac{1}{2}(\hat{y} - y)^2$$

Local Derivative(s)

$$\frac{\partial a_i}{\partial w_{ij}} = x_j$$

$$\frac{\partial z_i}{\partial a_i} = g'(a_i)$$

$$\frac{\partial \hat{y}}{\partial z_i} = w_i$$

$$\frac{\partial \hat{y}}{\partial w_i} = z_i$$

$$\frac{\partial E}{\partial \hat{y}} = \hat{y} - y$$

Global Derivative(s)

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_i} \frac{\partial z_i}{\partial a_i} \frac{\partial a_i}{\partial w_{ij}}$$

$$\frac{\partial E}{\partial a_i} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_i} \frac{\partial z_i}{\partial a_i}$$

$$\frac{\partial E}{\partial w_i} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_i}$$

$$\frac{\partial E}{\partial \hat{y}} = \hat{y} - y$$

Quick reference for neural network math

5.1 Forward Propagation

Forward propagation is the iterative application of the following two equations:


$$a_i^{(k)} = \sum_{j=1}^{N_{k-1}} w_{ij}^{(k)} z_j^{(k-1)}$$

$$z_i^{(k)} = \sigma(a_i^{(k)})$$

In matrix form those equations are:

$$\mathbf{a}^{(k)} = \mathbf{W}^{(k)} \mathbf{z}^{(k-1)}$$

$$\mathbf{z}^{(k)} = \sigma(\mathbf{a}^{(k)})$$



https://github.com/kylebradbury/neural-network-math/raw/master/neural_network_math.pdf

5.2 Backpropagation

Backpropagation begins with the calculation of the gradient of the error with respect to the final set of activations, which for mean square error with sigmoidal activation and K -layer neural network is:

$$\delta_i^{(K)} \triangleq \frac{\partial E_n}{\partial a_i^{(K)}} = (z_i^{(3)} - y_i) \sigma'(a_i^{(3)})$$

We then propagate that back through the neural network and calculate the gradients with respect to each weight along the way (for $k = K - 1, \dots, 1$):


$$\delta_i^{(k)} \triangleq \frac{\partial E_n}{\partial a_i^{(k)}} = \sigma'(a_i^{(k)}) \sum_{j=1}^{N_{k+1}} \delta_j^{(k+1)} w_{ji}^{(k+1)}$$

$$\frac{\partial E_n}{\partial w_{ij}^{(k)}} = \frac{\partial E_n}{\partial a_i^{(k)}} \frac{\partial a_i^{(k)}}{\partial w_{ij}^{(k)}} = \delta_i^{(k)} z_j^{(k-1)}$$

Or in matrix form:

$$\boldsymbol{\delta}^{(k)} \triangleq \frac{\partial E_n}{\partial \mathbf{a}^{(k)}} = \mathbf{W}^{(k+1)\top} \boldsymbol{\delta}^{(k+1)} \circ \sigma'(\mathbf{a}^{(k)})$$

$$\frac{\partial E_n}{\partial \mathbf{w}^{(k)}} = \boldsymbol{\delta}^{(k)} \mathbf{z}^{(k-1)\top}$$



https://github.com/kylebradbury/neural-network-math/raw/master/neural_network_math.pdf

Backpropagation

- 1 Run forward propagation on an input and calculate all the activations, a_i
- 2 Evaluate $\delta_i^{(k)} = \frac{\partial E}{\partial a_i^{(k)}}$ for all nodes in the network
- 3 Compute the weight derivatives: $\frac{\partial E}{\partial w_{ij}^{(k)}} = \delta_i^{(k)} z_j^{(k-1)}$ for all nodes in the network

Now we have all the derivatives we need, so we can run gradient descent

Gradient Descent

Batch gradient descent

1 Calculate the gradient for each training sample and average them

2 Update all the parameters based on that average gradient

3 Repeat 1 and 2 until stopping criteria met

$$\overline{\frac{\partial E}{\partial w_{ij}}} = \frac{1}{N} \sum_{n=1}^N \frac{\partial E_n}{\partial w_{ij}}$$

$$w_{ij} \leftarrow w_{ij} - \eta \overline{\frac{\partial E}{\partial w_{ij}}}$$

Our loss function (E_n) is calculated for EACH training sample $n = 1, 2, \dots, N$

$$E_n = \frac{1}{2} (\hat{y}_n - y_n)^2$$

The gradient also needs to be calculated for each sample (i.e. backprop needs to be run for each sample)

Stochastic gradient descent (SGD)

1 Randomly sort the list of training samples

2 Calculate the gradient from one training sample

3 Update all the parameters based on that error

4 Repeat 2 and 3 until all training samples have been used, then repeat 1-3 until stopping criteria met

$$\frac{\partial E_n}{\partial w_{ij}}$$

$$w_{ij} \leftarrow w_{ij} - \eta \frac{\partial E_n}{\partial w_{ij}}$$

Minibatch gradient descent

A tweak to SGD where you use a small batch of training samples rather than the whole dataset.

The average gradient across this minibatch is used for taking a gradient descent step

Other optimizers exist:
e.g. momentum, RMSprop, adam


Batch gradient descent can work well if the cost function is convex (rare) and the learning rate is properly tuned

Batch gradient descent tends to converge more slowly

Stochastic gradient descent (SGD) is better at avoiding local minima, but injects noise into the training process

Often **minibatch** gradient descent is used (a small batch of data is used instead of a single sample in SGD) and balances the two

Convex cost function parameter space



Batch gradient descent can work well if the cost function is convex (rare) and the learning rate is properly tuned

Batch gradient descent tends to converge more slowly


Stochastic gradient descent (SGD) is better at avoiding local minima, but injects noise into the training process

Often **minibatch** gradient descent is used (a small batch of data is used instead of a single sample in SGD) and balances the two

Cost function parameter space with multiple local minima




Batch Gradient Descent
Stochastic Gradient Descent




SGD Example

Gabor Model

$$f[x, \phi] = \sin[\phi_0 + 0.06 \cdot \phi_1 x] \cdot \exp\left(-\frac{(\phi_0 + 0.06 \cdot \phi_1 x)^2}{32.0}\right)$$




Training Data




Global cost function for different parameter values

Images from Prince, Understanding Deep Learning, 2023




Global cost function

Cost function with the entire training dataset



Minibatch cost function 2

Cost function with minibatch



What makes minibatch SGD stochastic? Different minibatches lead to **different cost functions**

Minibatch cost function 1

Cost function with minibatch

Minibatch cost function 3

Cost function with minibatch

SGD Example

Images from Prince, Understanding Deep Learning, 2023



SGD Example

Images from Prince, Understanding Deep Learning, 2023

Gradient descent
Batch size = N

**Minibatch stochastic
gradient descent**
 $1 < \text{batch size} < N$

N is the number of training samples



What is a neural network and **how does it work?**


How do we **choose model weights?**
(i.e. how do we fit our model to data)

What are the challenges of using neural networks?


Hyperparameter / Architectural choices

- Learning Rate
- Minibatch size
- Architecture (number of nodes, number of layers, types of layers)
- Activation functions
- Weight initialization
- Regularization
- Optimizer (SGD, SGD with momentum, ADAM, etc.)
- Stopping criteria


Activation Functions



Hyperbolic Tangent



Sigmoid





Rectified linear unit (ReLU)



Increases training/prediction speed, sparse activation, reduces vanishing gradient

Image from Danijar Hafner, Quora



Tanh Function



Sigmoid Function



Rectified Linear Unit (ReLU)



Images from: <https://ml-explained.com/blog/activation-functions-explained>

Weight initialization

Set all parameters to zeros

Bad idea: leads to too much symmetry causing many gradients to be the same and the parameters will tend to all update the same way


Random numbers

Need to be neither too small nor too big. A number of heuristics exist (Xavier, He, etc.)

Batch normalization

Ensures activations are unit Gaussian at each layer, improving optimization

(faster optimization, enables higher learning rates)



We calculate the activation at each layer for each of the training samples in each minibatch

- 1** Subtract the mean of that activation value averaged across the minibatch
- 2** Divide by the standard deviation of the activation value computed across the minibatch

Regularization

L2 Regularization


L1 Regularization

Dropout


While training, keep a neuron active with some probability p , or setting it to zero otherwise.

Early Stopping

Stop the model before it has a chance to overfit
(caution: can cause underfitting)



(a) Standard Neural Net



(b) After applying dropout.

Watch your learning curves

Questions to ask:

- Learning anything?
- Still learning?
- Overfitting?
- Learning rate too low/high?




Image by Jason Brownlee

Successfully training neural networks

Advice from Andrej Karpathy: <http://karpathy.github.io/2019/04/25/recipe/>

Challenges

1. Neural network training is not plug and play. You need to understand the methods.
2. It is difficult to tell when there is a mistake

Recipe for training

1. Understand your data
2. **Setup an end-to-end training/evaluation pipeline and test simple baselines**
3. Overfit your model to the data to make sure you can do it (ensure you CAN learn)
4. Regularize the model
 1. Add data
 2. Augmentation
 3. Use dropout
 4. Early stopping
5. Tune your model (optimize hyperparameters)


Hyperparameter tuning

- Learning Rate
- Minibatch size
- Architecture (number of nodes, number of layers, types of layers)
- Activation functions
- Weight initialization
- Regularization
- Optimizer (SGD, SGD with momentum, ADAM, etc.)
- Stopping criteria

Hyperparameter Tuning


Grid Search

- Can search thoroughly
- Computationally expensive



Bayesian Optimization

- Guided sampling
- Has hyperparameters of its own



Hutter, F., Kotthoff, L. and Vanschoren, J., 2019. Automated machine learning: methods, systems, challenges. Springer.

Advice on using practical neural networks

Deep Learning Tuning Playbook

Deep Learning Tuning Playbook

This is not an officially supported Google product.

Varun Godbole[†], George E. Dahl[†], Justin Gilmer[†], Christopher J. Shallue[‡], Zachary Nado[†]

[†] Google Research, Brain Team

[‡] Harvard University

Table of Contents

- Who is this document for?
- Why a tuning playbook?
- Guide for starting a new project
 - Choosing the model architecture
 - Choosing the optimizer
 - Choosing the batch size
 - Choosing the initial configuration
- A scientific approach to improving model performance
 - The incremental tuning strategy
 - Exploration vs exploitation
 - Choosing the goal for the next round of experiments
 - Designing the next round of experiments
 - Determining whether to adopt a training pipeline change or hyperparameter configuration
 - After exploration concludes
- Determining the number of steps for each training run
 - Deciding how long to train when training is not compute-bound