

Deep Learning

Deep learning

Representation learning with a hierarchy of concepts

Those concepts are represented by layers in a neural network model

Unsupervised models

- Autoencoders

Supervised models

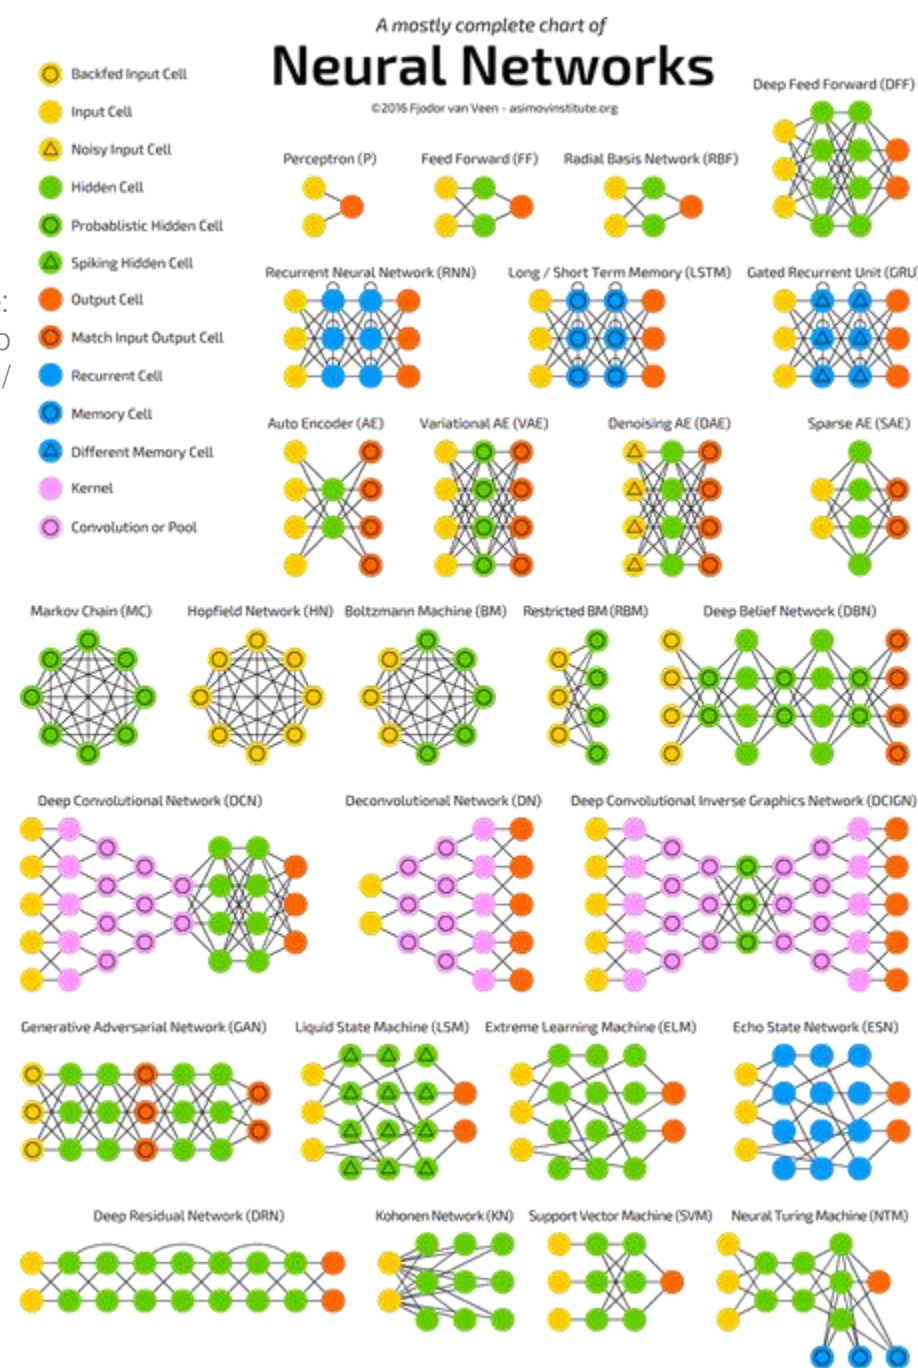
- Image analysis:
 - Convolutional Neural Networks (CNNs)
 - ...and Vision Transformers (ViTs)
- Text analysis and NLP
 - Transformers
- Timeseries analysis:
 - Transformers
 - Recurrent Neural Networks (RNNs)

Generative models

- Generative Adversarial Networks (GANs)
- Diffusion Models (e.g. DALL-E 2, Stable Diffusion)
- Generative Pre-trained Transformer (GPT)

Types of Deep Learning Tools

Azimov Institute:
<http://www.asimovinstitute.org/neural-network-zoo/>



Autoencoders

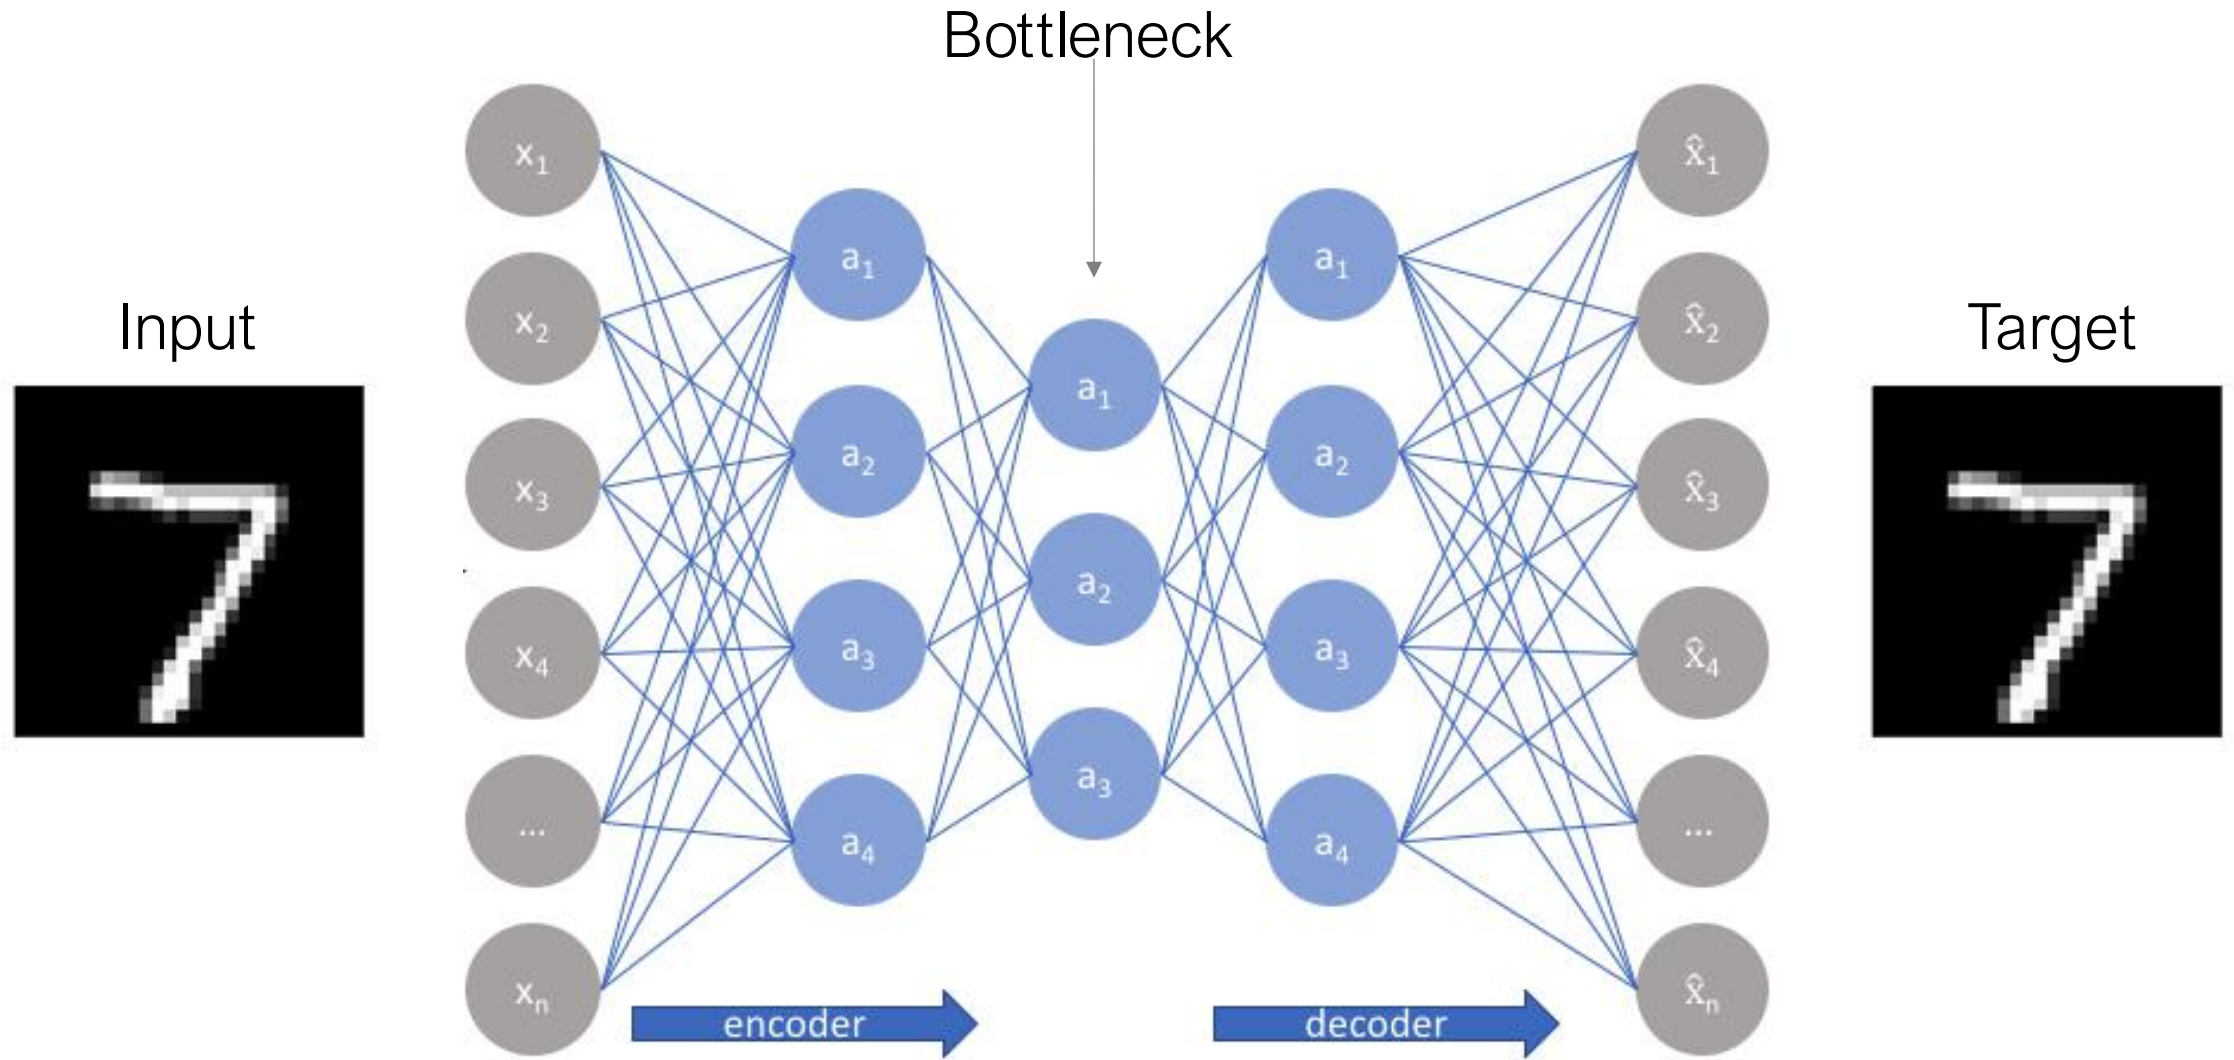



Image from: <https://www.jeremyjordan.me/autoencoders/>



Autoencoders

Our goal is often to develop a good **encoder** that represents our features well




Convolutional Neural Networks

AlexNet



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

Key



Convolutional Neural Networks

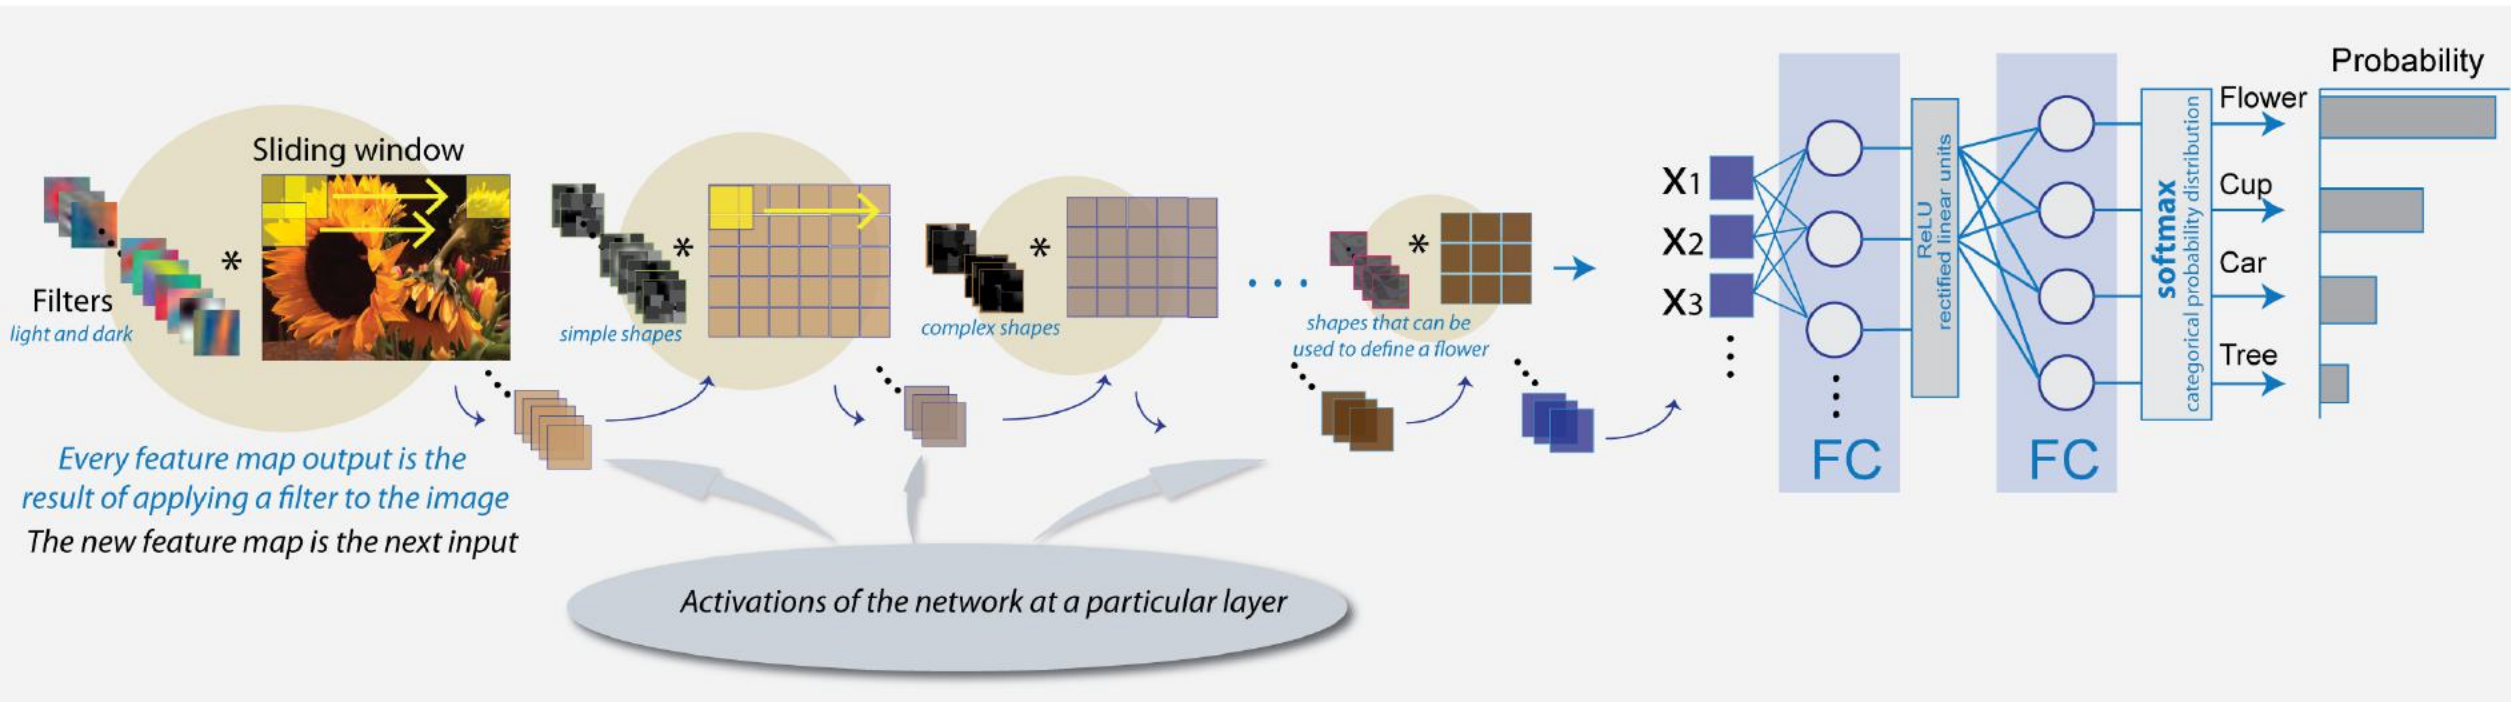
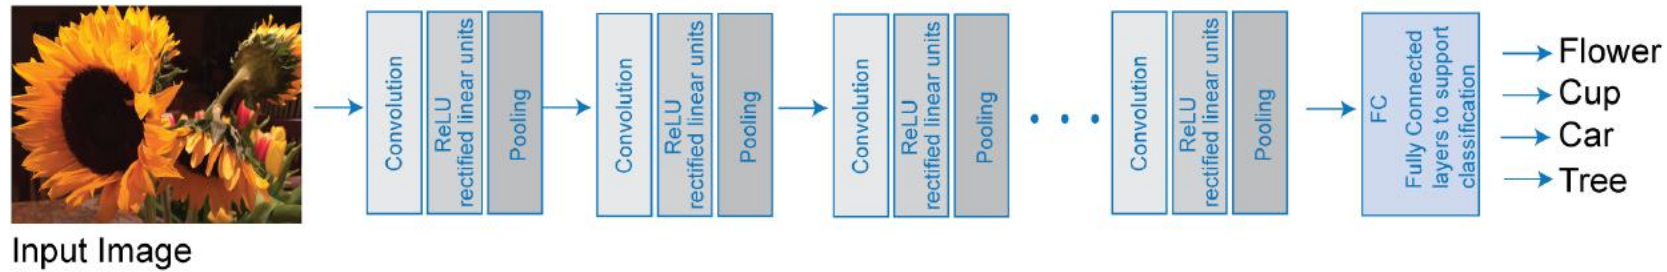




Image from the Mathworks

Convolution in 1 dimension


a)




b)



c)



d)



$$z_2 = x_1 w_1 + x_2 w_2 + x_3 w_3$$

$$z_3 = x_2 w_1 + x_3 w_2 + x_4 w_3$$

$$z_1 = 0 w_1 + x_1 w_2 + x_2 w_3$$

$$z_1 = x_1 w_1 + x_2 w_2 + x_3 w_3$$

Data: \mathbf{x}

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6



Weights: \mathbf{w}

1	1	1
0	0	0
-1	-1	-1

=

Output: $\mathbf{x} * \mathbf{w}$

2D Convolution

Data: x

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

$*$

$=$

Output: $x * w$

Computing
one output
value:

$$\begin{aligned} & 1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\ & 0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\ & (-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5 \end{aligned}$$

2D Convolution

Data: x

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $x * w$

-6			

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5 = -6
 \end{aligned}$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11		

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 2 + 1 \cdot 5 + 1 \cdot 1 + \\
 &0 \cdot 2 + 0 \cdot 3 + 0 \cdot 2 + \\
 &(-1) \cdot 5 + (-1) \cdot 5 + (-1) \cdot 9 = -11
 \end{aligned}$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11	-12	

Computing
one output
value:

$$1 \cdot 5 + 1 \cdot 1 + 1 \cdot 4 +$$

$$0 \cdot 3 + 0 \cdot 2 + 0 \cdot 0 +$$

$$(-1) \cdot 5 + (-1) \cdot 9 + (-1) \cdot 8 = \mathbf{-12}$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11	-12	-11

Computing
one output
value:

$$1 \cdot 1 + 1 \cdot 4 + 1 \cdot 2 +$$

$$0 \cdot 2 + 0 \cdot 0 + 0 \cdot 0 +$$

$$(-1) \cdot 9 + (-1) \cdot 8 + (-1) \cdot 1 = -11$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11	-12	-11
-7			

Computing
one output
value:

$$1 \cdot 0 + 1 \cdot 2 + 1 \cdot 3 +$$

$$0 \cdot 4 + 0 \cdot 5 + 0 \cdot 5 +$$

$$(-1) \cdot 6 + (-1) \cdot 3 + (-1) \cdot 4 = \mathbf{-7}$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

6 x 6

Weights: w

1	1	1
0	0	0
-1	-1	-1

3 x 3



Output: $X * w$

-6	-11	-12	-11
-7	-2	-2	-4
4	1	-2	1
3	-4	-6	-10

4 x 4

2D Convolution

What features do layers respond to?



Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>

Features


Dataset Examples show us what neurons respond to in practice



Optimization isolates the causes of behavior from mere correlations. A neuron may not be detecting what you initially thought.



Baseball—or stripes?
mixed4a, Unit 6



Animal faces—or snouts?
mixed4a, Unit 240



Clouds—or fluffiness?
mixed4a, Unit 453



Buildings—or sky?
mixed4a, Unit 492

Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>


Resources on Visualization of Features

Feature visualization: <https://distill.pub/2017/feature-visualization/>


Building blocks of interpretability: <https://distill.pub/2018/building-blocks/>

Activation Activation Atlases: <https://distill.pub/2019/activation-atlas/>

Convolution Example




Convolution Layer



From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions




$$\text{Parameters} = (5 \times 5 \times 3) \times 6 = 450$$

$$(5 \times 5 \times 6) \times 10 = 1,500$$


From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

1 x 1 Convolution Explained




1x1 CONV
with 32 filters

→
(each filter has size
1x1x64, and performs a
64-dimensional dot
product)




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

Max Pooling




max pool with 2x2 filters
and stride 2




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


Fully Connected Layer




AlexNet




Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.



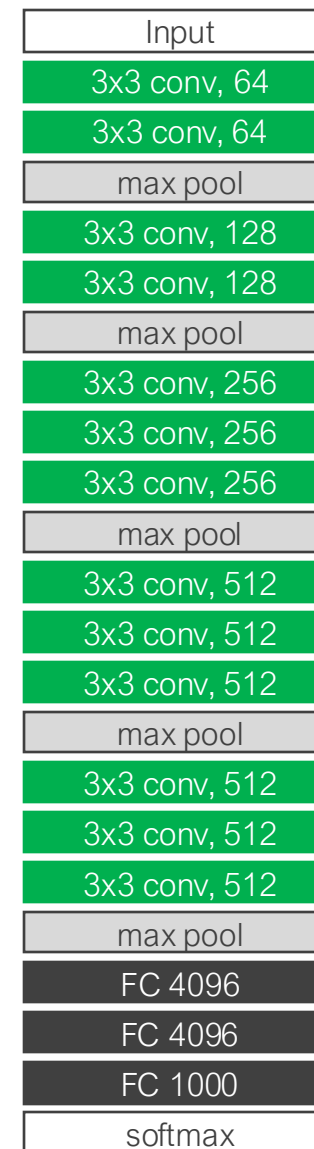
Key



AlexNet
(2012)



VGG16
(2014)




VGG19
(2014)



Note: an activation function is applied to the output of each layer

Fewer layers, larger filters


Key



CNN Architectures

Adapted from Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


Residual Networks (ResNet)



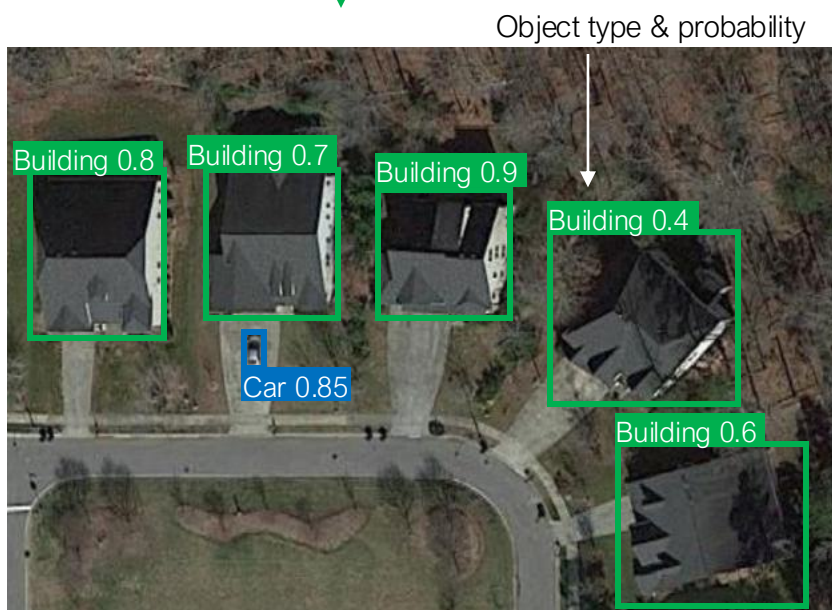
Skip Connection enable faster convergence, more effectively backpropagate the error signal (avoiding vanishing gradients)



Scene
classification



Object
detection



Deep Learning

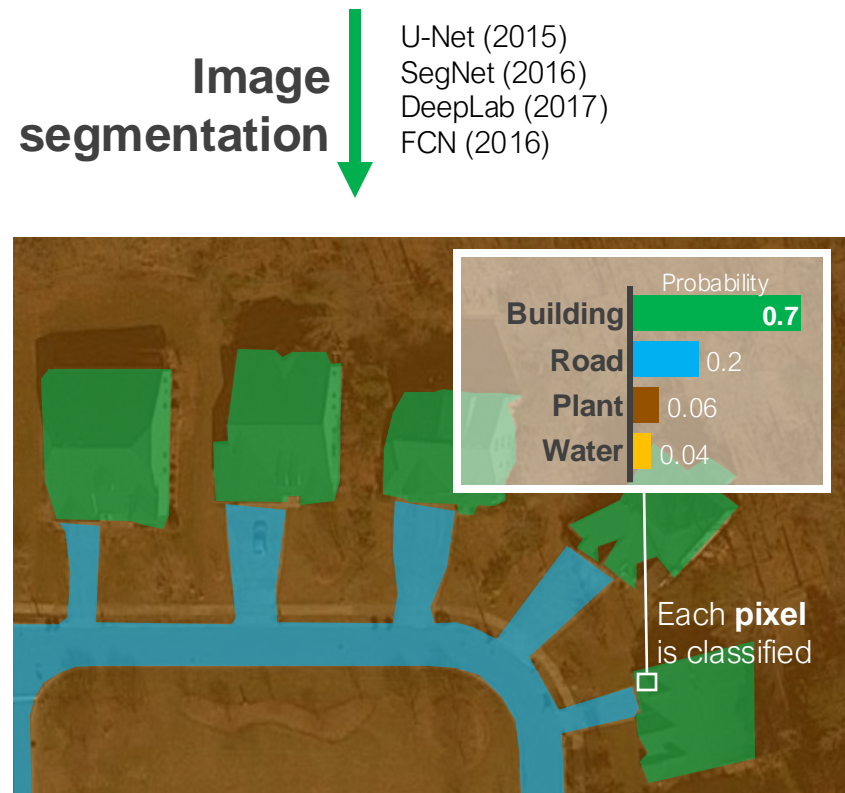
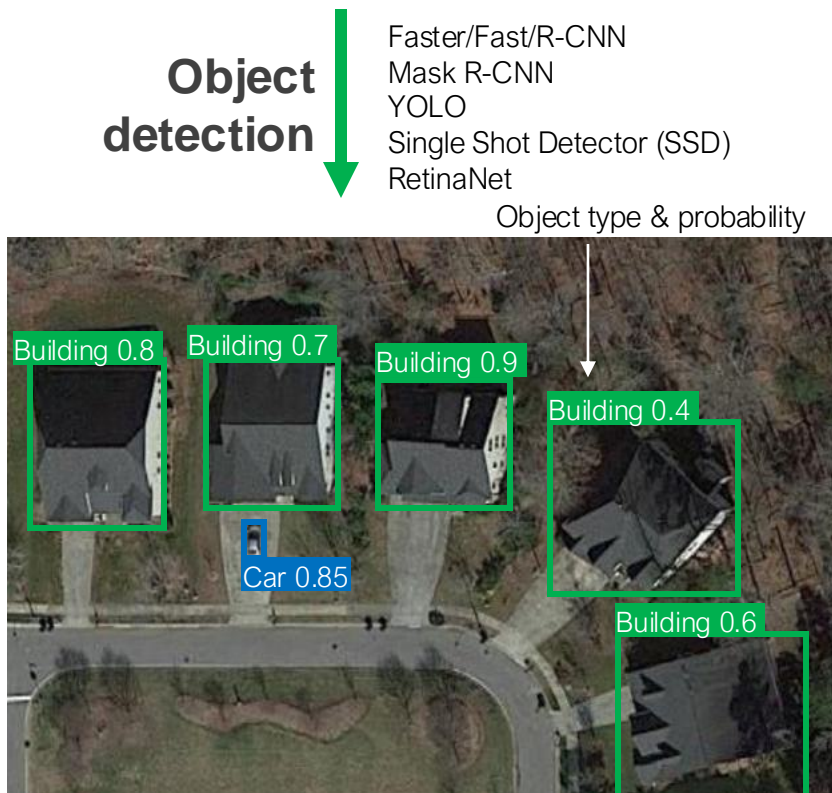



Image
segmentation



Lecture 14

32



ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

Fei-Fei Li et al. 2010 ([link](#))


Competition at:
Conference on Computer Vision and Pattern Recognition (CVPR)

USED FOR MODEL PRETRAINING

AlexNet
(Hinton, Sutskever, and Krizhevsky)

Human error ~5%

perfect





Source: Quartz, [link](#)

David Yanofsky / Quartz

Data: ImageNet

Deep Learning Models Compared



Models compared for ImageNet
Many of these models are available through Keras ([link](#))

A. Canziani, E. Culurciello and A. Paszke, "Evaluation of neural network architectures for embedded systems," *2017 IEEE International Symposium on Circuits and Systems (ISCAS)*, Baltimore, MD, 2017, pp. 1-4.

Deep learning frameworks

Tensorflow ([link](#))

Framework for implementing graphical models, such as neural networks



Keras ([link](#))

Wrapper for Tensorflow to make coding easier: higher level and excellent API



PyTorch ([link](#))

Framework for implementing graphical models, such as neural networks



KERAS DEMO

Generative Adversarial Networks




Image from: <https://skymind.ai/wiki/generative-adversarial-network-gan>