

# Deep Learning II

The restaurant refused to serve me a ham sandwich because it only cooks vegetarian food. In the end, they just gave me two slices of bread. Their ambiance was just as good as the food and service.

The restaurant refused to serve me a ham sandwich because **it** only cooks vegetarian food. In the end, **they** just gave me two slices of bread. **Their** ambiance was just as good as the food and service.

What does each **bolded** word refer to?

Is this a positive review or a negative one?

What kind of food do they serve at this restaurant?

# Types of NLP problems

## Text Classification

- Sentiment Analysis
- Topic categorization
- Extractive question answering

## Sequence-to-sequence

- Machine translation (text-to-text)
- Summarization (text-to-text)
- Speech-to-text, text-to-speech

## • Generative models

- Chatbots
- Generative question answering
- Code generation

### Encoder models

**BERT** (Bidirectional Encoder Representations from Transformers)

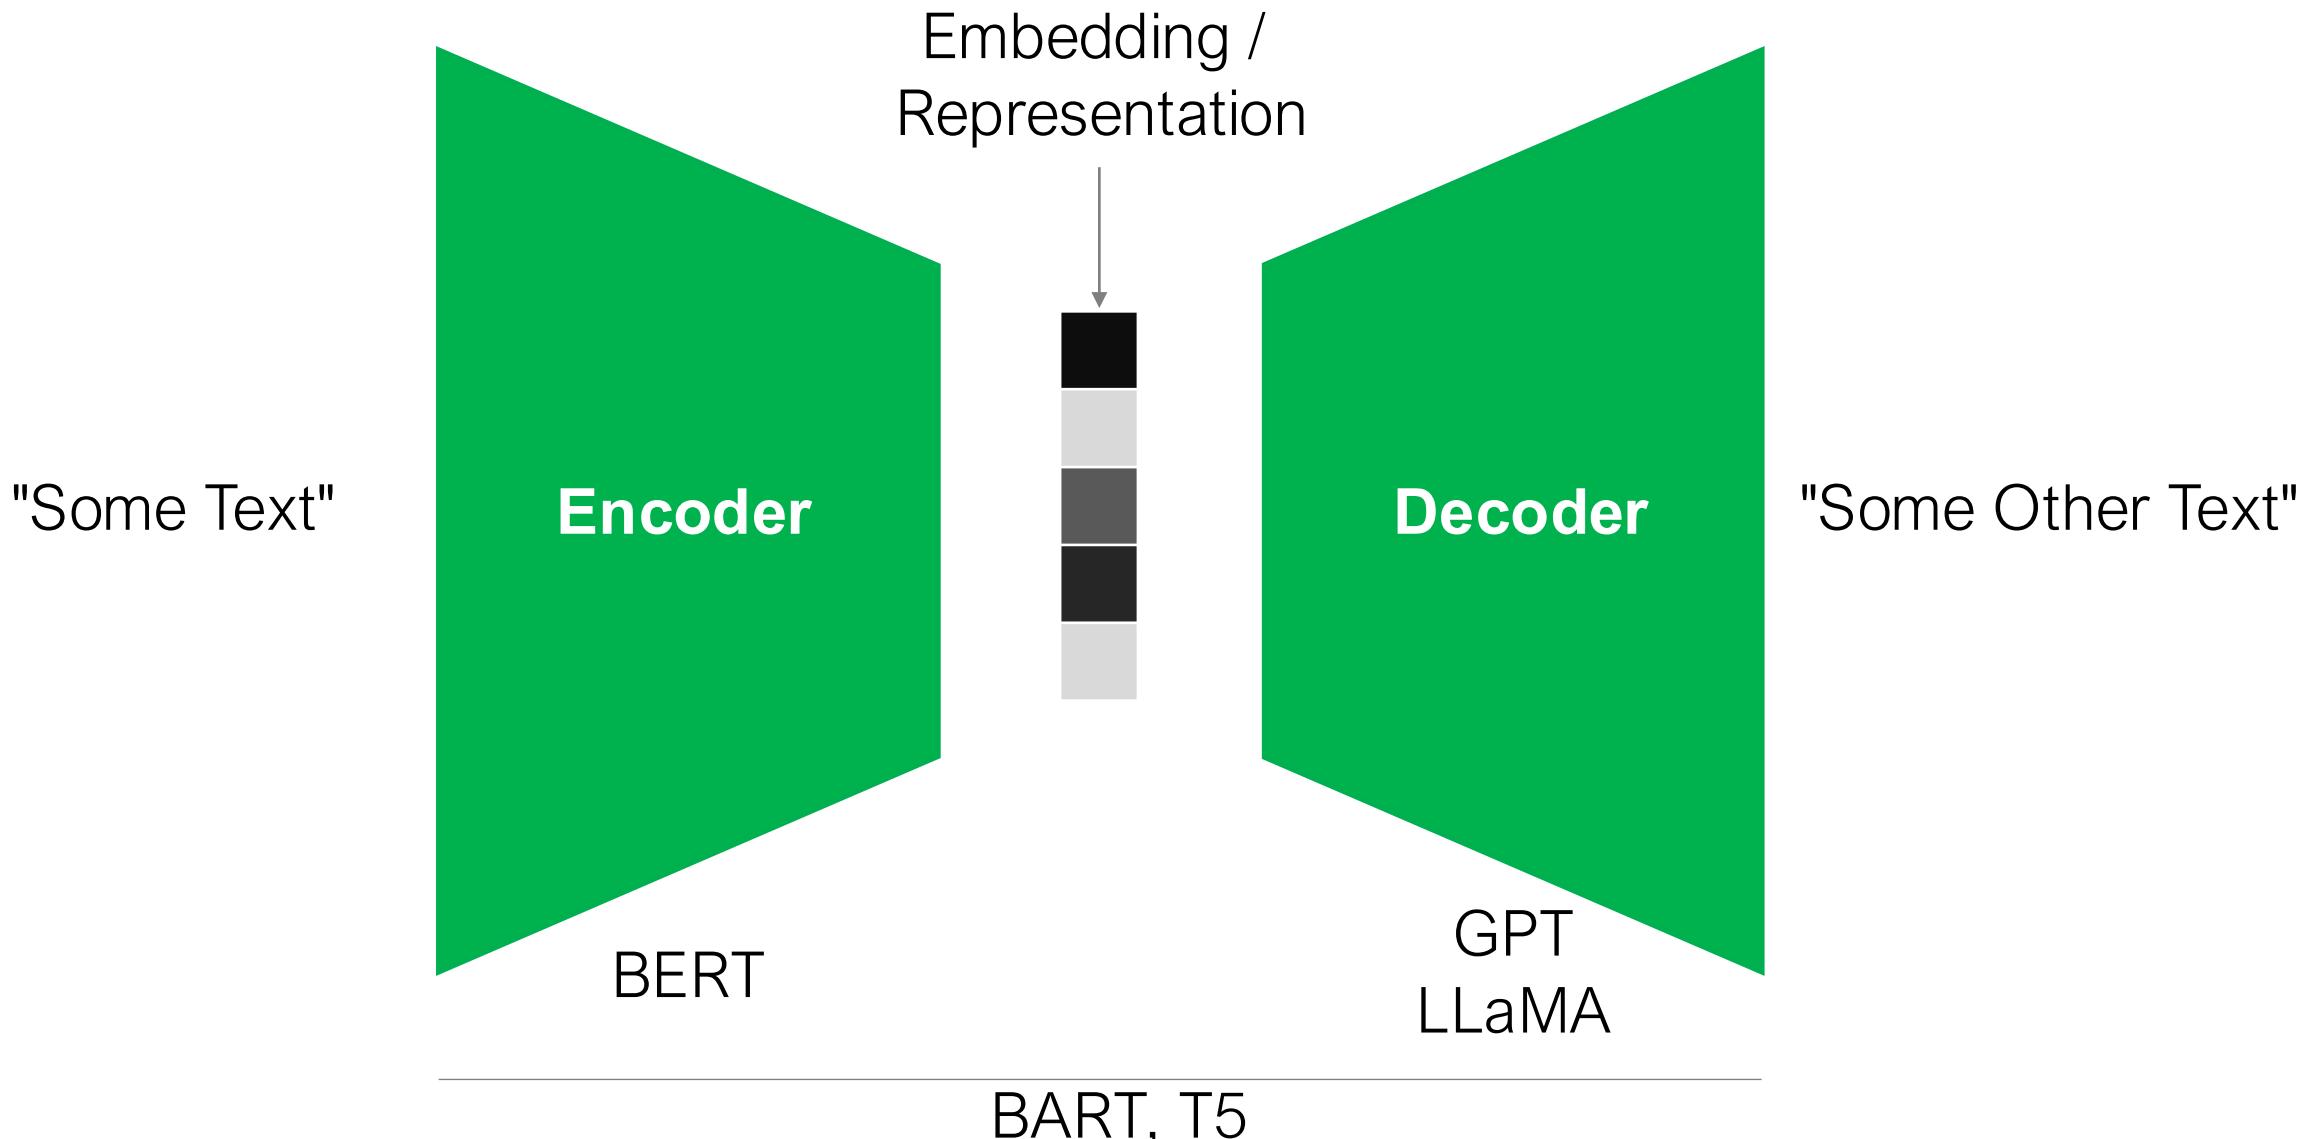
### Encoder-decoder models

**BART** (Bidirectional and Auto-Regressive Transformers)  
**T5** (Text-to-Text Transfer Transformer)

### Decoder models

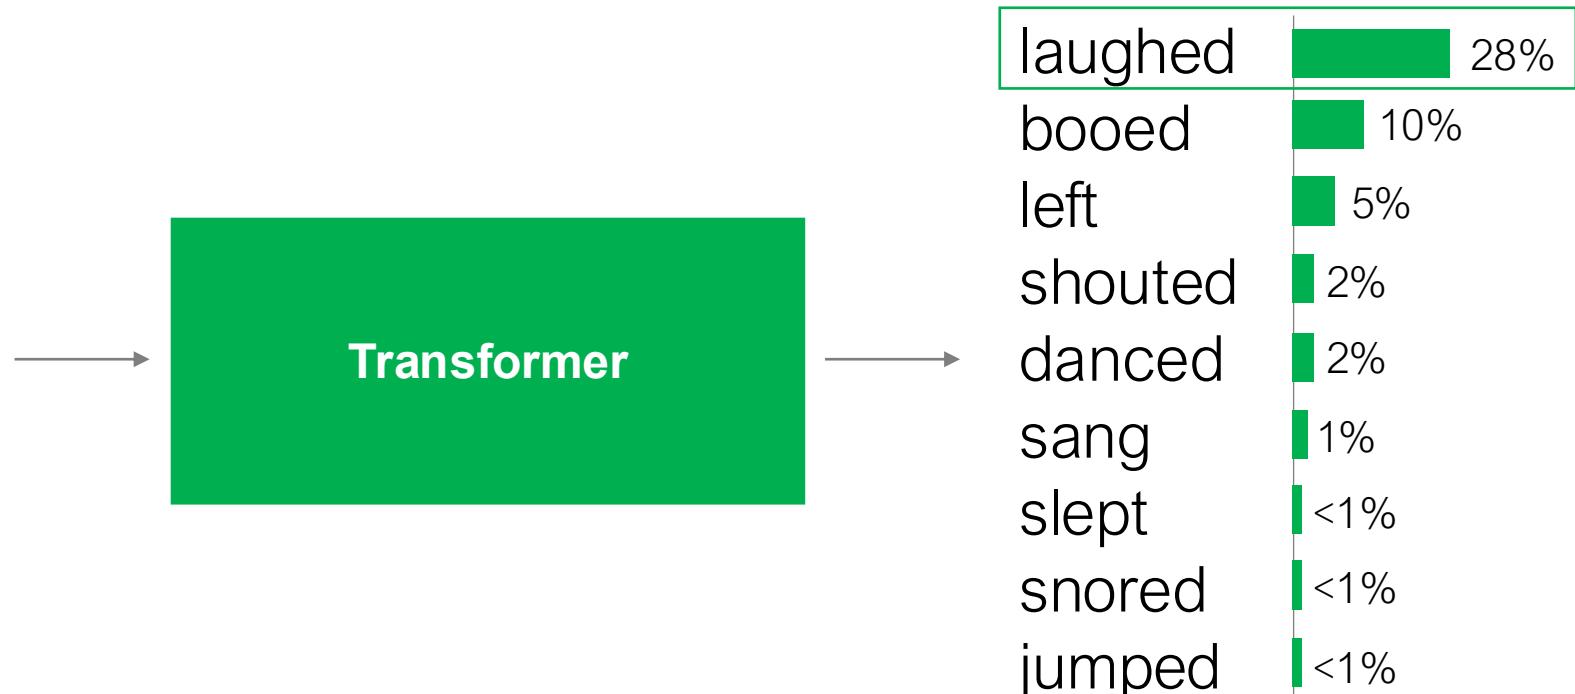
**GPT** (Generative Pre-trained Transformer)  
**LLaMA** (Large Language Model Meta AI)

# Types of Transformers



# Example: Predict the next word

After hearing the terrible pun, the data science students, to the surprise of their instructor, all suddenly \_\_\_\_\_



If we repeat this again and again, we can generate text

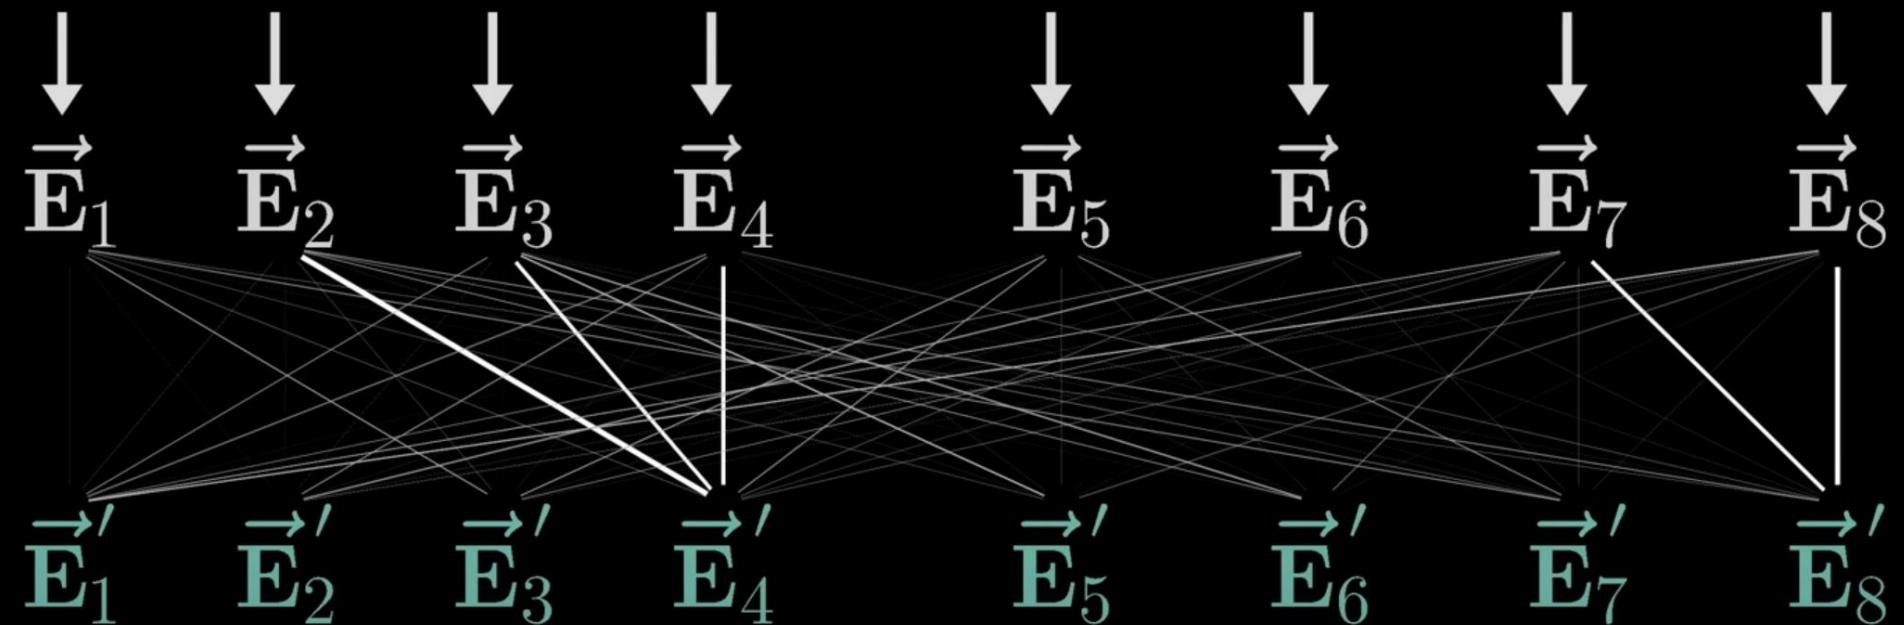
# Transformer intuition

The quick brown fox jumped over the lazy dog





a **fluffy** **blue** **creature** roamed the **verdant** **forest**



Create improved embeddings by incorporating meaning from other words



Image from 3Blue1Brown ([link](#))

# Transformer steps and components

1. Tokenization – convert text to numbers
2. Position embedding – encode the order of the tokens
3. Self-attention – enable learning from context
  1. Queries
  2. Keys
  3. Values
4. Layer normalization
5. Multilayer perceptron

# Tokenization

Converts text to numerical IDs for each token

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h	l	u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h	l	u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

c) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

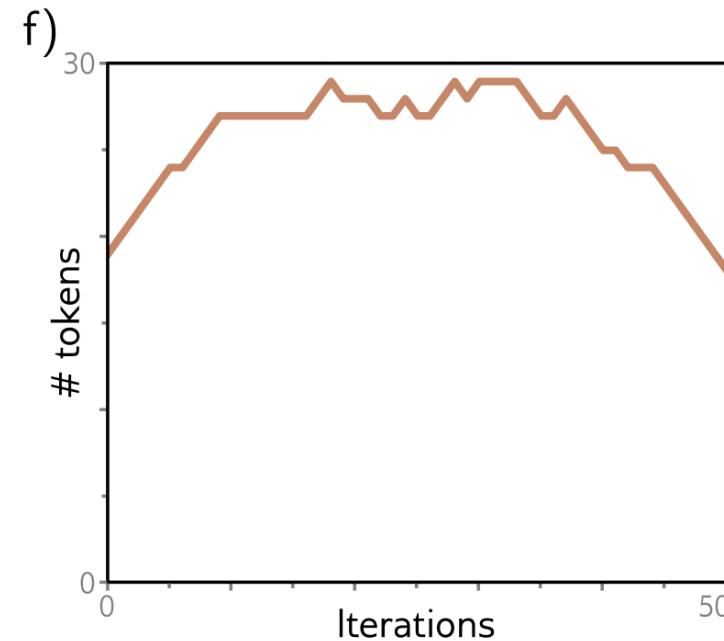
_	se	a	e	t	o	h	l	u	b	d	e	w	c	s	f	i	m	n	p	r
21	13	12	12	11	8	6	6	4	3	3	3	3	2	2	1	1	1	1	1	1

⋮ ⋮

d) see\_sea\_e\_b\_l\_w\_a\_could\_hat\_he\_o\_t\_t\_the\_to\_u\_a\_d\_f\_m\_n\_p\_s\_sailor\_to

⋮ ⋮ ⋮

e) see\_sea\_could\_he\_the\_a\_all\_blue\_bottom\_but\_deep\_of\_sailor\_that\_to\_was\_went\_what\_



Token indices are transformed into a sequence of embeddings (the embeddings are learned)

"an aardvark  
ate an ant"

# Producing embeddings

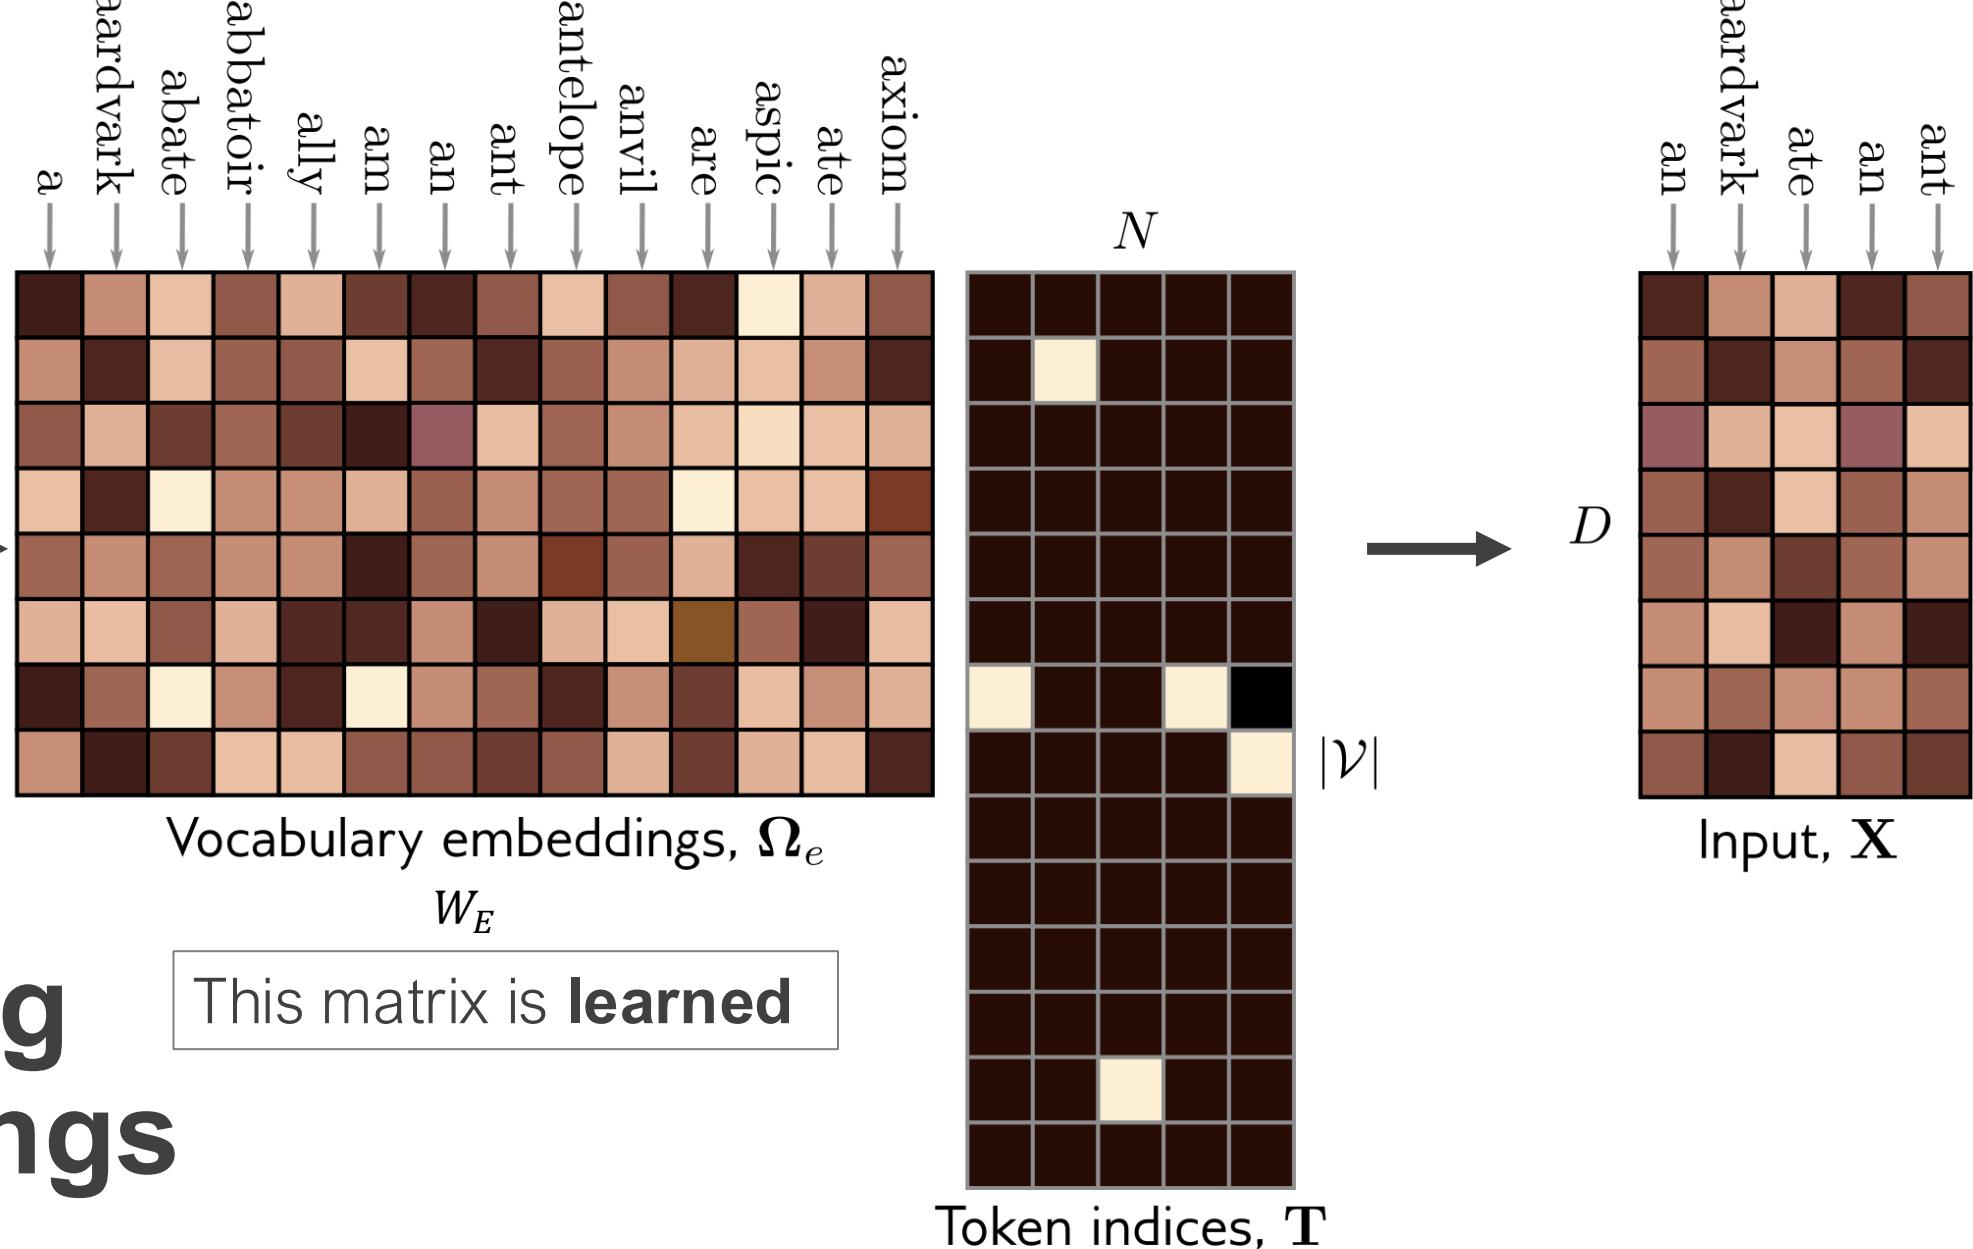
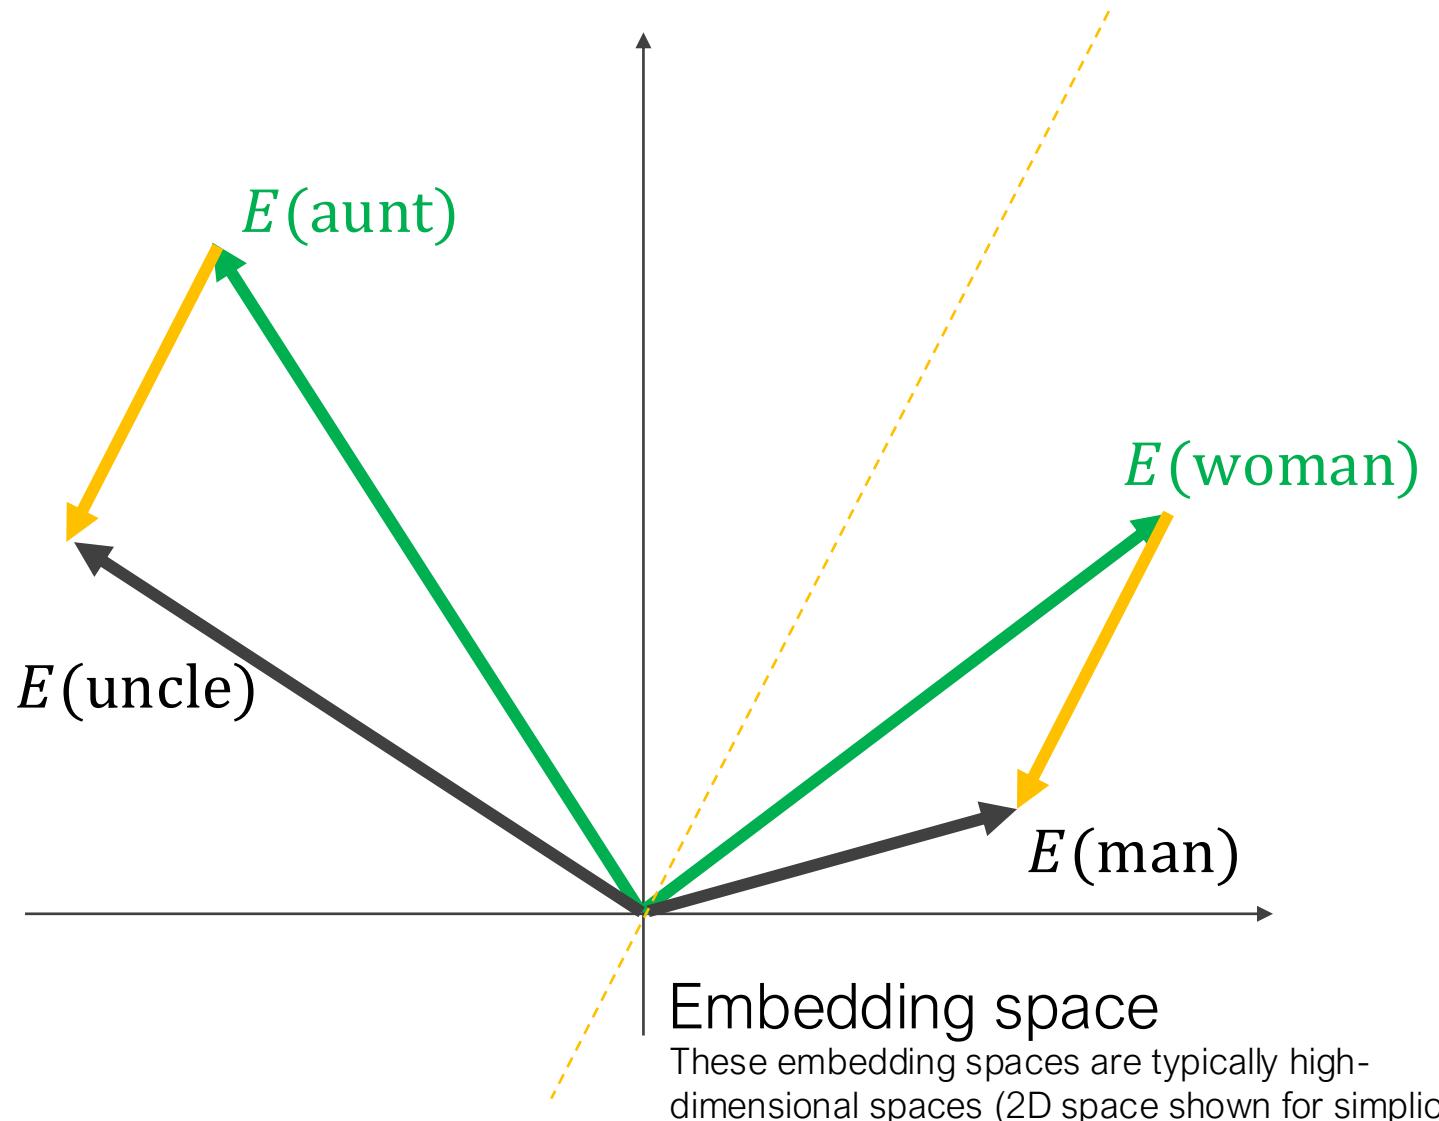


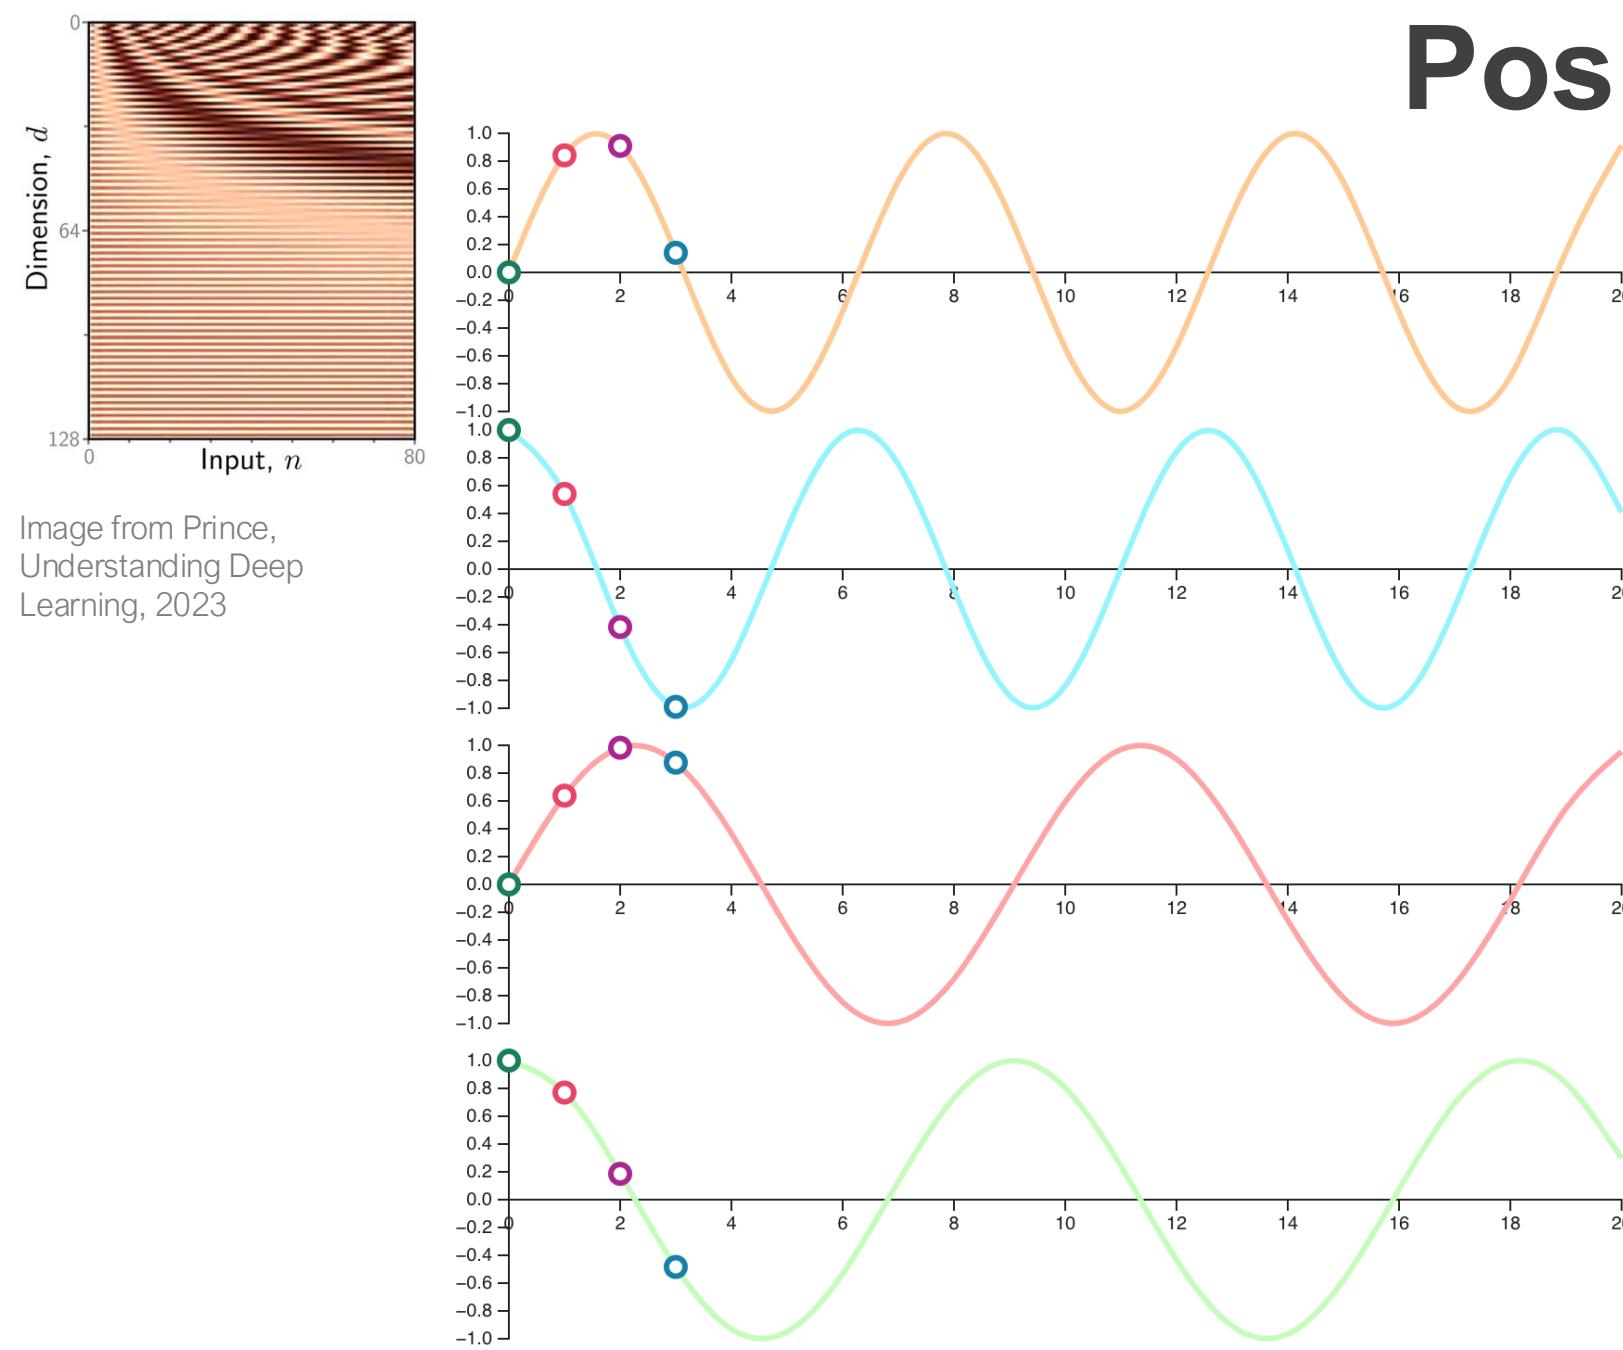
Image from Prince, Understanding Deep Learning, 2023



These embeddings spaces start off random and through training gain semantic meaning

Directions in these high dimensional spaces may begin to represent concepts

# Positional Encoding



$p0$	$p1$	$p2$	$p3$	$i=0$
0.000	0.841	0.909	0.141	
1.000	0.540	-0.416	-0.990	$i=1$
0.000	0.638	0.983	0.875	$i=2$
1.000	0.770	0.186	-0.484	$i=3$

## Positional Encoding

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

## Settings: $d = 50$

The value of each positional encoding depends on the *position* (*pos*) and *dimension* (*d*). We calculate result for every *index* (*i*) to get the whole vector.

Image source: <https://erdem.pl/2021/05/understanding-positional-encoding-in-transformers>

"quick brown fox"

## Tokenization

[423,23,342] →

## Token (word) embeddings

1.6
3.2
-0.4
-4.9
0.42
1.6

+

-1.0
0.50
-0.10
0.0
0.35
-0.21

## Token and position embedding

0.6
3.7
-0.50
-4.9
0.77
1.39

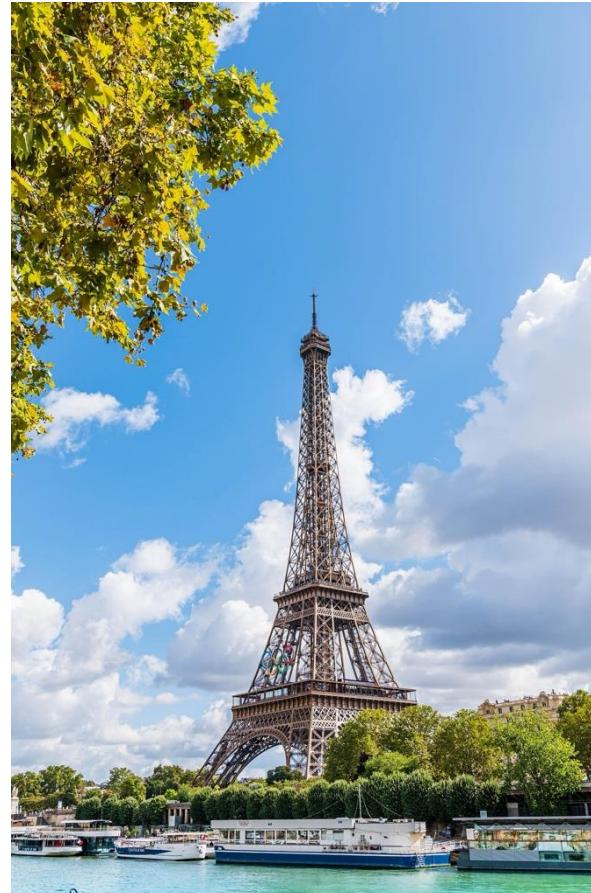
What it is and where it is

# Context typically adds meaning

tower



Eiffel tower

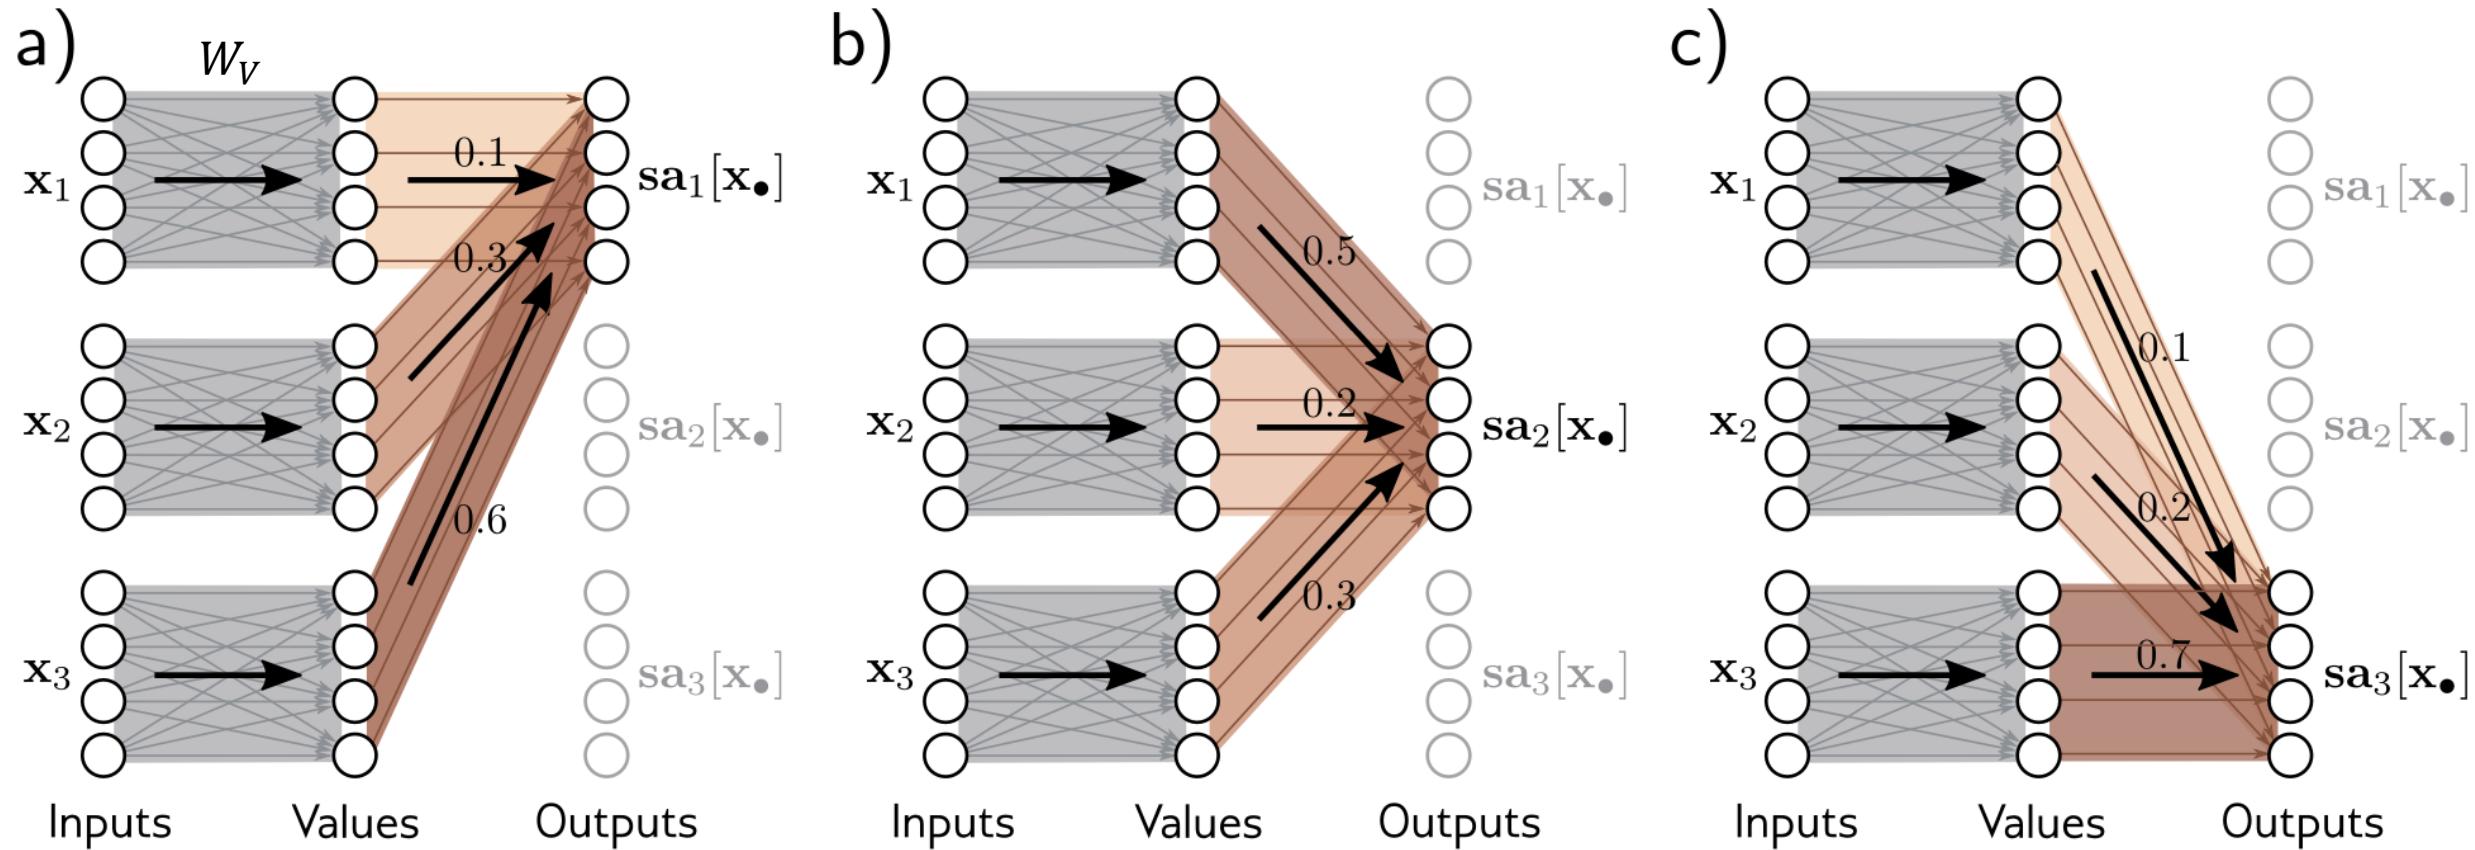


miniature Eiffel tower



Concept from 3Brown1Blue  
Images from Pixabay

# Self-attention



Inputs are the word/token embeddings

# Self-attention

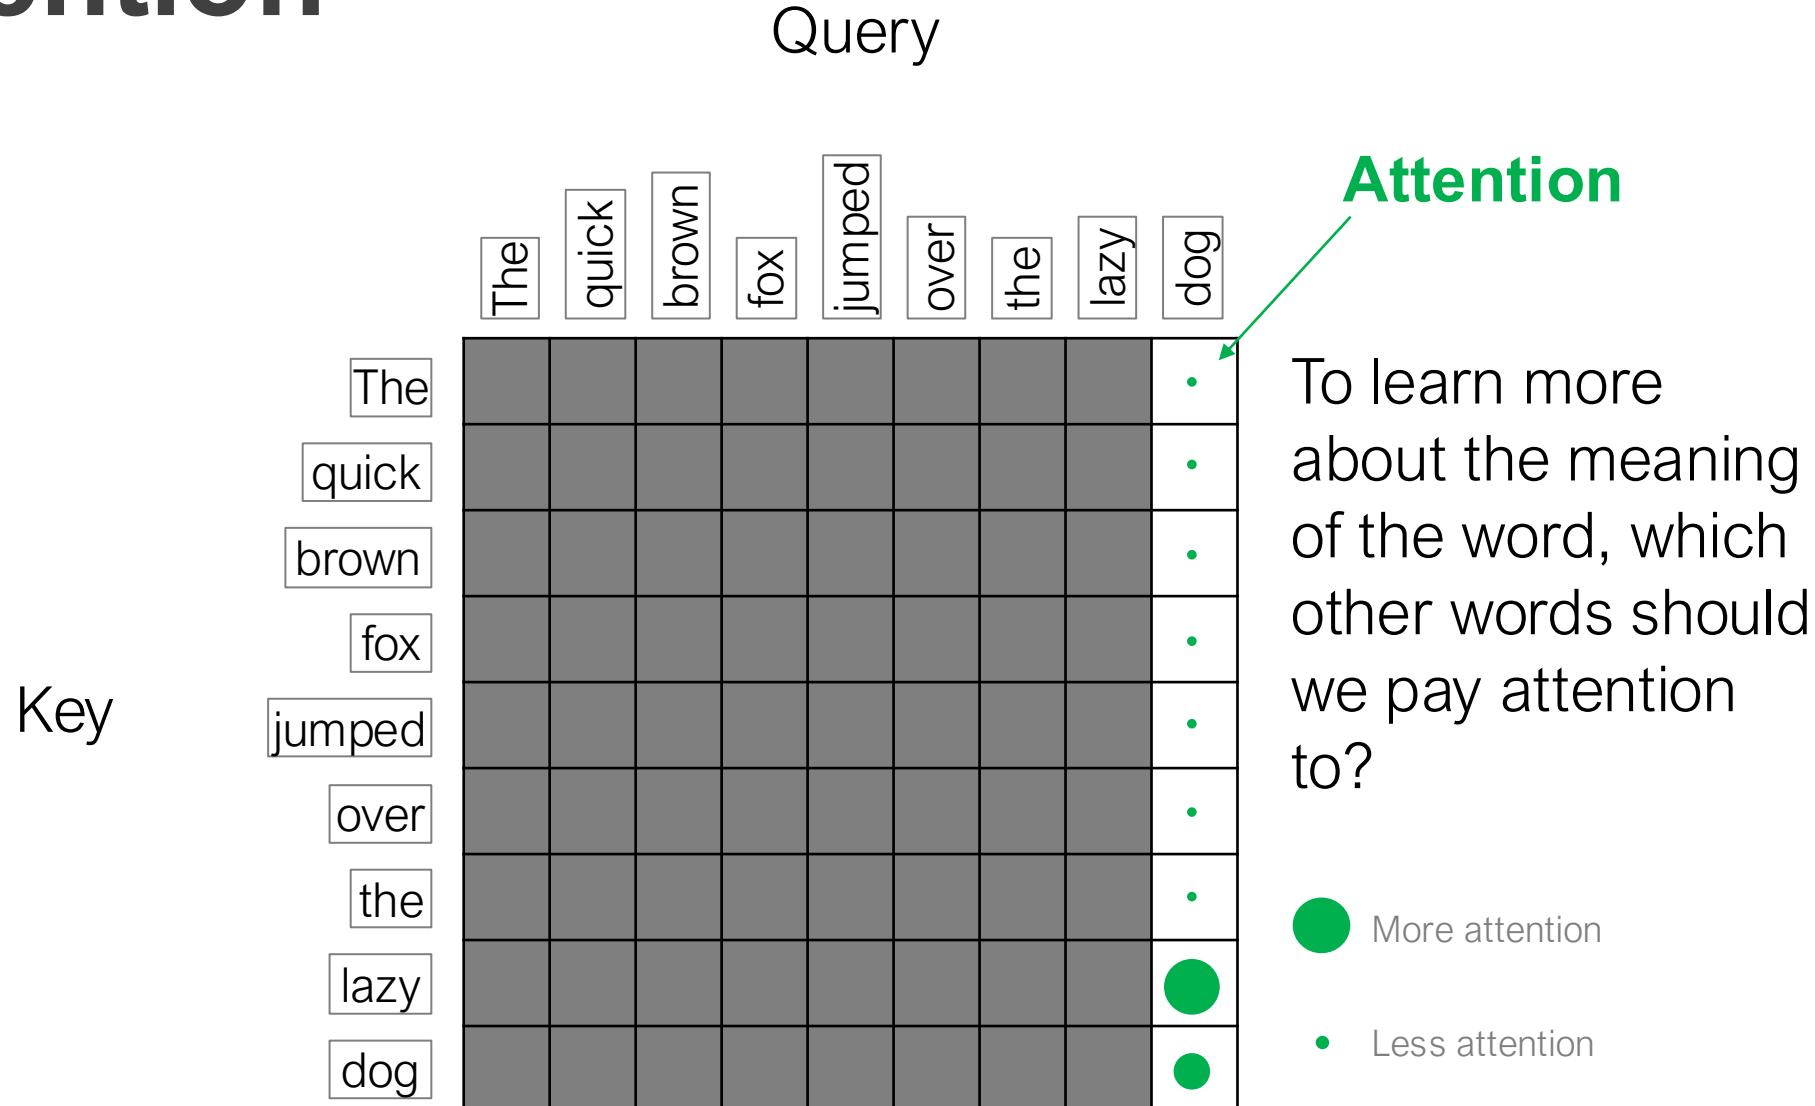
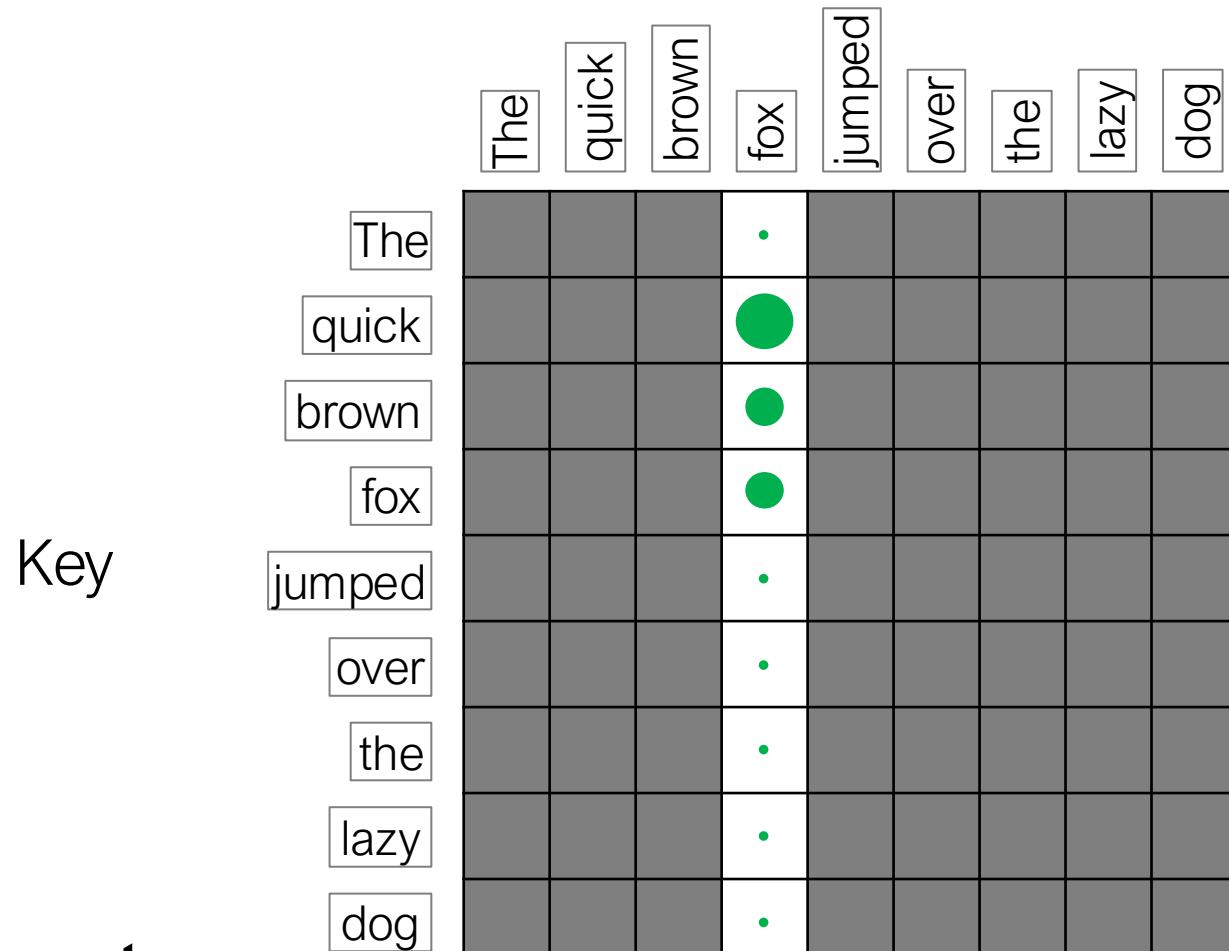


Figure adapted from 3Blue1Brown ([link](#))

# Self-attention

Query



To learn more about the meaning of the word, which other words should we pay attention to?

● More attention

• Less attention

Figure adapted from 3Blue1Brown ([link](#))

How do we compute self-attention?

# Query

The quick brown fox jumped over the lazy dog

$\mathbf{x}_1 \quad \mathbf{x}_2 \quad \mathbf{x}_3 \quad \mathbf{x}_4 \quad \mathbf{x}_5 \quad \mathbf{x}_6 \quad \mathbf{x}_7 \quad \mathbf{x}_8 \quad \mathbf{x}_9$

$\downarrow W_q \quad \downarrow W_q$

$\mathbf{q}_1 \quad \mathbf{q}_2 \quad \mathbf{q}_3 \quad \mathbf{q}_4 \quad \mathbf{q}_5 \quad \mathbf{q}_6 \quad \mathbf{q}_7 \quad \mathbf{q}_8 \quad \mathbf{q}_9$

Keys  $\mathbf{k}_i = W_k \mathbf{x}_i$

$K = [\mathbf{k}_1, \mathbf{k}_2, \dots, \mathbf{k}_9]$

The  $\mathbf{x}_1 \xrightarrow{W_k} \mathbf{k}_1$   
 quick  $\mathbf{x}_2 \xrightarrow{W_k} \mathbf{k}_2$   
 brown  $\mathbf{x}_3 \xrightarrow{W_k} \mathbf{k}_3$   
 fox  $\mathbf{x}_4 \xrightarrow{W_k} \mathbf{k}_4$   
 jumped  $\mathbf{x}_5 \xrightarrow{W_k} \mathbf{k}_5$   
 over  $\mathbf{x}_6 \xrightarrow{W_k} \mathbf{k}_6$   
 the  $\mathbf{x}_7 \xrightarrow{W_k} \mathbf{k}_7$   
 lazy  $\mathbf{x}_8 \xrightarrow{W_k} \mathbf{k}_8$   
 dog  $\mathbf{x}_9 \xrightarrow{W_k} \mathbf{k}_9$

$\mathbf{k}_1^T \mathbf{q}_1$	$\mathbf{k}_1^T \mathbf{q}_2$	$\mathbf{k}_1^T \mathbf{q}_3$							
$\mathbf{k}_2^T \mathbf{q}_1$	$\mathbf{k}_2^T \mathbf{q}_2$	$\mathbf{k}_2^T \mathbf{q}_3$							
$\mathbf{k}_3^T \mathbf{q}_1$	$\mathbf{k}_3^T \mathbf{q}_2$	$\mathbf{k}_3^T \mathbf{q}_3$							
			..						

Queries  $\mathbf{q}_i = W_q \mathbf{x}_i$

$Q = [\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_9]$

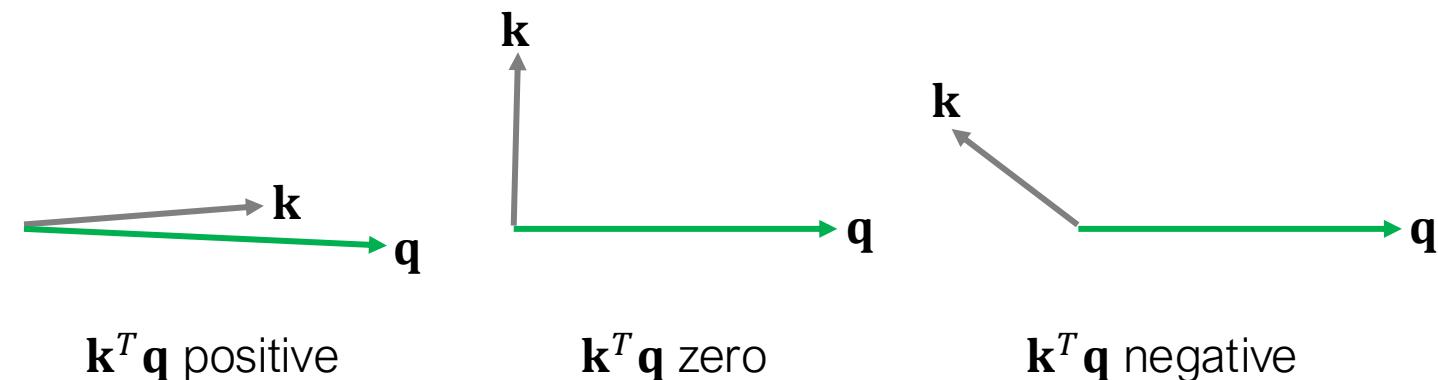
Figure adapted from 3Blue1Brown ([link](#))

# Queries and Keys

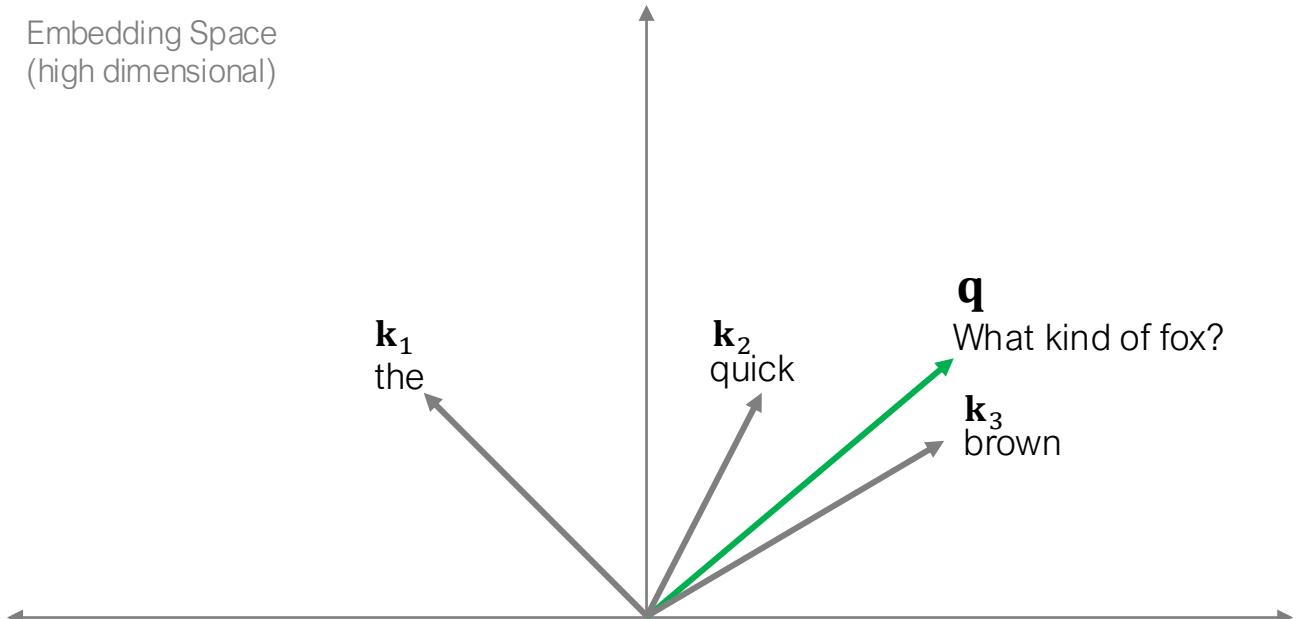
The quick brown fox jumped over the lazy dog

We measure the similarity between the key ( $\mathbf{k}$ ) and the query ( $\mathbf{q}$ ) through the dot product:

$$\mathbf{k} \cdot \mathbf{q} = \mathbf{k}^T \mathbf{q}$$

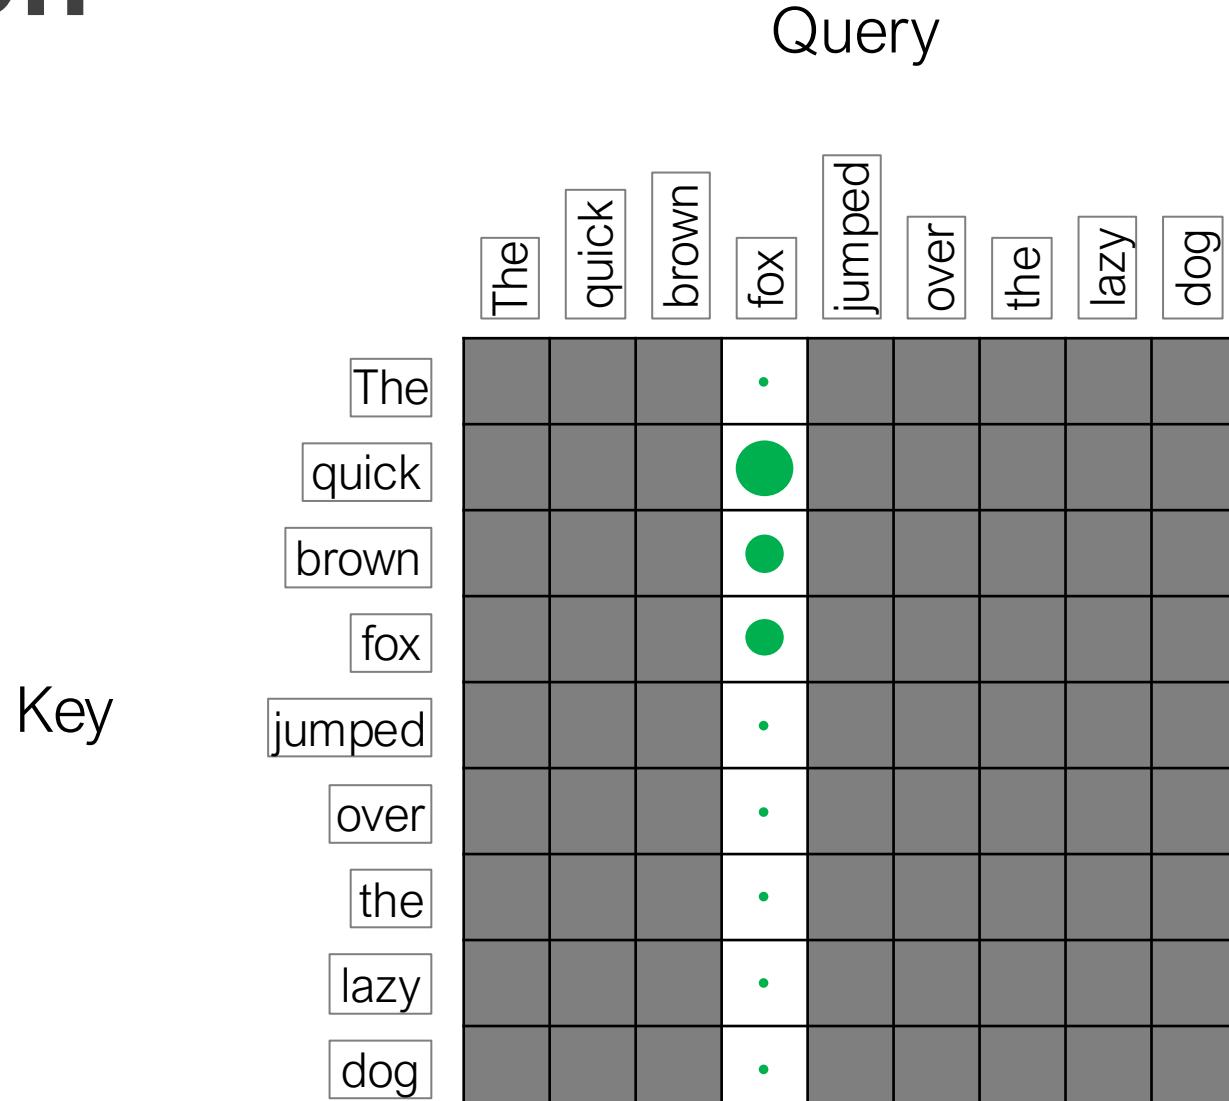


Embedding Space  
(high dimensional)



This tells the model what to pay attention to for a particular concept

# Attention



To learn more about the meaning of the word, which other words should we pay attention to?

● More attention

• Less attention

Figure adapted from 3Blue1Brown ([link](#))

# Attention

Keys  $\mathbf{k}_i = W_K \mathbf{x}_i$

$$K = [k_1, k_2, \dots, k_9]$$

The	$\mathbf{x}_1 \xrightarrow{W_K} \mathbf{k}_1$
quick	$\mathbf{x}_2 \xrightarrow{W_K} \mathbf{k}_2$
brown	$\mathbf{x}_3 \xrightarrow{W_K} \mathbf{k}_3$
fox	$\mathbf{x}_4 \xrightarrow{W_K} \mathbf{k}_4$
jumped	$\mathbf{x}_5 \xrightarrow{W_K} \mathbf{k}_5$
over	$\mathbf{x}_6 \xrightarrow{W_K} \mathbf{k}_6$
the	$\mathbf{x}_7 \xrightarrow{W_K} \mathbf{k}_7$
lazy	$\mathbf{x}_8 \xrightarrow{W_K} \mathbf{k}_8$
dog	$\mathbf{x}_9 \xrightarrow{W_K} \mathbf{k}_9$

## Query

The	quick	brown	fox	jumped	Over	the	lazy	dog
$\mathbf{x}_1$	$\mathbf{x}_2$	$\mathbf{x}_3$	$\mathbf{x}_4$	$\mathbf{x}_5$	$\mathbf{x}_6$	$\mathbf{x}_7$	$\mathbf{x}_8$	$\mathbf{x}_9$
$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$	$\downarrow W_Q$
$\mathbf{q}_1$	$\mathbf{q}_2$	$\mathbf{q}_3$	$\mathbf{q}_4$	$\mathbf{q}_5$	$\mathbf{q}_6$	$\mathbf{q}_7$	$\mathbf{q}_8$	$\mathbf{q}_9$
$\mathbf{k}_1^T \mathbf{q}_1$	$\mathbf{k}_1^T \mathbf{q}_2$	$\mathbf{k}_1^T \mathbf{q}_3$		0				
$\mathbf{k}_2^T \mathbf{q}_1$	$\mathbf{k}_2^T \mathbf{q}_2$	$\mathbf{k}_2^T \mathbf{q}_3$	0.4					
$\mathbf{k}_3^T \mathbf{q}_1$	$\mathbf{k}_3^T \mathbf{q}_2$	$\mathbf{k}_3^T \mathbf{q}_3$	0.4					
			0.2					
			0					
			0					
			0					
			0					

Queries  $\mathbf{q}_i = W_Q \mathbf{x}_i$

$$Q = [q_1, q_2, \dots, q_9]$$

## Attention weights

$$a[\mathbf{x}_i, \mathbf{x}_i] = \text{softmax}_i(K^T \mathbf{q}_j)$$

Note: we apply softmax to each column so it sums to 1

Figure adapted from 3Blue1Brown ([link](#))

# Self-attention

Weighted sum of values  
(based on self-attention weights)

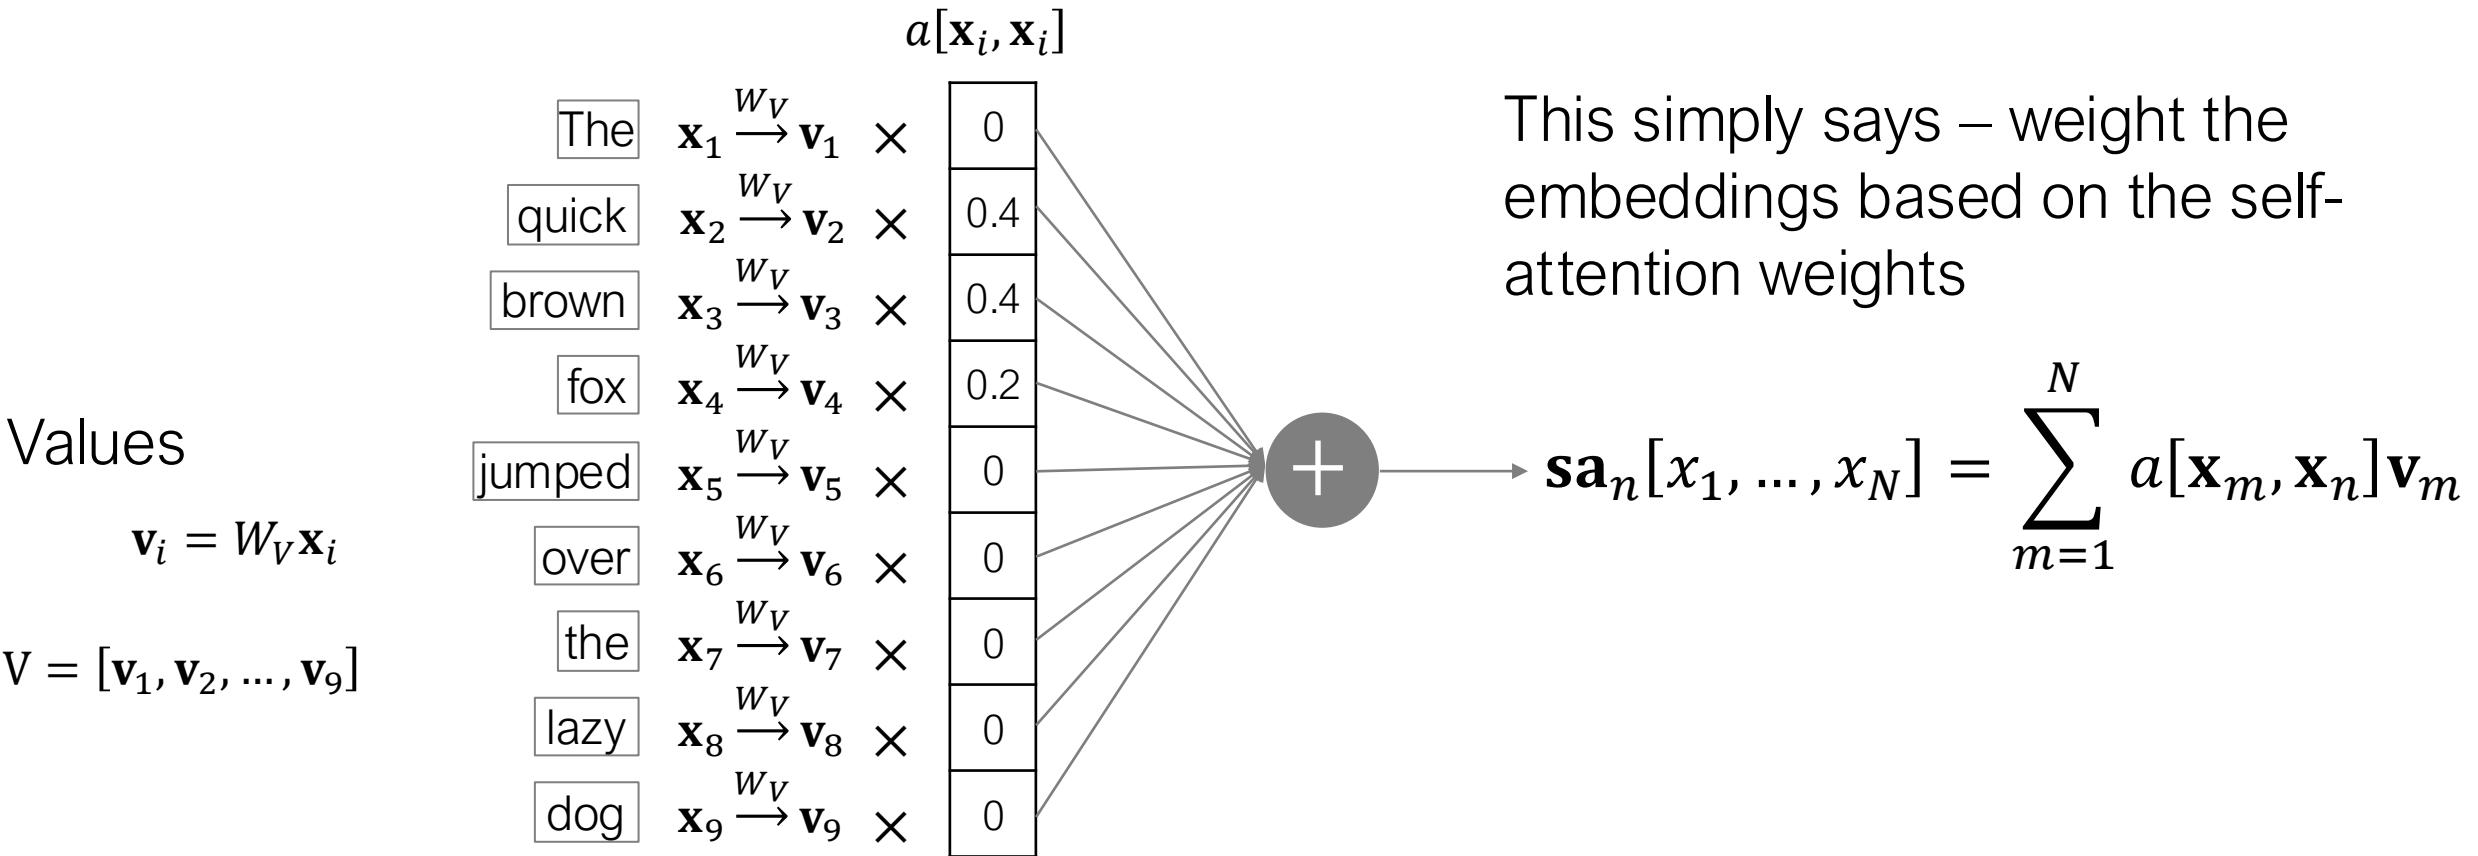


Figure adapted from 3Blue1Brown ([link](#))

# Self-attention in matrix form

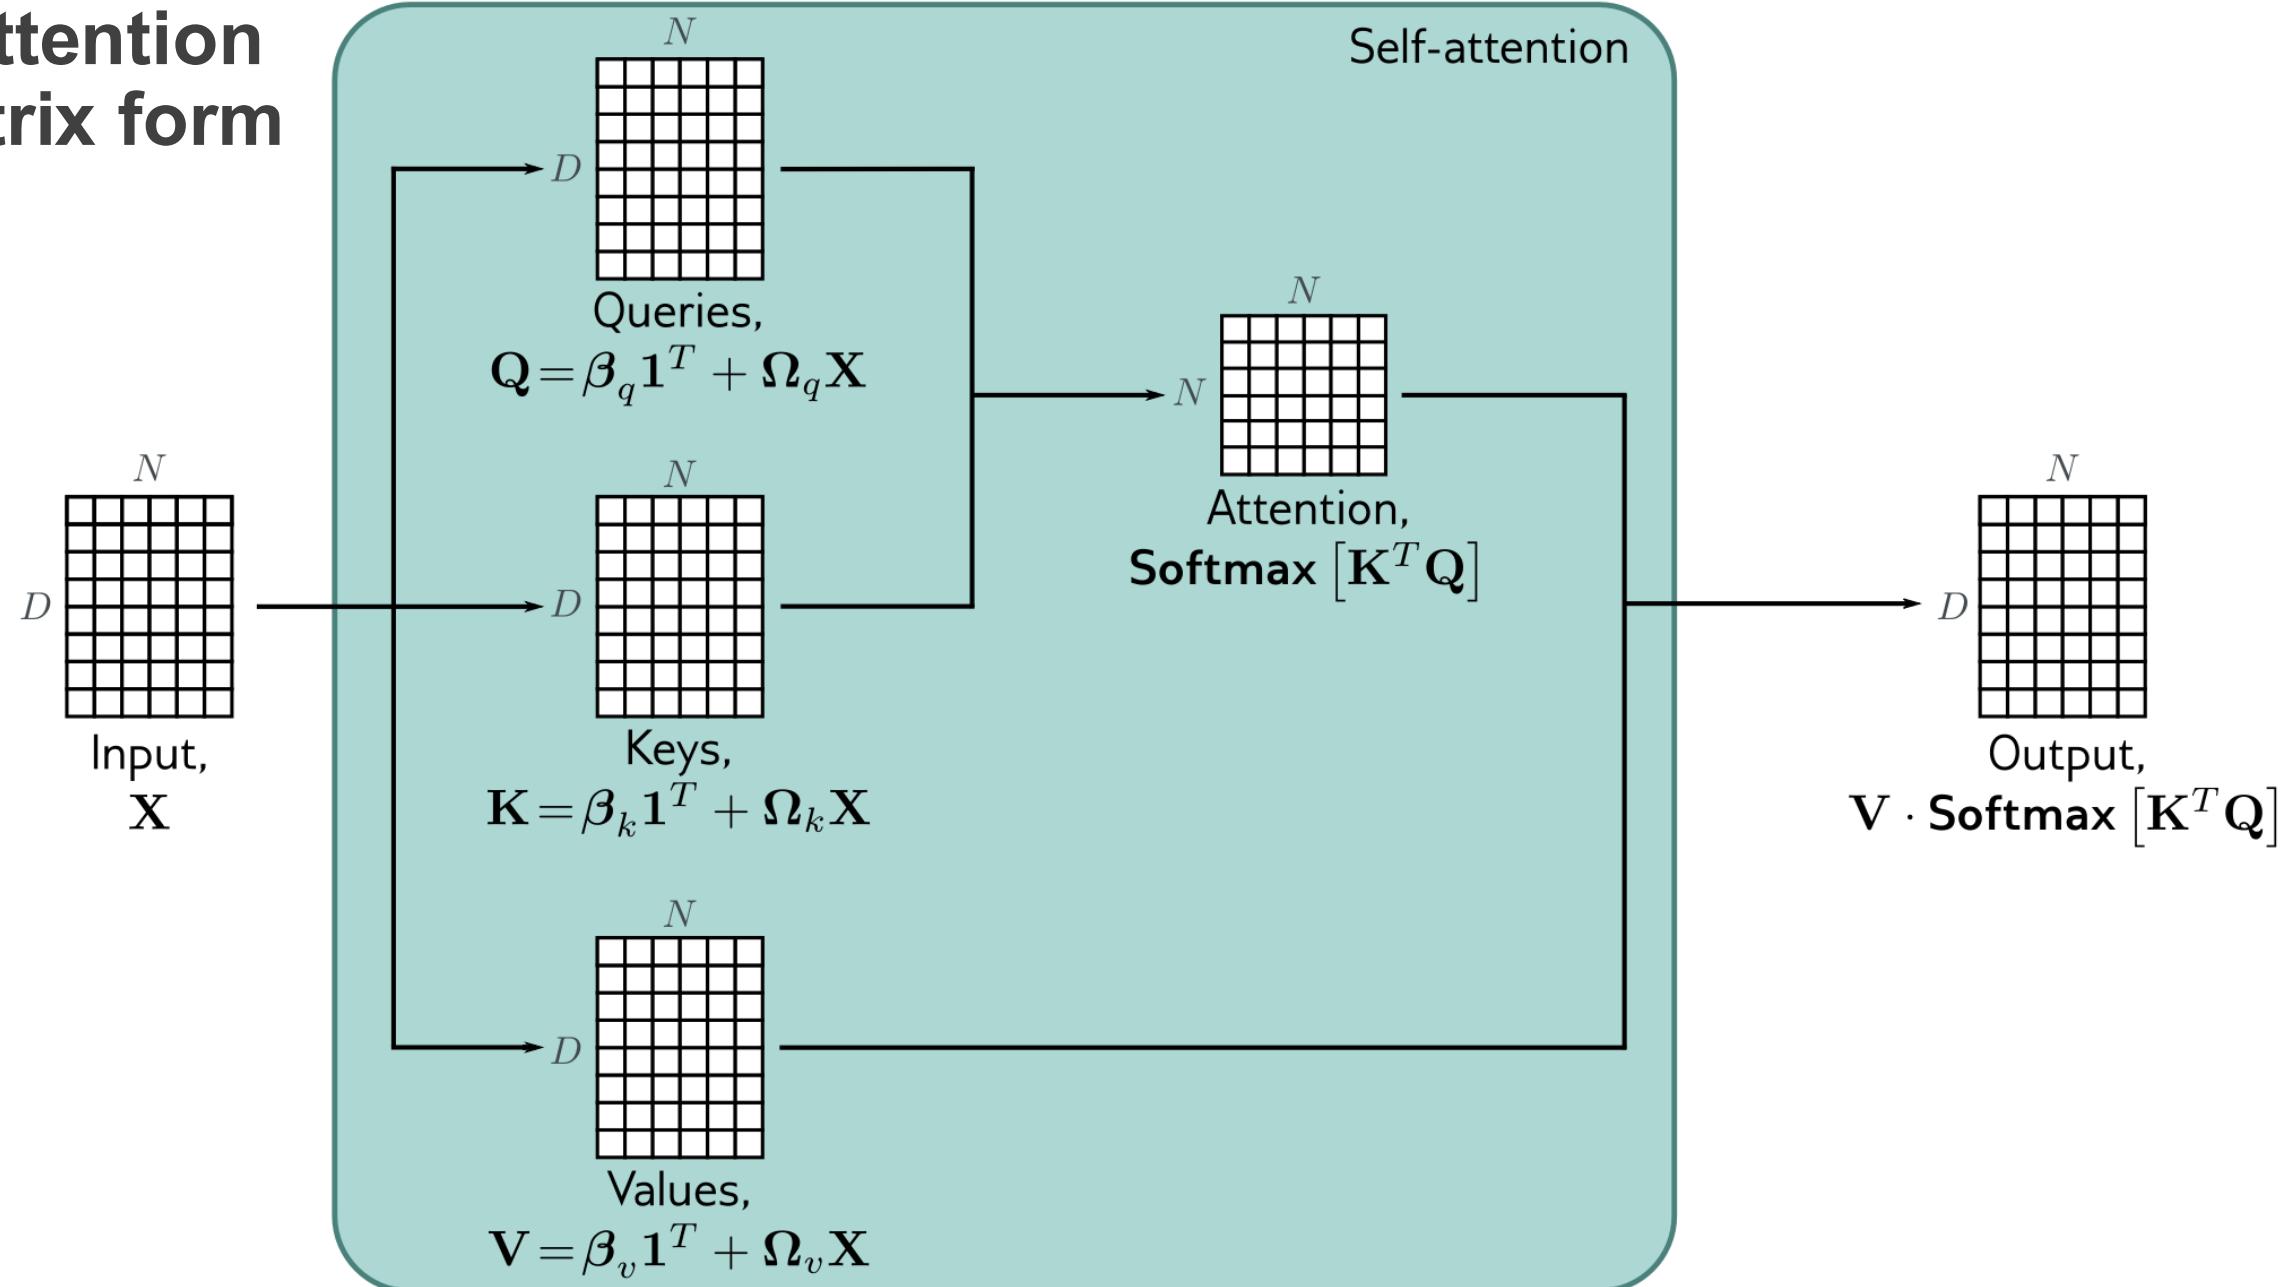


Image from Prince, Understanding Deep Learning, 2023

# Multiheaded attention

Each head may explore different types of relevant information in the rest of the text

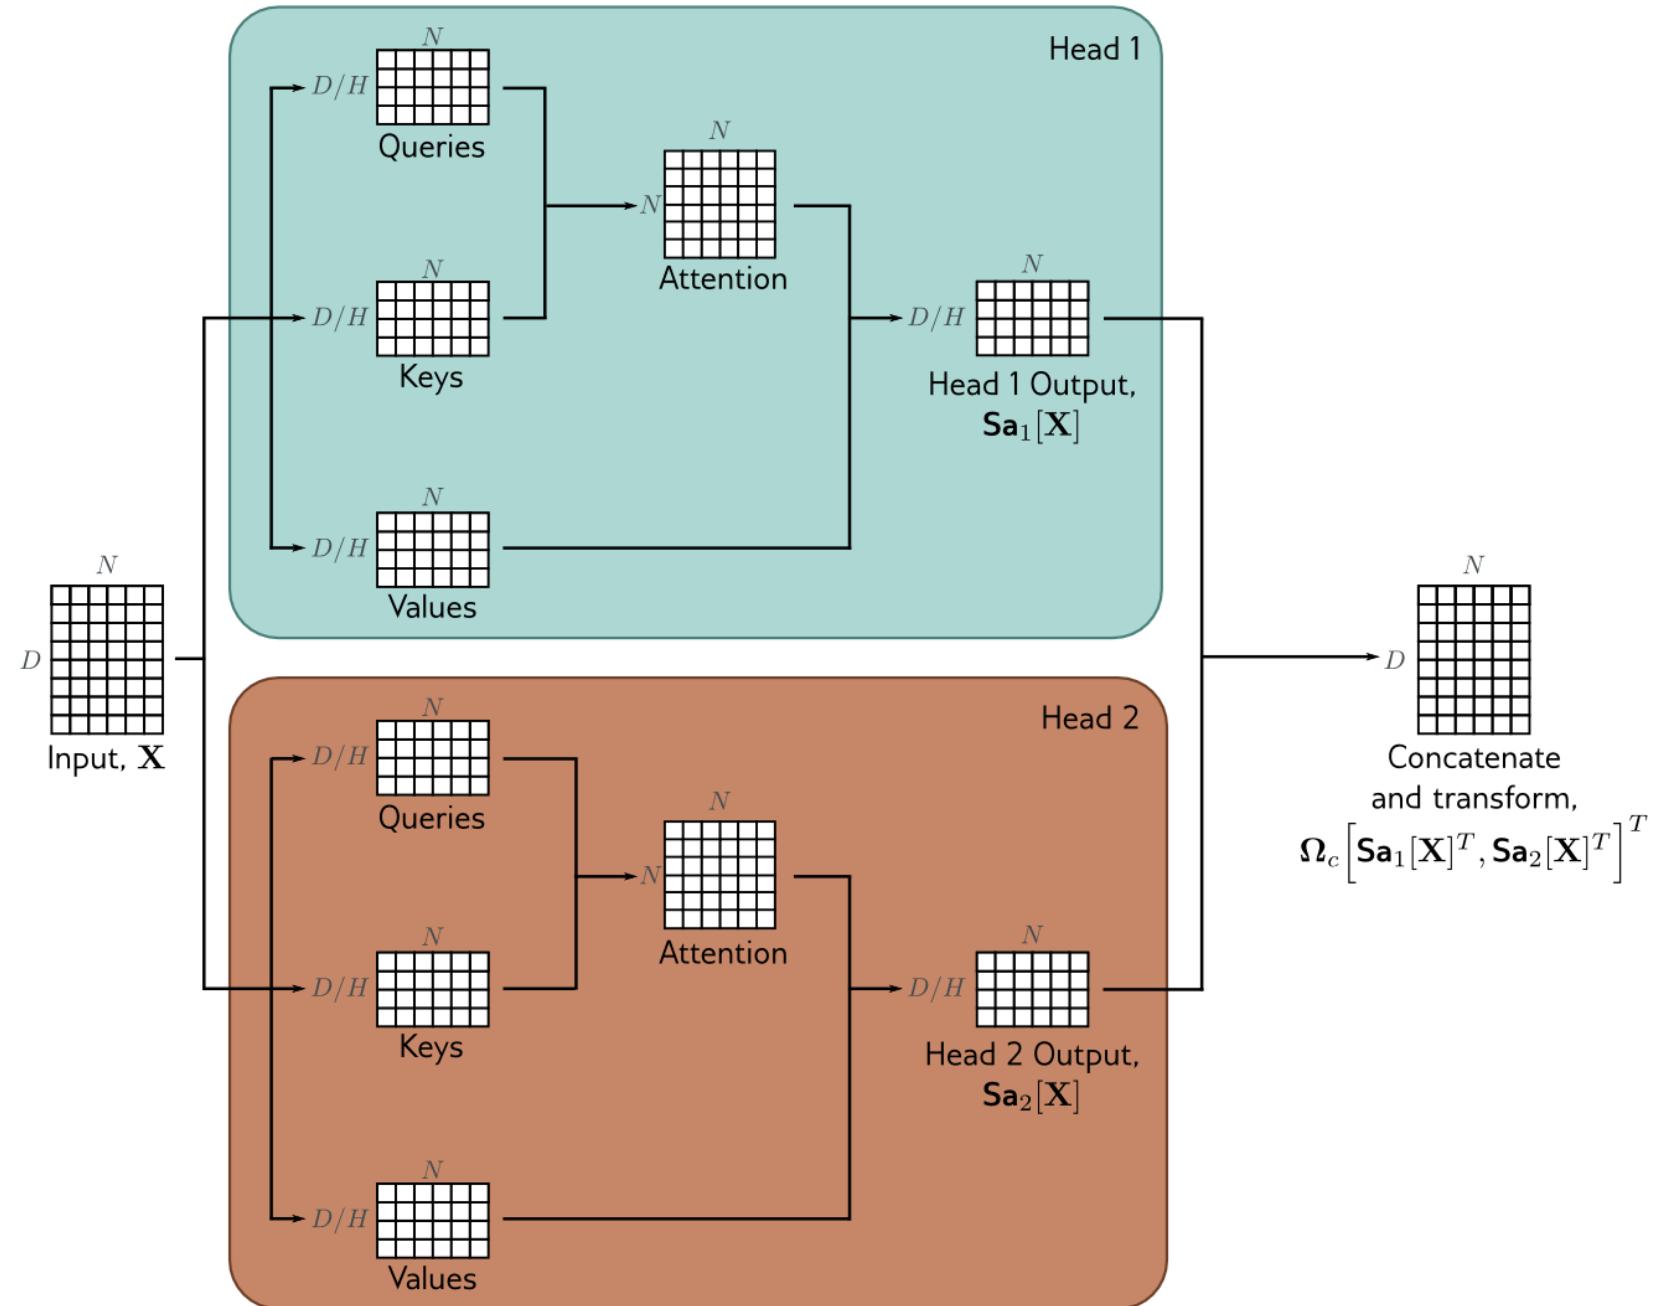
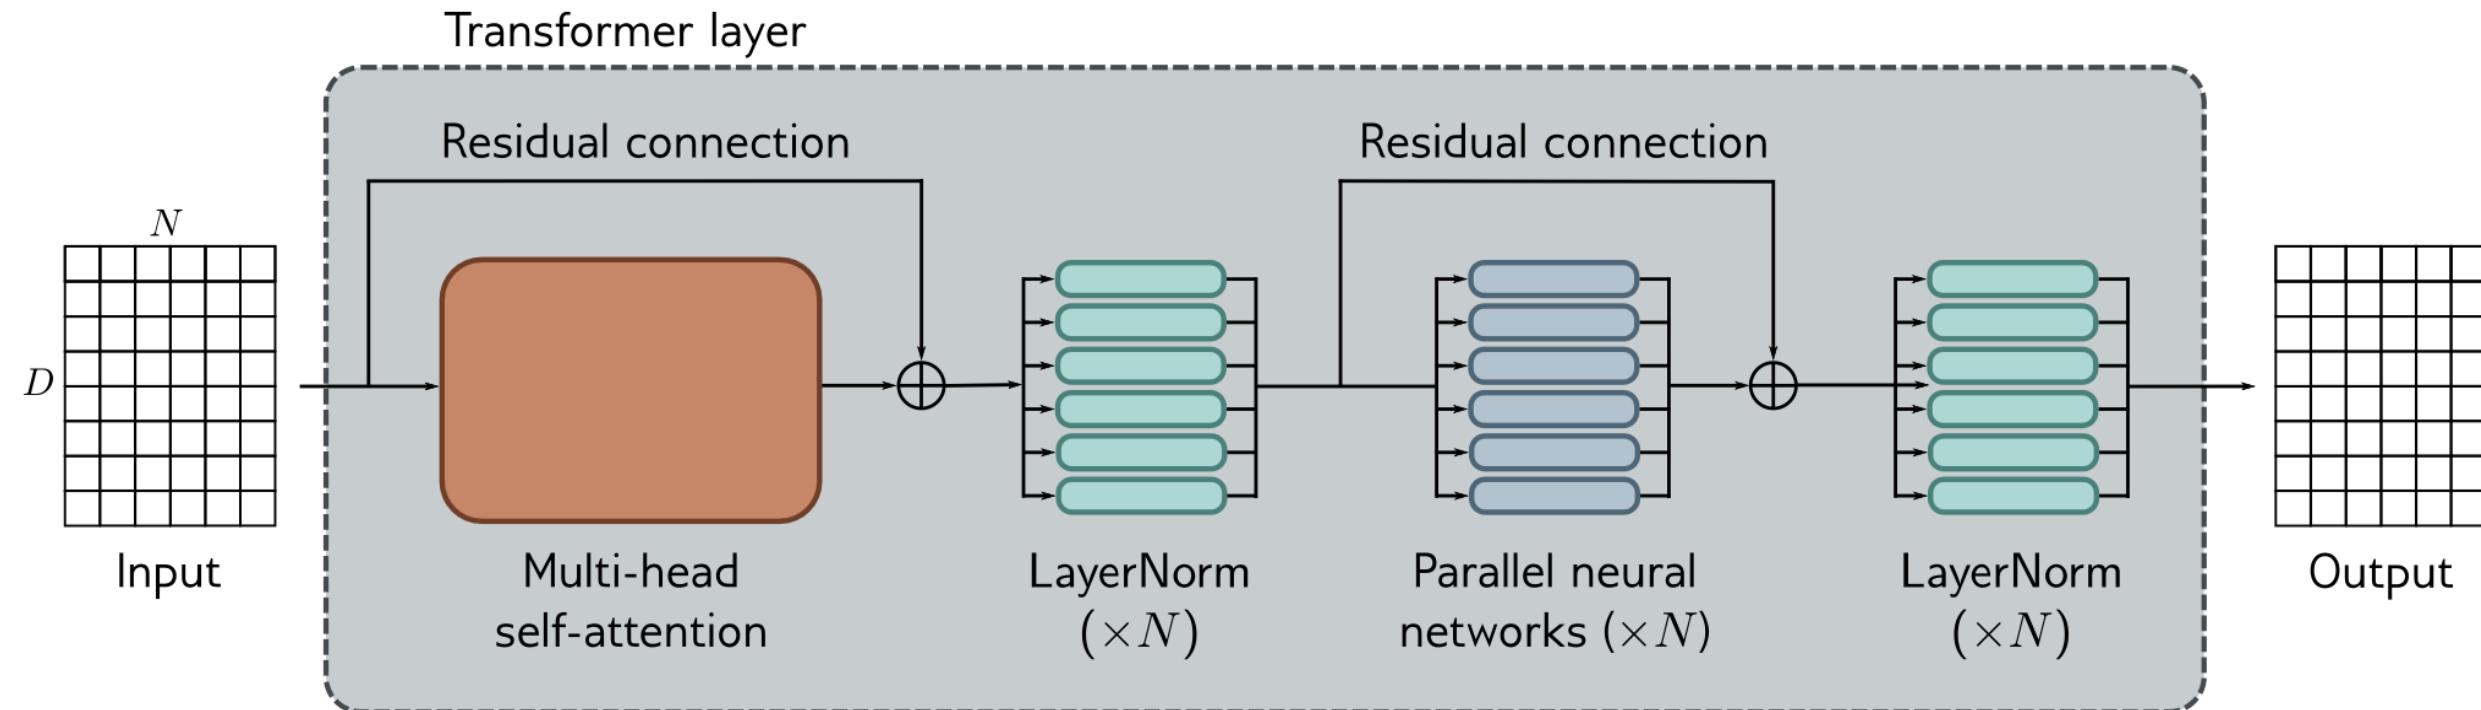


Image from Prince, Understanding Deep Learning, 2023

# The Transformer



**Residual connections can be thought of as:**

new representation = old representation + correction/improvement  
(token + position embedding) (self-attention output)

# Vision Transformer (ViT)

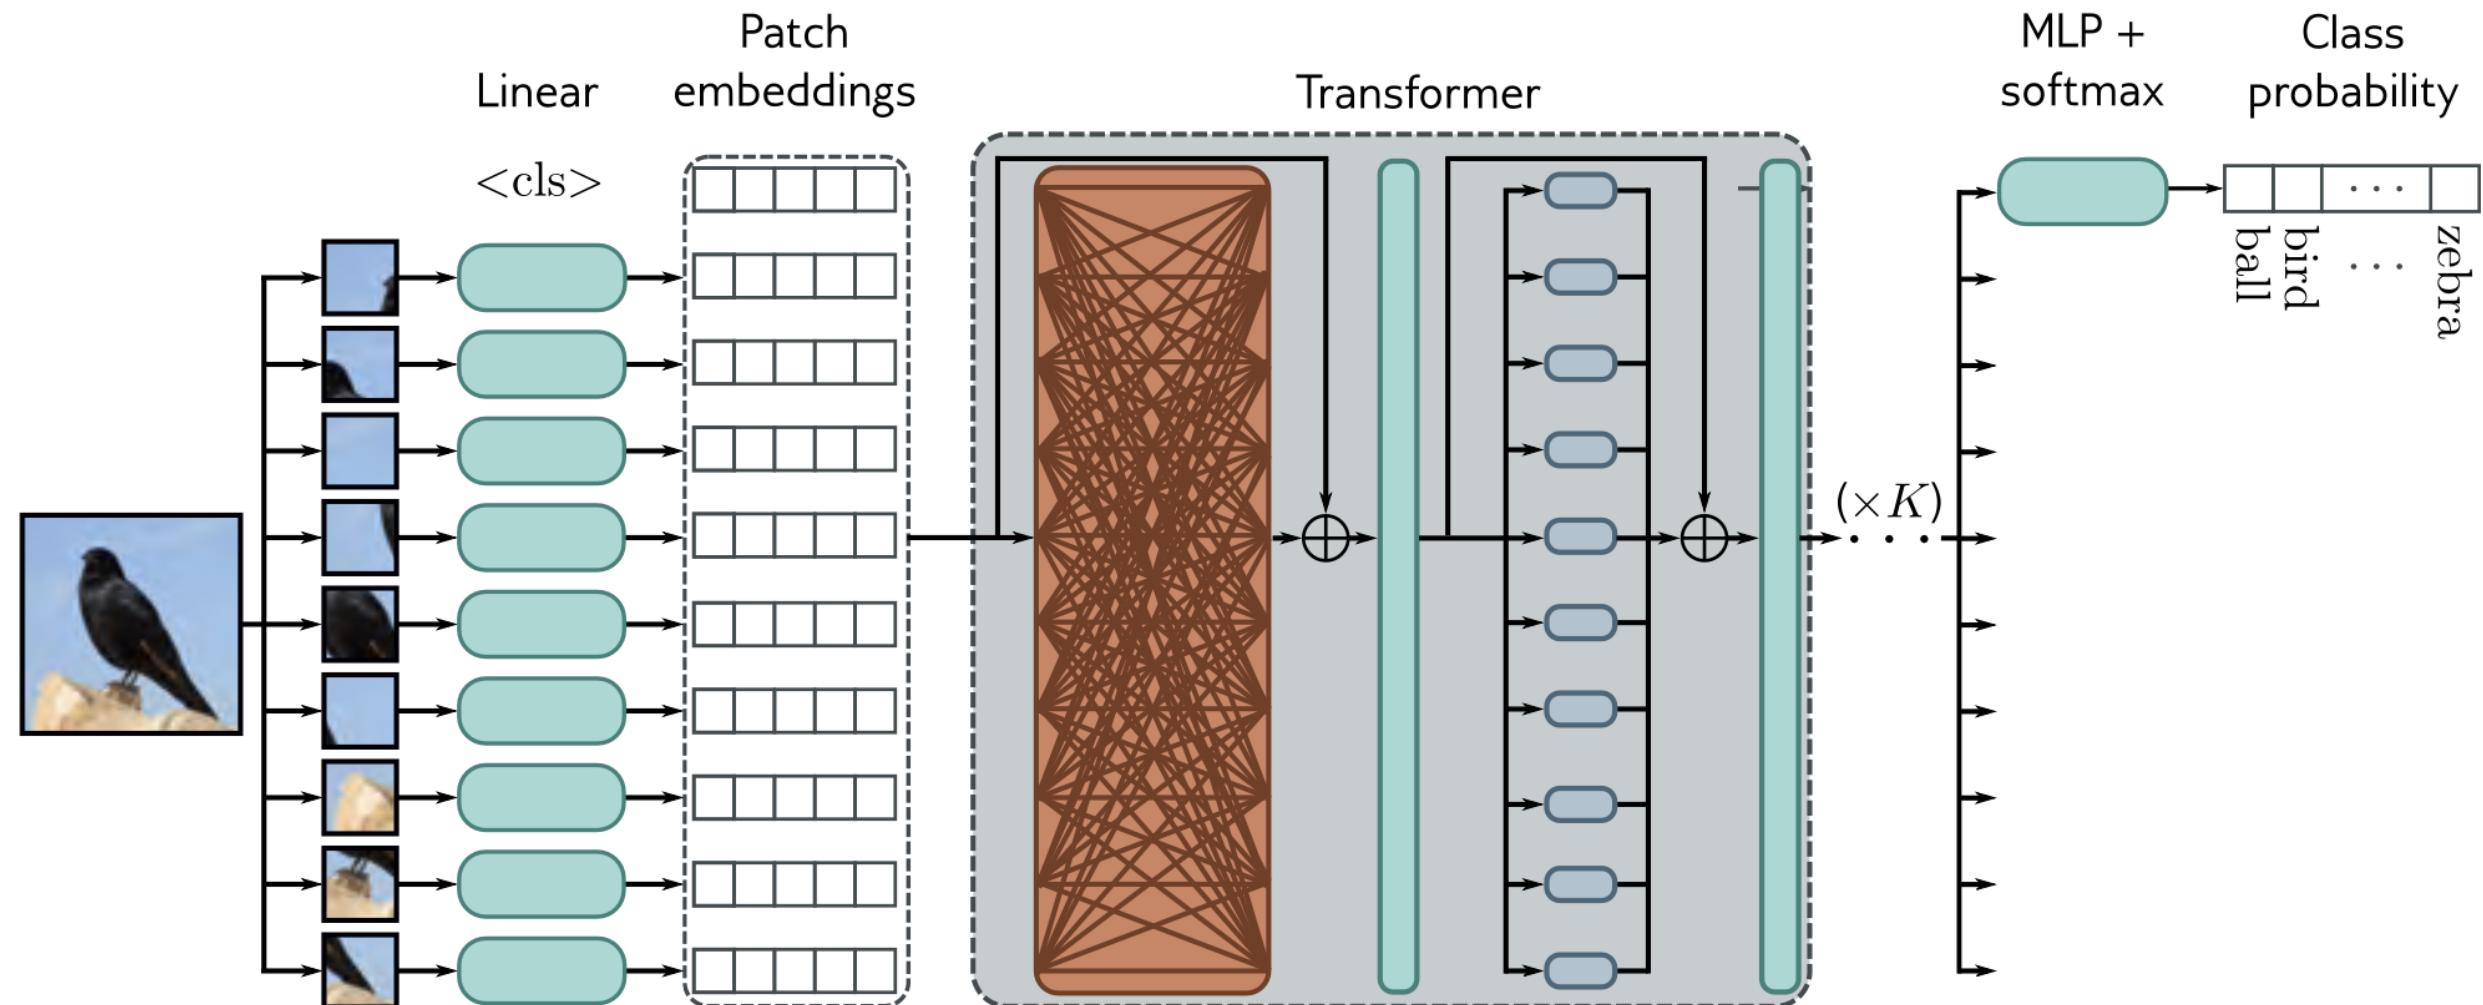


Image from Prince, Understanding Deep Learning, 2023

# Supervised Learning Techniques

Covered so far

- Linear Regression
- ● K-Nearest Neighbors
- ● Logistic Regression
- ● Linear/Quadratic Discriminant Analysis
- ● Naïve Bayes
- ● Decision Trees
- ● Random Forests
- ● Gradient Boosted Decision Trees
- ● Neural networks and deep learning

Ensemble approaches, including bagging, boosting, and stacking, can be used with numerous machine learning techniques, often CART

Appropriate for:  
● Classification  
● Regression

# Supervised learning in practice

