

# Reinforcement Learning II

# Reinforcement Learning Roadmap

- 1 Core concepts in reinforcement learning  
Actions, Rewards, Value, Environments, and Policies

- 2 Markov decision processes  
...and Markov chains and Markov reward processes

- 3 Dynamic Programming  
How do we find optimal policies?  
(Bellman equations)

- 4 Monte Carlo Control  
How do we estimate our value functions?  
How do we use the value functions to choose actions?  
How do we learn optimal policies from experience?

Knowledge of **Environment**

## Perfect knowledge


Known Markov  
Decision Process



## No knowledge

Must learn from  
experience

# Reinforcement Learning Components



**Policy** (agent behavior),  $\pi(s)$

**Reward function** (the goal),  $r_t$

**Value functions** (expected returns),  
 $v_\pi(s)$  State value  
 $q_\pi(s, a)$  Action value

## Policy $\pi(s)$

(which actions to take in each state)

Start			
	↓		
→	→	↓	←
↑		↓	
	→	→	↓
	↑		↓
→	↑		↓
Exit			

## Reward $r_t$

(rewards are received after actions are taken)

Start			
	-1		
-1	-1	-1	-1
-1		-1	
	-1	-1	-1
	-1		-1
-1	-1		-1
Exit			

## State Value $v_\pi(s)$

(expected cumulative rewards starting from current state **if** we follow the policy)

Start			
	-8		
-8	-7	-6	-7
-9		-5	
	-5	-4	-3
	-6		-2
-8	-7		-1
Exit			


## Action Value $q_\pi(s, a)$

(expected cumulative rewards starting from current state **if** we take action  $a$  then follow the policy)

↑	-9
→	-7
←	-9

↑	-4
↓	-2

# Value functions



## State Value function, $v_\pi(s)$

- How “good” is it to be in a state,  $s_t$  then follow policy  $\pi$  to choose actions
- Total expected rewards

$$v_\pi(s) = E_\pi[G_t | s_t = s]$$


## Action Value function, $q_\pi(s, a)$

- How “good” is it to be in a state,  $s$ , take action  $a$ , then follow policy  $\pi$  to choose actions
- Total expected rewards

$$q_\pi(s, a) = E_\pi[G_t | s_t = s, a_t = a]$$

Where  $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$

# Policy



## Policy, $\pi(s)$

- Selects an action to choose based on the state
- Determines an agent's "behavior"

Deterministic policy:

$$a = \pi(s)$$

Stochastic policy:

$$\pi(a|s) = P(A_t = a | S_t = s)$$

Helps us "explore" the state space

RL tries to learn the "best" policy

# Returns / cumulative reward

**Episodic** tasks (finite number,  $T$ , of steps, then reset)

$$G_t = r_{t+1} + r_{t+2} + \dots + r_T$$

**Continuing** tasks with discounting ( $T \rightarrow \infty$ )

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

where  $0 \leq \gamma \leq 1$  is the discount rate

$\gamma$  makes the agent care more/less about immediate rewards

# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)


**MDPs form the framework for most reinforcement learning environments**

Adapted from David Silver, 2015



# History

The record of all that has happened in this system



Step 0:  $s_0, a_0$

Step 1:  $r_1, s_1, a_1$

Step 2:  $r_2, s_2, a_2$

$\vdots$

Step T:  $r_t, s_t, a_t$

History at time  $t$  :  $H_t = \{s_t, a_t, r_{t-1}, s_{t-1}, a_{t-1}, \dots, r_1, s_1, a_1, s_0, a_0\}$

# Markov property

Instead of needing the full history:

$$H_t = \{s_t, a_t, r_{t-1}, s_{t-1}, a_{t-1}, \dots, r_1, s_1, a_1, s_0, a_0\}$$

We can summarize everything in the current state

$$H_t = \{s_t, a_t\}$$

**The future is independent of the past given the present**

Another way of saying this is:

$$P(s_{t+1}|s_t) = P(s_{t+1}|s_t, s_{t-1}, \dots, s_1, s_0)$$

Yet another way of saying this is that as long as we know our current state, how we got there is irrelevant with respect to planning about the future

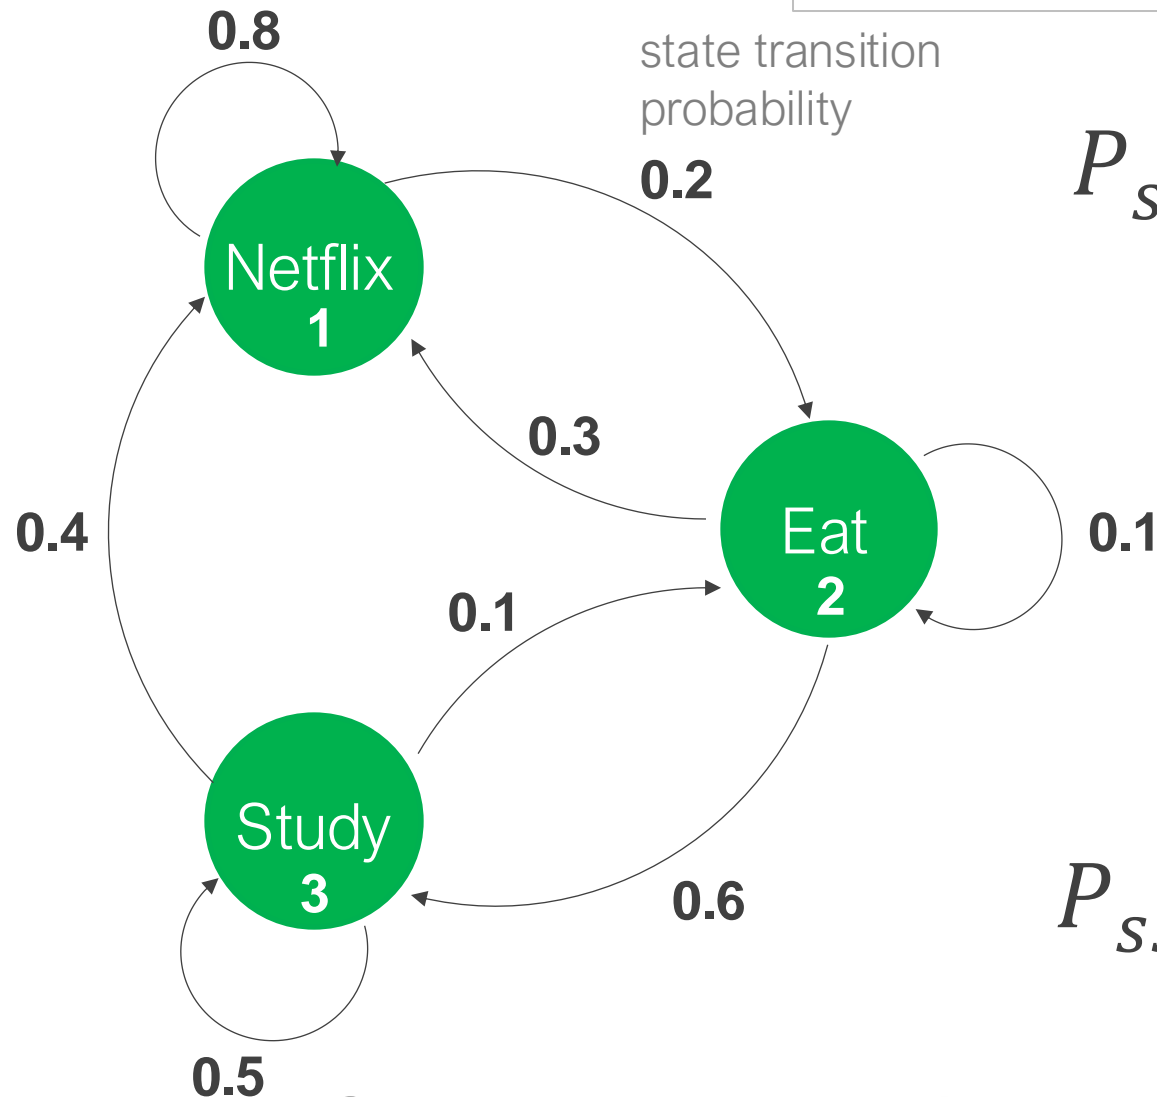
# Markov Chains

Example: student life

Two components:  $\{S, P\}$

State space,  $S$

Transition matrix,  $P$



State-space representation

$$P_{SS'} =$$

$$P_{SS'} =$$

State transition probabilities

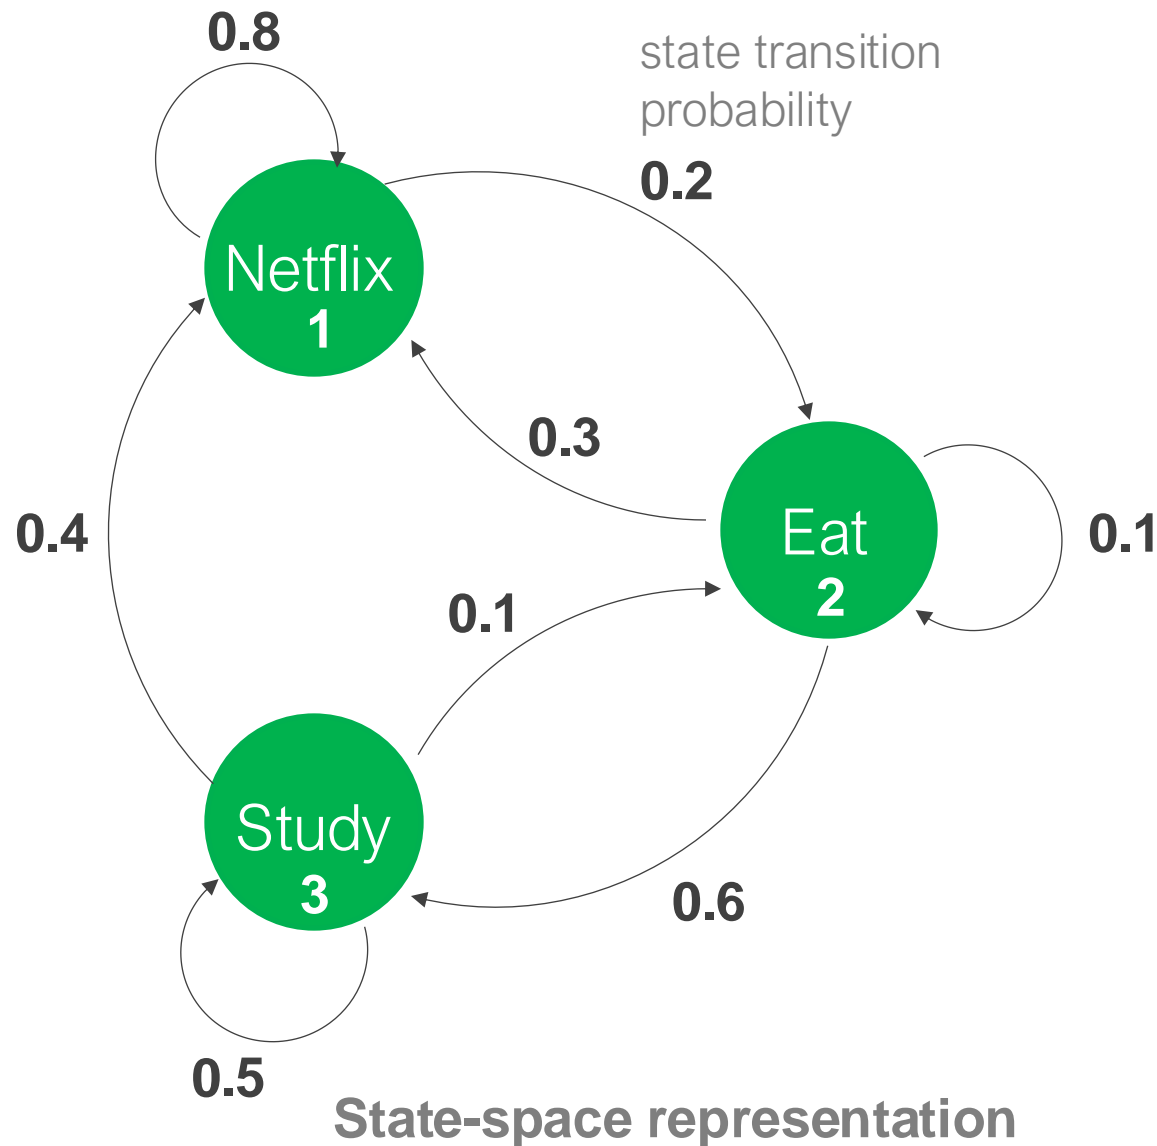
		To state		
		1	2	3
From state	1	$p_{11}$	$p_{12}$	$p_{13}$
	2	$p_{21}$	$p_{22}$	$p_{23}$
	3	$p_{31}$	$p_{32}$	$p_{33}$

Transitions out of each state sum to 1

		To state		
		Netflix	Eat	Study
From state	Netflix	0.8	0.2	0
	Eat	0.3	0.1	0.6
	Study	0.4	0.1	0.5

# Markov Chains

Example: student life



If we start in state 1, what's the probability we'll be in each state after one step?

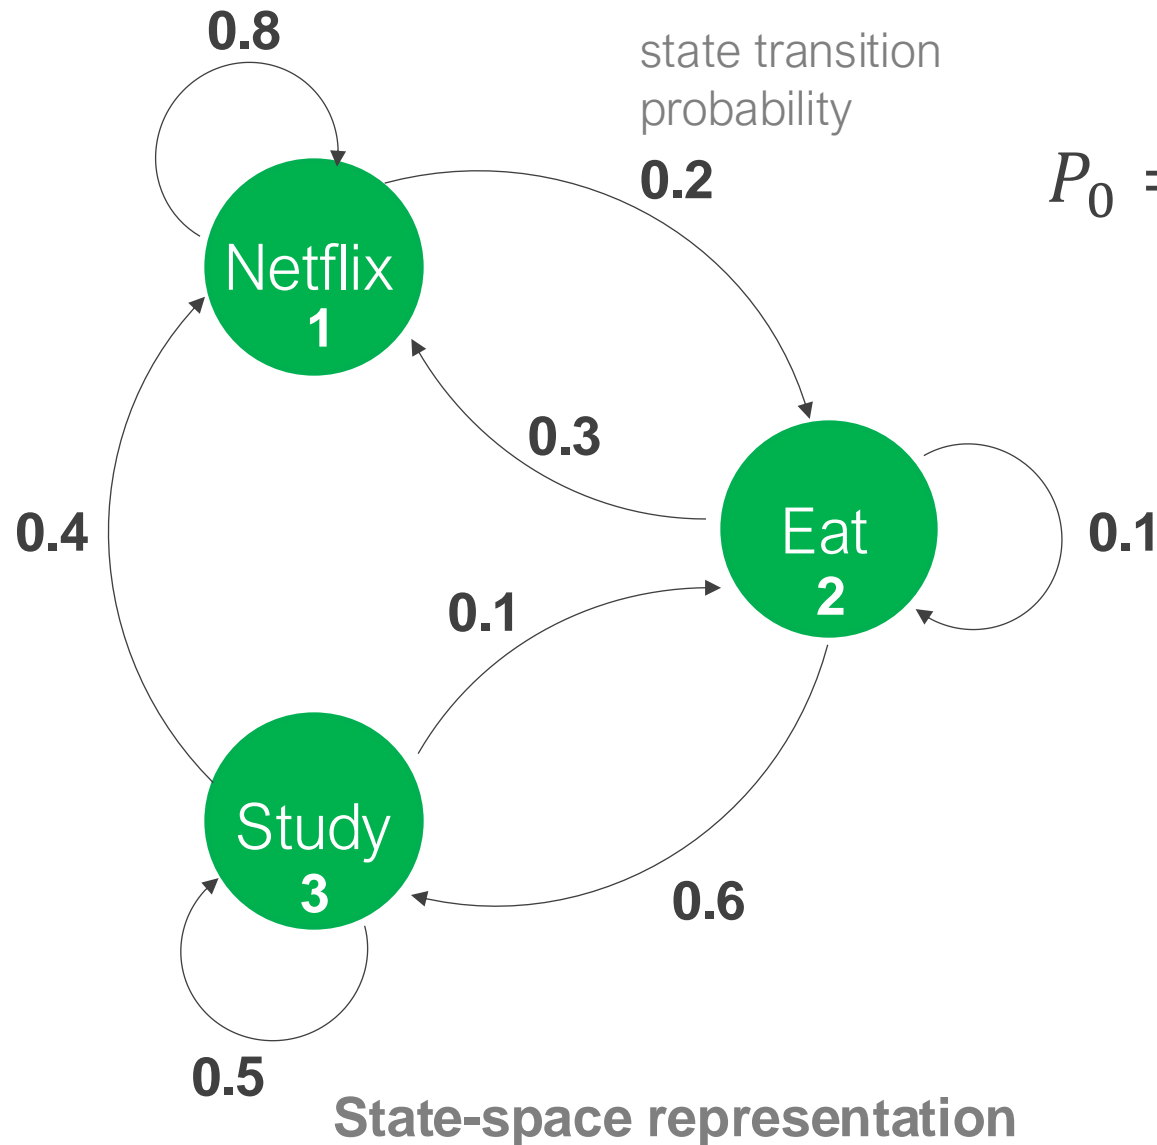
$$P_1 = \begin{bmatrix} 0.8 & 0.2 & 0 \end{bmatrix}$$

This is the first row of the state transition probability matrix

$$P_{ss'} = \begin{array}{c} \text{From state} \\ \text{Netflix} \\ \text{Eat} \\ \text{Study} \end{array} \begin{array}{c} \text{To state} \\ \text{Netflix} \\ \text{Eat} \\ \text{Study} \end{array} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

# Markov Chains

Example: student life



If we started in state 1, we can calculate the probabilities of being in each state at step 1 as:

$$P_0 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}^T$$

$$P_1 = P_0 P_{ss'}$$

$$P_1 = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

$$P_1 = \begin{bmatrix} 0.8 & 0.2 & 0 \end{bmatrix}$$

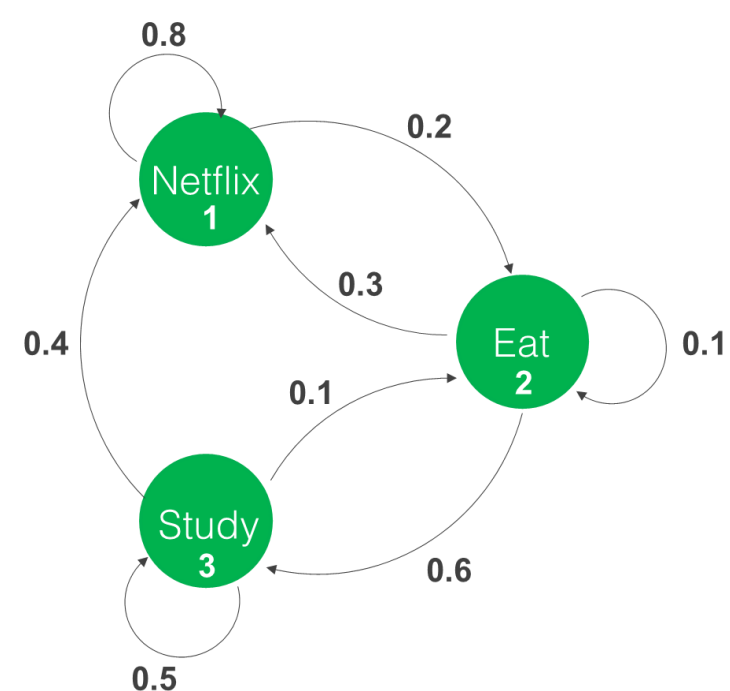
$$P_{ss'} = \begin{array}{c} \text{From state} \\ \text{Netflix} \\ \text{Eat} \\ \text{Study} \end{array} \begin{array}{c} \text{To state} \\ \text{Netflix} \\ \text{Eat} \\ \text{Study} \end{array} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

$$\textbf{1} \quad P_1 = P_0 P_{ss'}$$

$$P_1 = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

$$P_1 = [0.8 \quad 0.2 \quad 0]$$

$$P_0 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}^T$$



$$\textbf{2} \quad P_2 = P_1 P_{ss'} = P_0 P_{ss'} P_{ss'} = P_0 P_{ss'}^2$$

$$P_2 = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix} \begin{bmatrix} 0.8 & 0.2 & 0 \\ 0.3 & 0.1 & 0.6 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

$$P_2 = [0.7 \quad 0.18 \quad 0.12]$$

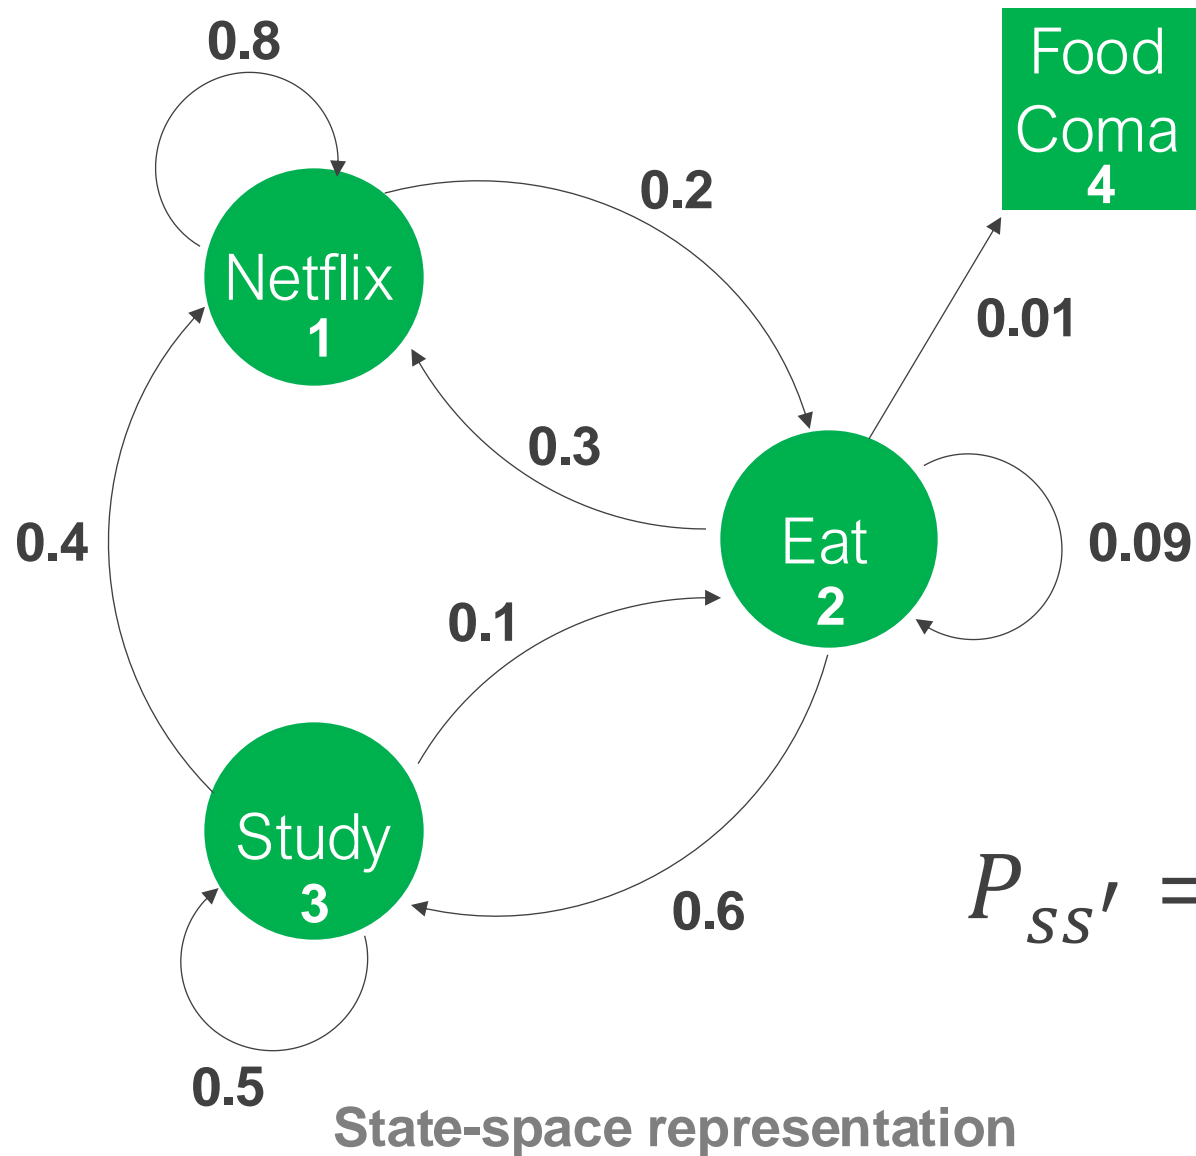
$$\textbf{n} \quad P_n = P_0 P_{ss'}^n$$

As  $n \rightarrow \infty$ , we identify our steady state probabilities

$$P_\infty = [0.64 \quad 0.16 \quad 0.20]$$

# Markov Chains with absorbing state

Example: student life




$$P_{ss'} =$$

		To state			
		Netflix	Eat	Study	Food Coma
From state	Netflix	0.8	0.2	0	0
	Eat	0.3	0.09	0.6	0.01
	Study	0.4	0.1	0.5	0
	Food				
	Coma	0	0	0	1

# Markov Chains

Example: student life



Markov chains can be used to represent sequential discrete-time data

Can estimate long-term state probabilities

Can simulate state sequences based on the model

Markov property applies  
(current state gives you all the information you need about future states)

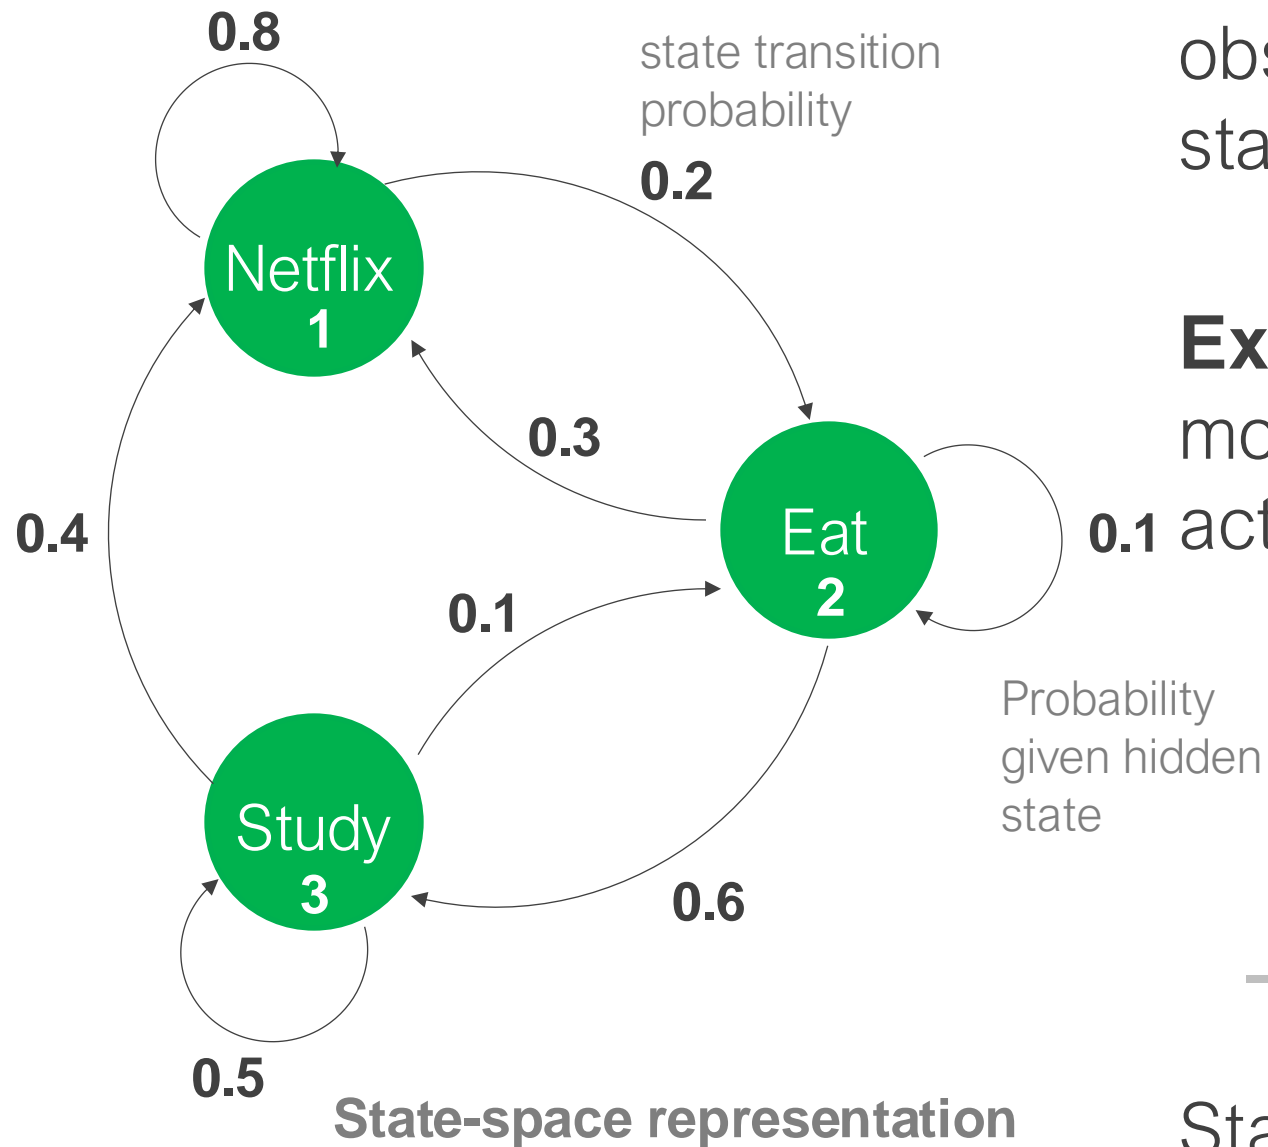
$$P(s_{t+1}|s_t) = P(s_{t+1}|s_t, s_{t-1}, \dots, s_1, s_0)$$

Valid if the system is **autonomous** and the states are **fully observable**




# Hidden Markov Models

Example: student life



What if we **don't directly observe what state** the system is in, but instead observe a quantity that depends on the state?

**Example:** the student wears an EEG monitor, and we see readings of brain activity.



States are **hidden** or **latent variables**

# Markov Models

	States are <b>Fully Observable</b>	States are <b>Partially Observable</b>
<b>Autonomous</b> (no actions; make predictions)	Markov Chain, Markov Reward Process	Hidden Markov Model (HMM)
<b>Controlled</b> (can take actions)	Markov Decision Process (MDP)	Partially Observable Markov Decision Process (POMDP)

## Applications

HMMs: time series ML, e.g. speech + handwriting recognition

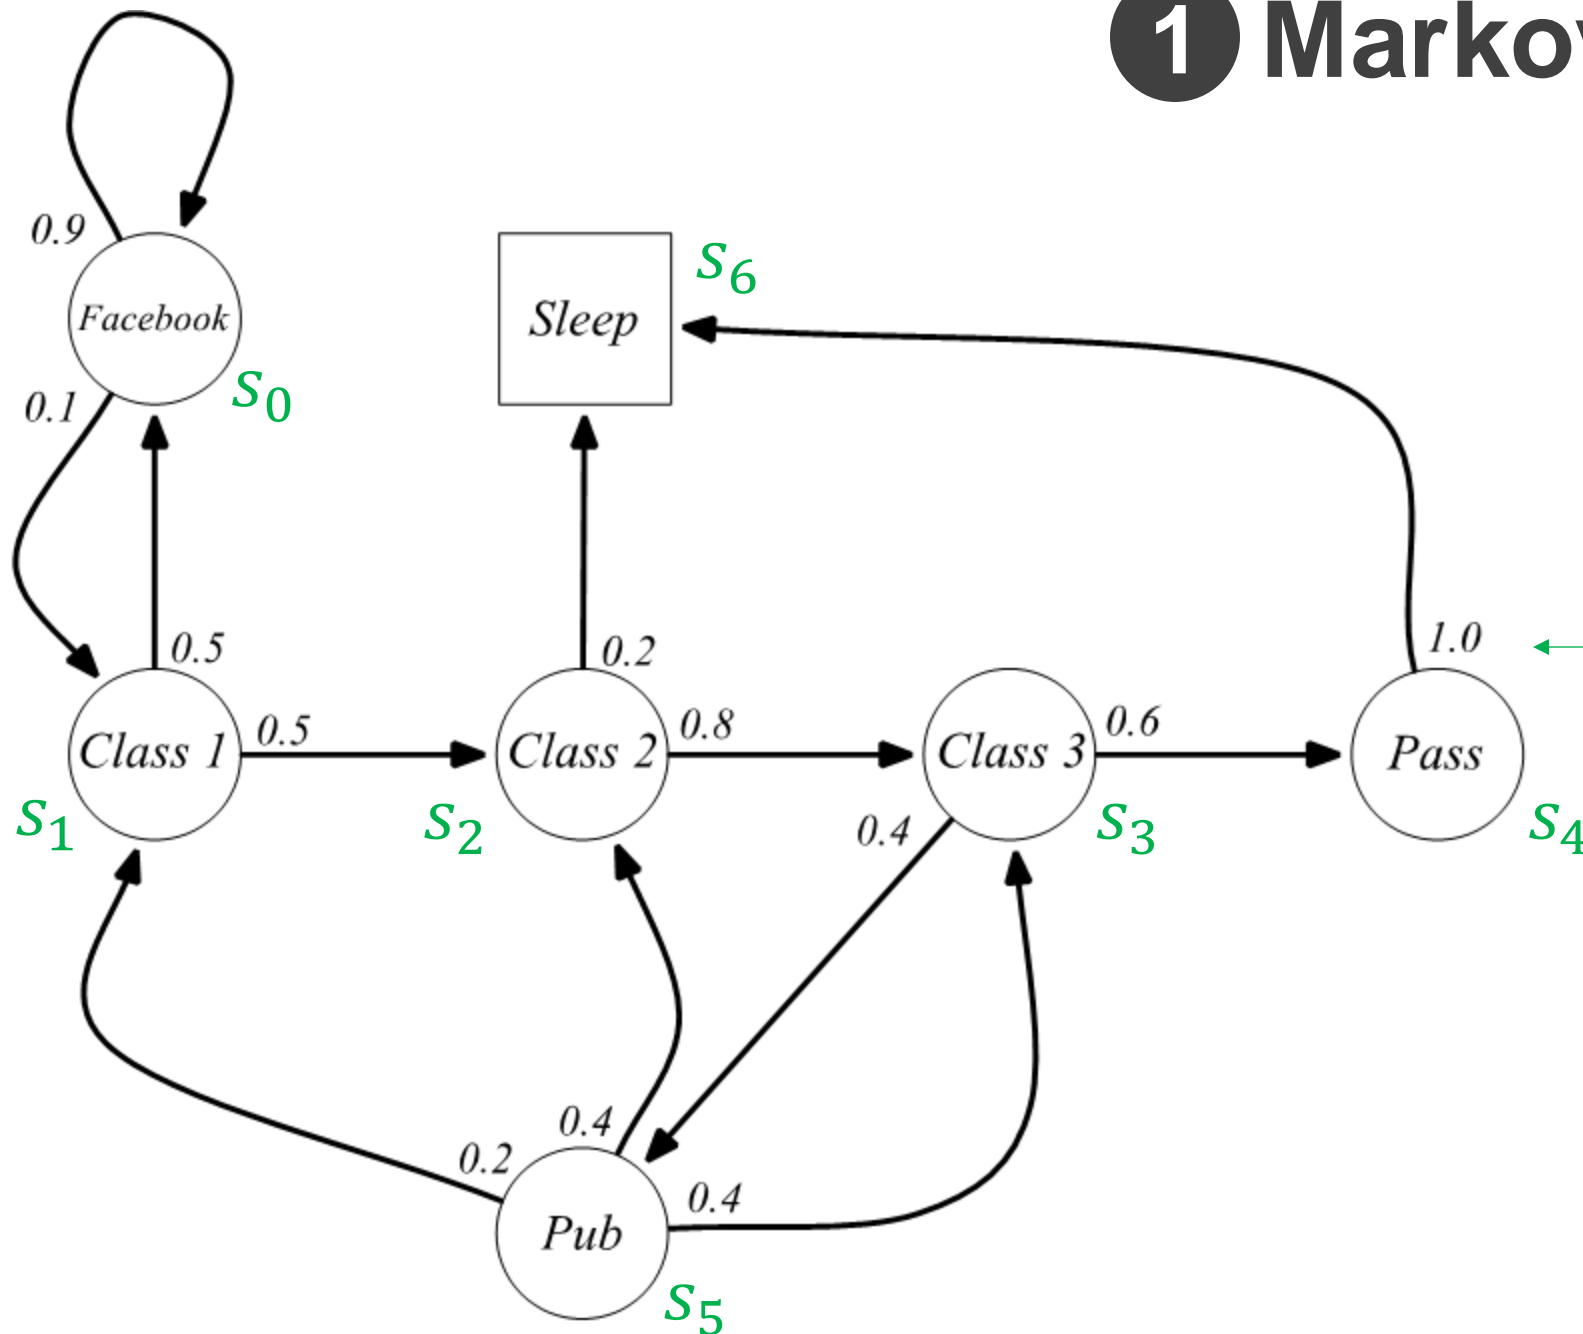
MDPs: framework for reinforcement learning

# 1 Markov Chain Example

**Components:**

State space  $S$ ,

Transition probabilities  $P$



$$P_{46} = P_{ss'}$$

**Sample Episodes:**

C1,C2,Sleep

C1,FB,FB,FB,C1,C2,C3,Pass,Sleep

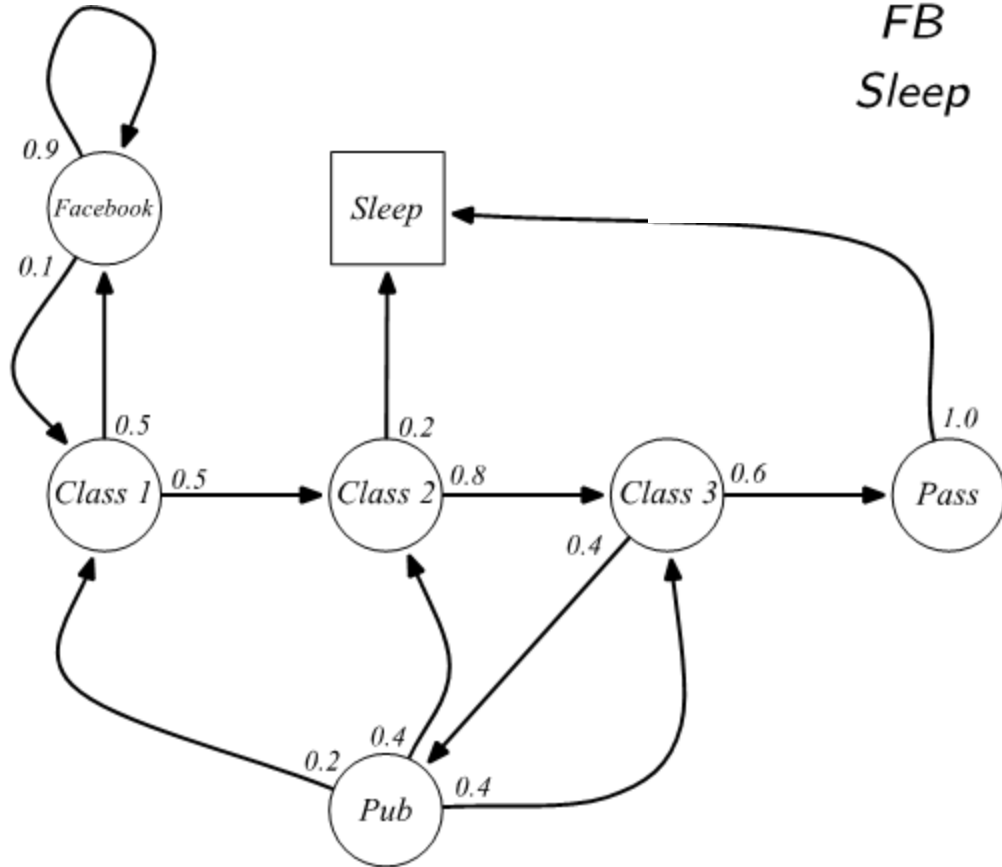
Example from David Silver, UCL, 2015

# Markov Chain

$\mathcal{P} =$   
 $\begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix}$

	C1	C2	C3	Pass	Pub	FB	Sleep
C1		0.5				0.5	
C2			0.8				0.2
C3				0.6	0.4		
Pass	0.2	0.4	0.4				1.0
Pub	0.1					0.9	
FB							
Sleep							1

State transition probability matrix,  $P_{ss'}$



Example from David Silver, UCL, 2015

# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)

**MDPs form the framework for most reinforcement learning environments**

Adapted from David Silver, 2015

## 2 Markov Reward Process

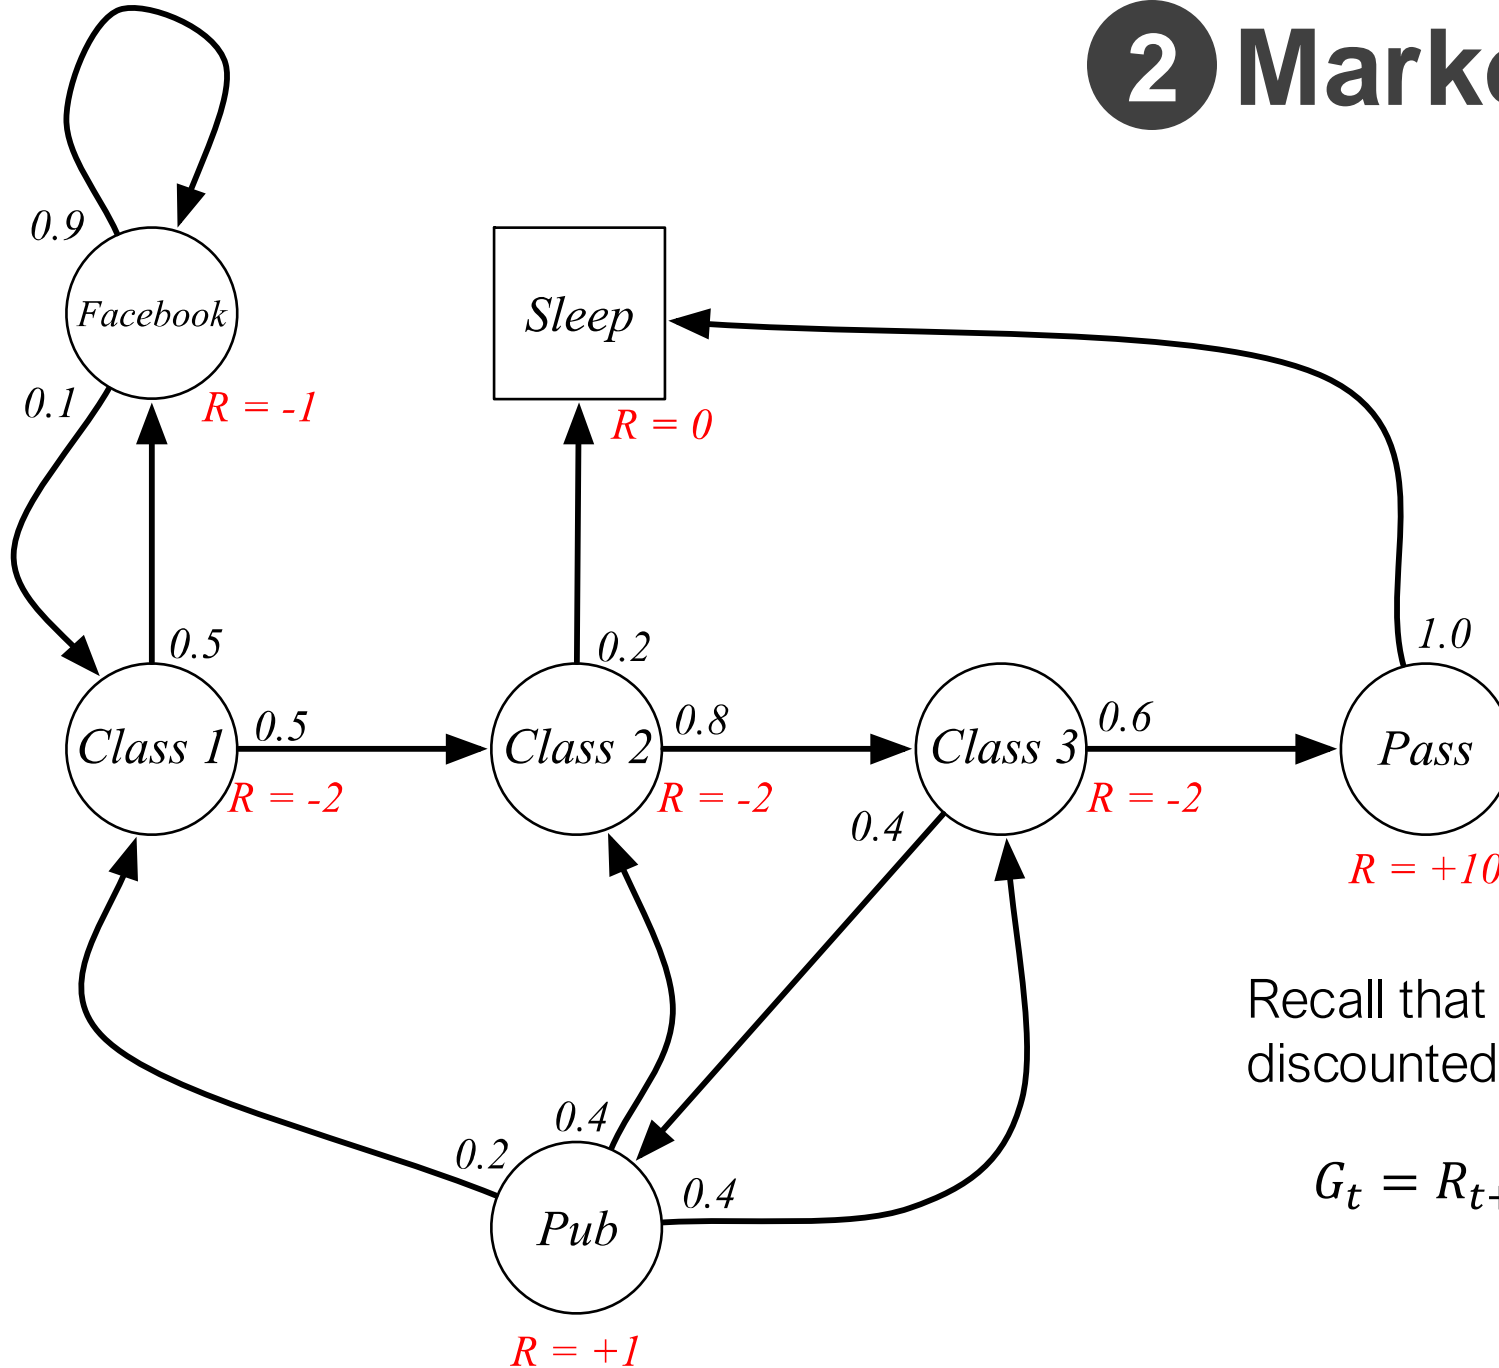
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



Recall that returns, let's call  $G_t$ , are the total discounted rewards from time  $t$ :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

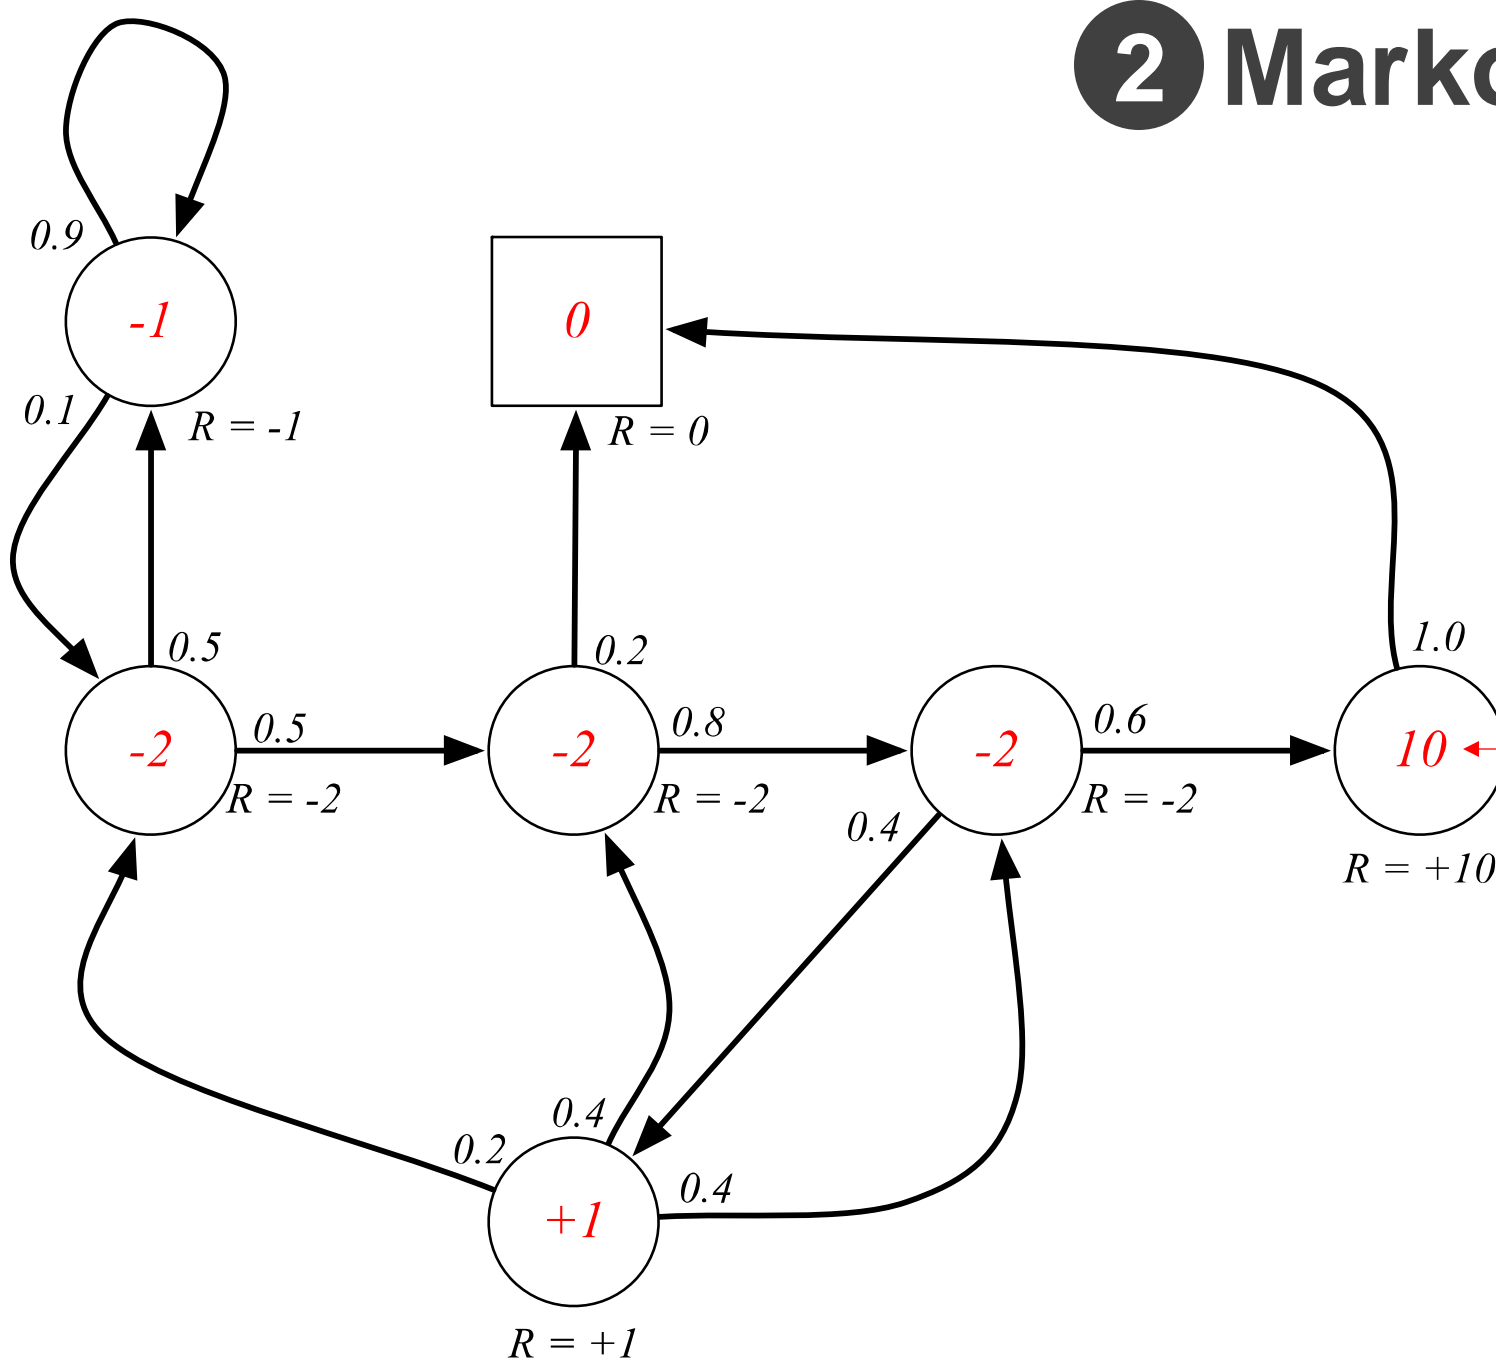
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

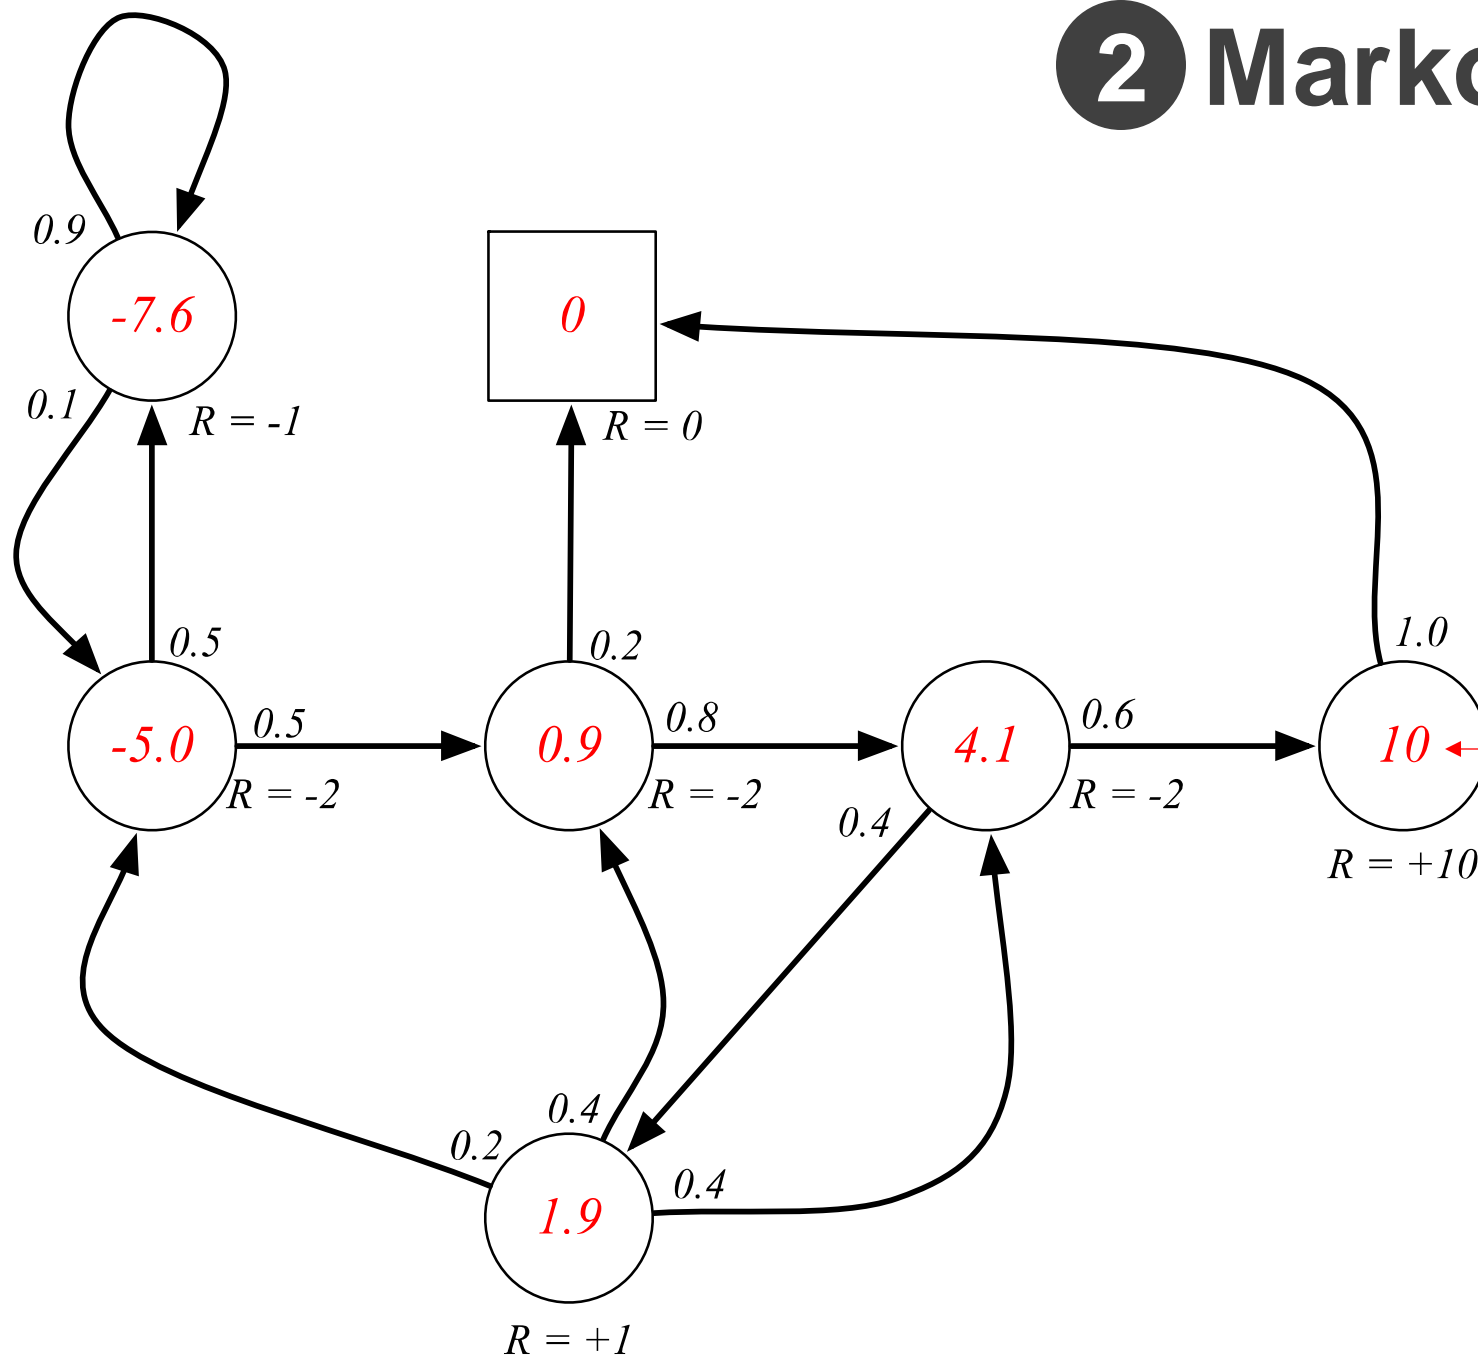
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0.9$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015



# “Backup” property of state value functions

$$v(s) \triangleq E[G_t | S_t = s] \quad \text{where } G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$$

$$= E[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots | S_t = s]$$

$$= E[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} \dots) | S_t = s]$$

$$= E[R_{t+1} + \gamma G_{t+1} | S_t = s]$$


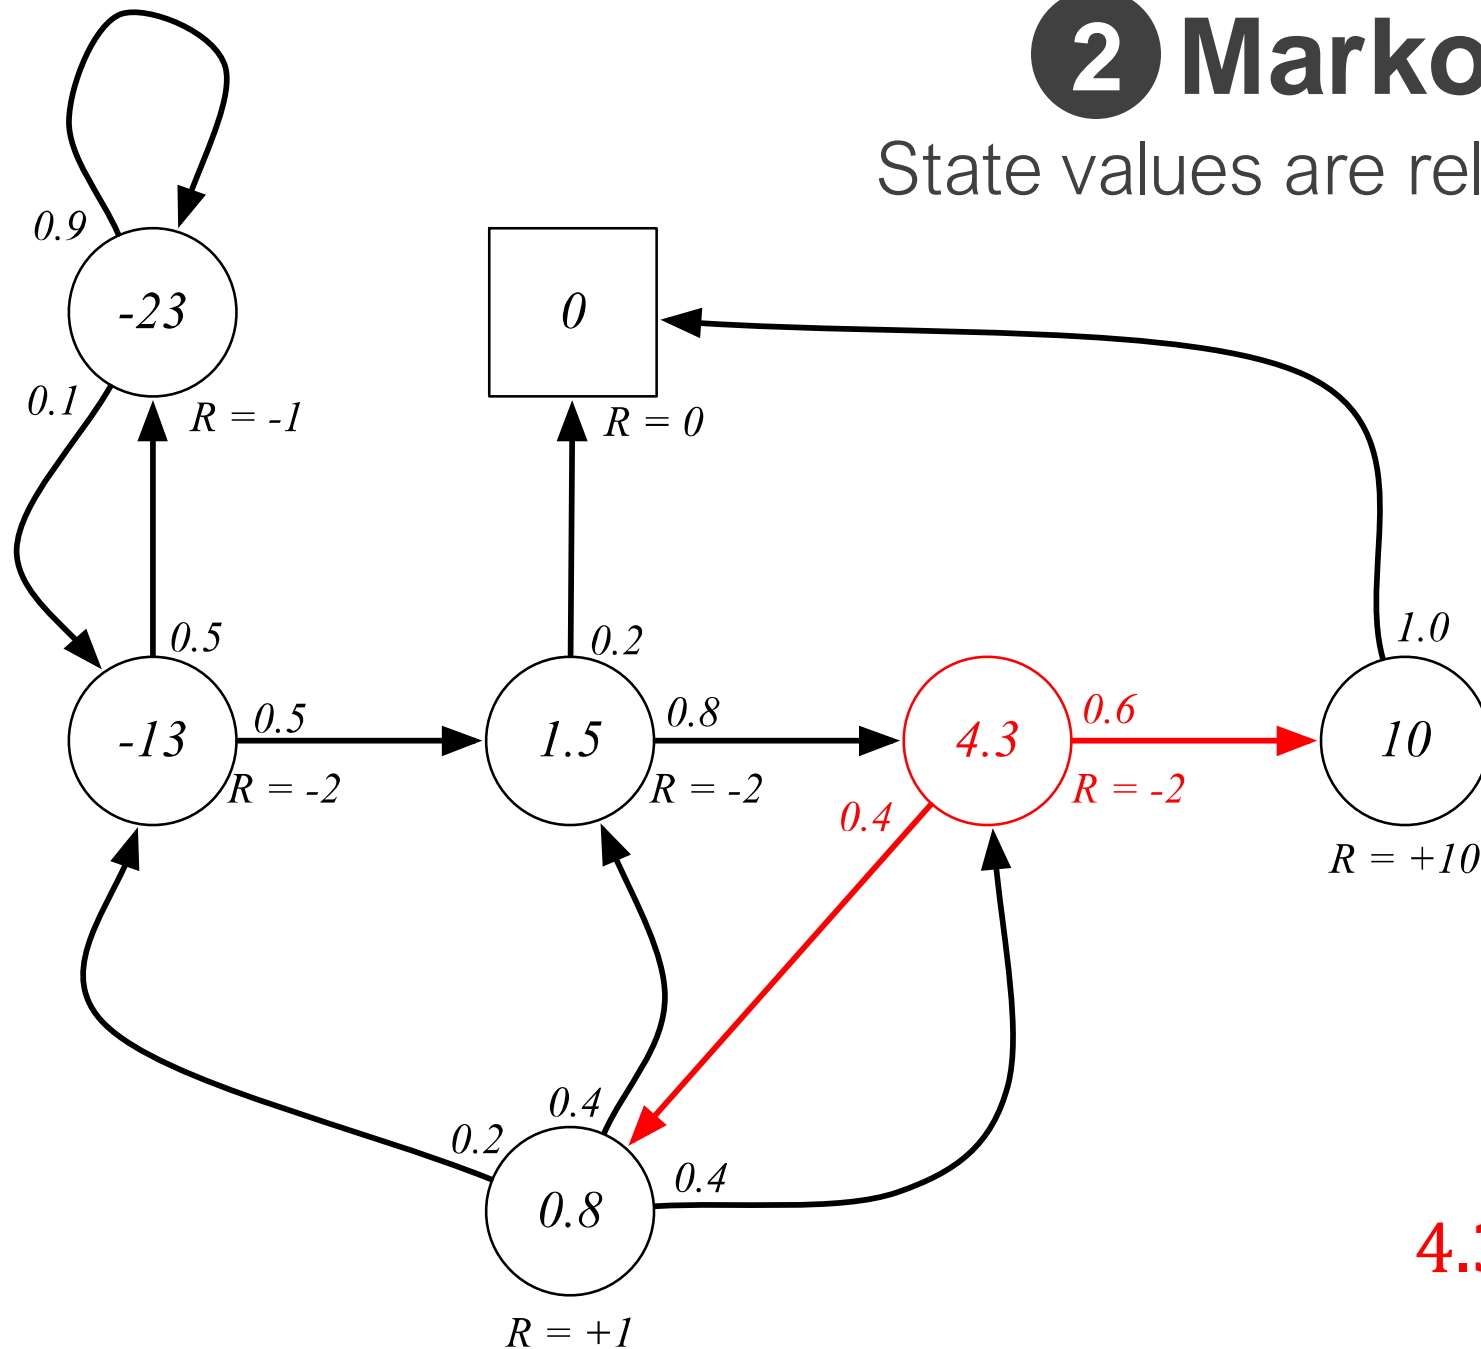
$$= E[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$$

This recursive relationship is a version of the **Bellman Equation**

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

State values are related to neighboring states



possible states we could transition to from  $s$

$$v(s) = E[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$$

$$v(s) = r_s + \gamma \sum_{s'} P_{ss'} v(s')$$

$$r_s = E[R_{t+1} | S_t = s]$$

Assume  $\gamma = 1$

$$4.3 = -2 + 0.6 \times 10 + 0.4 \times 0.8$$

Example from David Silver, UCL, 2015

# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)

**MDPs form the basis for most reinforcement learning environments**

Adapted from David Silver, 2015

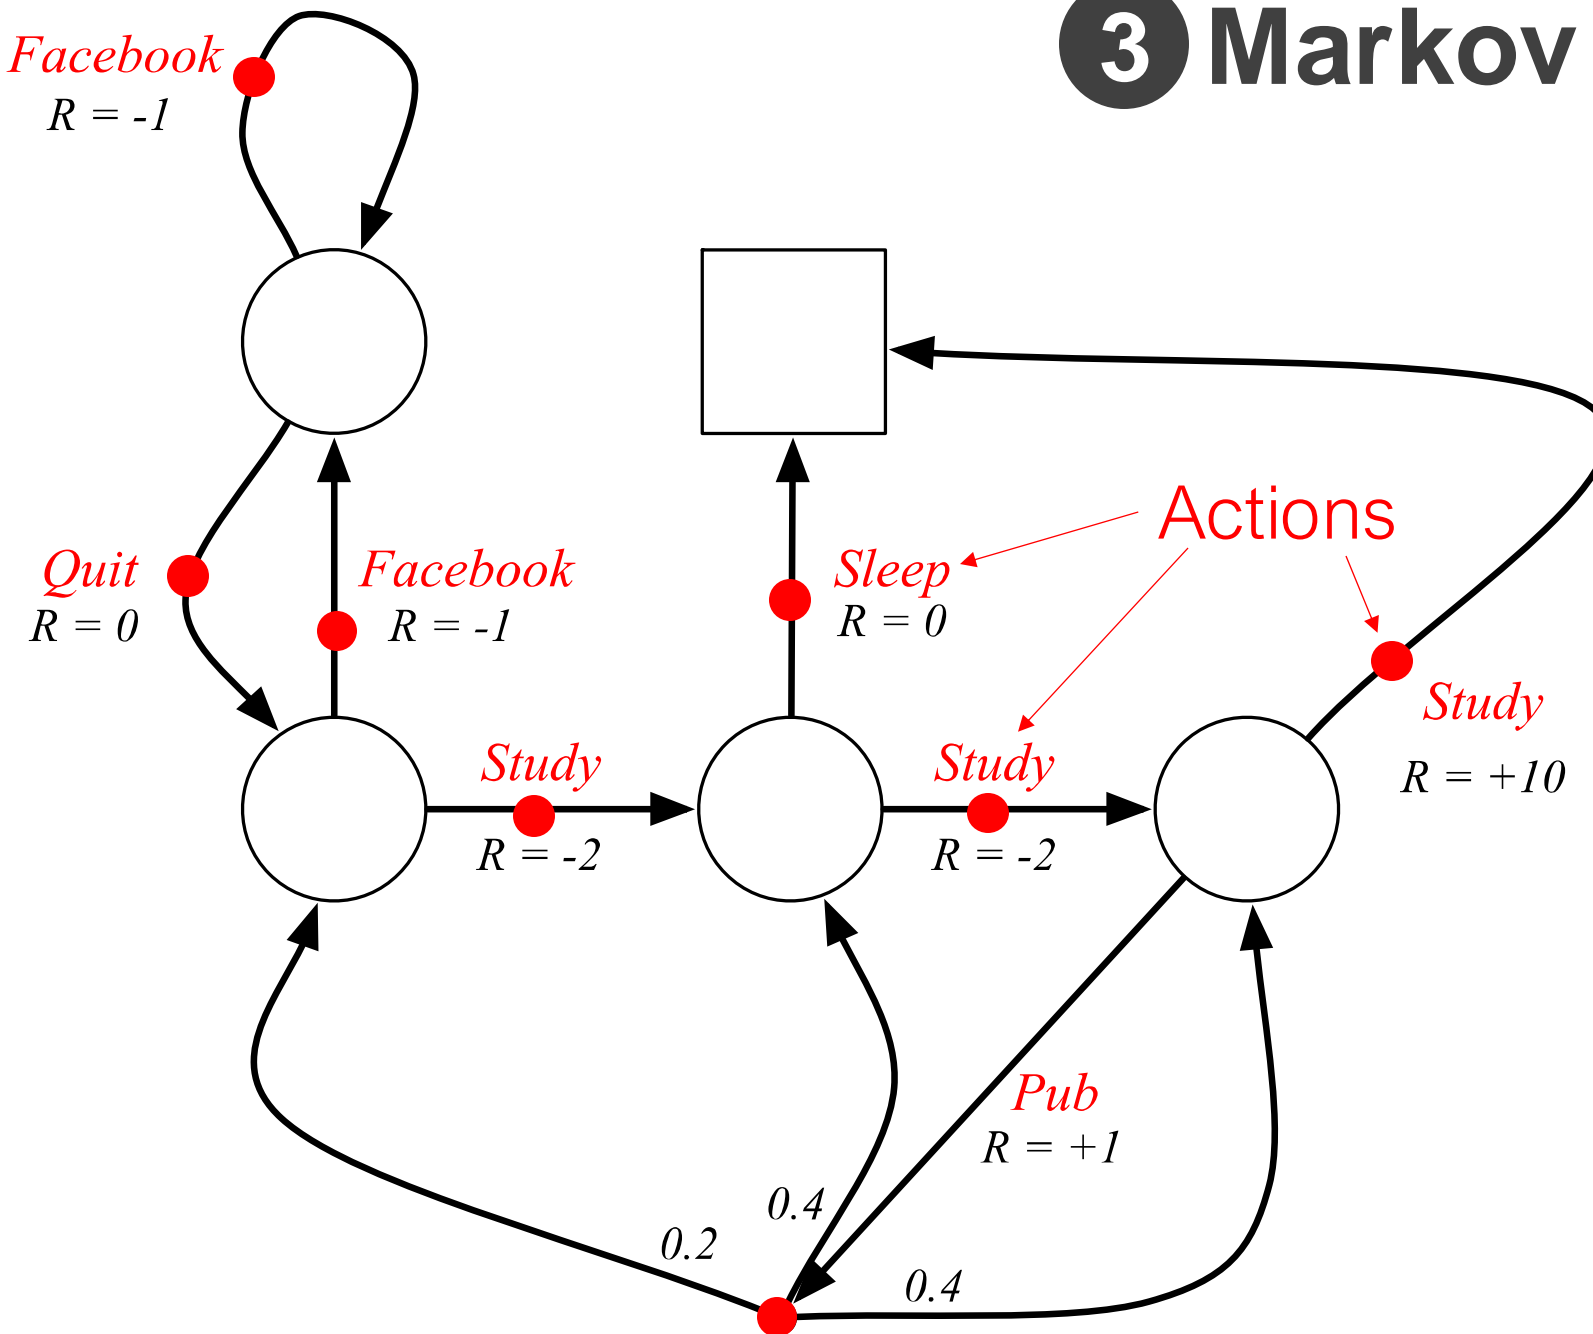
# 3 Markov Decision Process

## Components:

State space  $S$ ,  
Transition probabilities,  $P$   
Rewards,  $R$   
Discount rate,  $\gamma$   
Actions,  $A$

Adds interaction with the environment

An agent in a state chooses an action, the environment (the MDP) provides a reward and the next state



Example from David Silver, UCL, 2015

# 3 Markov Decision Process

Policy (how we choose actions)

(can be stochastic or deterministic)

$$\pi(a|s) = P(a|s)$$

State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t|s]$$

$$v_{\pi}(s) = E[R_{t+1} + \gamma v_{\pi}(S_{t+1})|S_t = s]$$


Action value function

(expected return from state  $s$ , taking action  $a$ , and following policy  $\pi$ )

$$q_{\pi}(s, a) = E_{\pi}[G_t|s, a]$$

$$q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1})|s, a]$$

Example from David Silver, UCL, 2015



# “Backup” property of state value functions

$$\begin{aligned} v_{\pi}(s) &\triangleq E_{\pi}[G_t | S_t = s] \quad \text{where } G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots \\ &= E_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots | S_t = s] \\ &= E_{\pi}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} \dots) | S_t = s] \\ &= E_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= E_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \end{aligned}$$


This recursive relationship is a version of the **Bellman Equation**

Example from David Silver, UCL, 2015

# Bellman Expectation Equations for the **state value** function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) \leftarrow s$$



$$v_{\pi}(s) = E_{\pi}[G_t | S_t = s]$$

$$\begin{aligned} v_{\pi}(s) &= E_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= E_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \end{aligned}$$

Expectation over the possible actions

Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

$$v_{\pi}(s) = \sum_{\underline{a}} \pi(a|s) \sum_{\underline{s'}} \sum_{\underline{r}} p(s', r | s, a) [r + \gamma v_{\pi}(s')]$$

$p(s', r | s, a)$  is the joint distribution of state transitions and rewards given you start in state  $s$  and take action  $a$


# Bellman Expectation Equations for the **action value** function

(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )

$$q_{\pi}(s, a) = E_{\pi}[G_t | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) \leftarrow s, a$$



Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

Expectation over the possible actions

$$q_{\pi}(s, a) = \sum_{\underline{s'}} \sum_{\underline{r}} p(s', r | s, a) \left[ r + \gamma \sum_{\underline{a'}} \pi(a' | s') q_{\pi}(s', a') \right]$$




# Bellman Expectation Equations

## State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E_{\pi}[G_t | S_t = s]$$

$$v_{\pi}(s) = E_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]$$




$$v_{\pi}(s) = \sum_a \pi(a|s) \underbrace{\sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_{\pi}(s')]}_{q_{\pi}(s, a)}$$

## Action value function

(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )

$$q_{\pi}(s, a) = E_{\pi}[G_t | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$




$$q_{\pi}(s, a) = \sum_{s'} \sum_r p(s', r | s, a) \left[ r + \gamma \underbrace{\sum_{a'} \pi(a' | s') q_{\pi}(s', a')}_{v_{\pi}(s')} \right]$$

# Example

Policy: randomly choose an action

$$\pi(a_1|s_0) = \pi(a_2|s_0) = 0.5$$



$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r|s, a) [r + \gamma v_\pi(s')]$$

$$v_\pi(s_0) = \underbrace{\pi(a_1|s_0)}_{0.5} \underbrace{(\underbrace{5}_{r} - \underbrace{1}_{v_\pi(s_1)})}_{q_\pi(s_0, a_1)} + \underbrace{\pi(a_2|s_0)}_{0.5} \underbrace{[(\underbrace{0.9}_{p})(\underbrace{1}_{r}) + (\underbrace{0.1}_{p})(\underbrace{-50}_{r}) + (\underbrace{0.9}_{p})(\underbrace{10}_{v_\pi(s_2)}) + (\underbrace{0.1}_{p})(\underbrace{0}_{v_\pi(s_3)})]}_{q_\pi(s_0, a_2)} \quad \gamma = 1$$

# 3 Markov Decision Process

State value function,  $v_\pi(s)$

Assumptions:

$$\pi(a|s) = \frac{1}{\text{\# of reachable states}}$$

(randomly select an action)

$$\gamma = 1$$

(no discounting)


Bellman Expectation Equation

$$v_\pi(s) = E_\pi[R_{t+1} + \gamma v_\pi(S_{t+1})|s]$$

$$-1.3 \approx 0.5 \times (-1 - 2.3) + 0.5 \times (-2 + 2.7)$$

Assume  $\gamma = 1$

Example from David Silver, UCL, 2015




# 3 Markov Decision Process

Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies

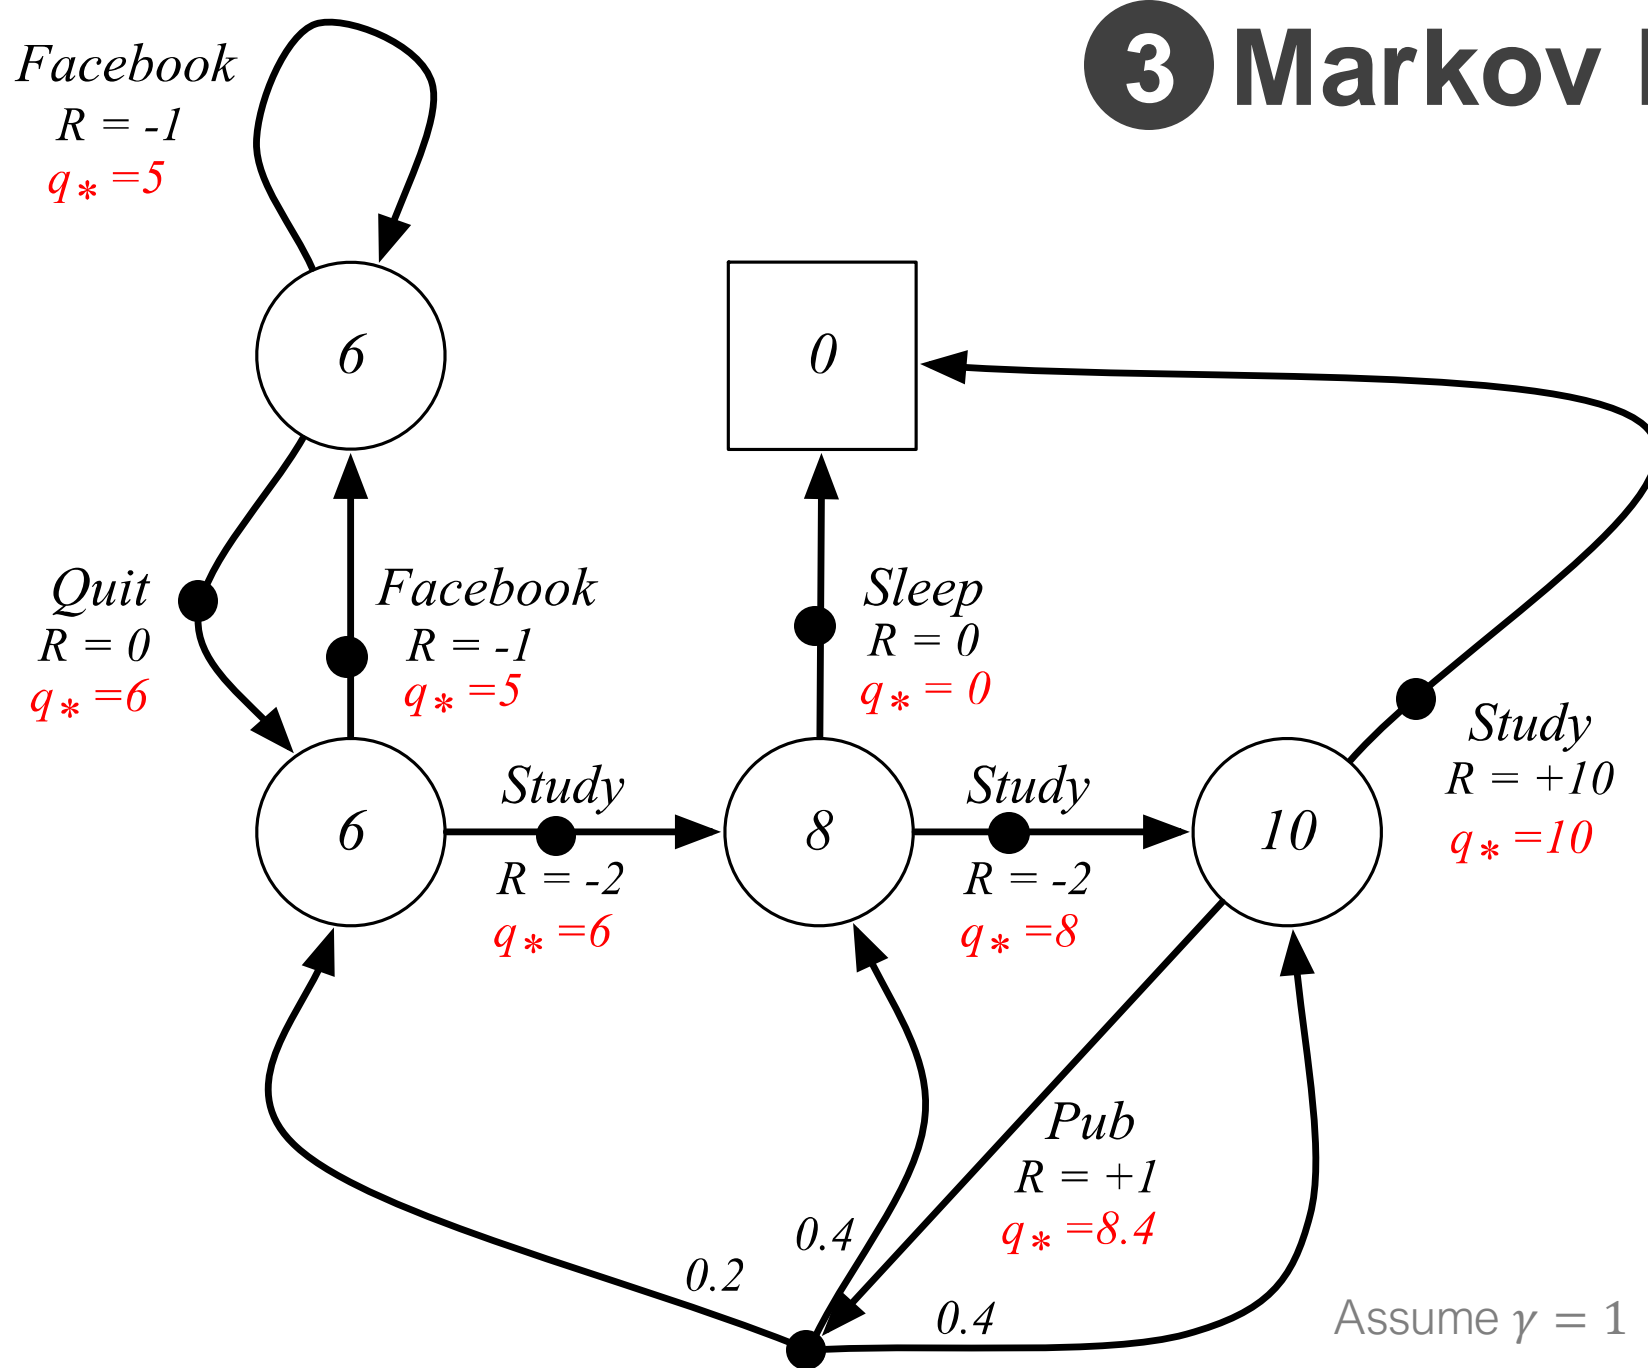
$$v_*(s) = \max_{\pi} v_{\pi}(s)$$



Assume  $\gamma = 1$

Example from David Silver, UCL, 2015

# 3 Markov Decision Process



Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

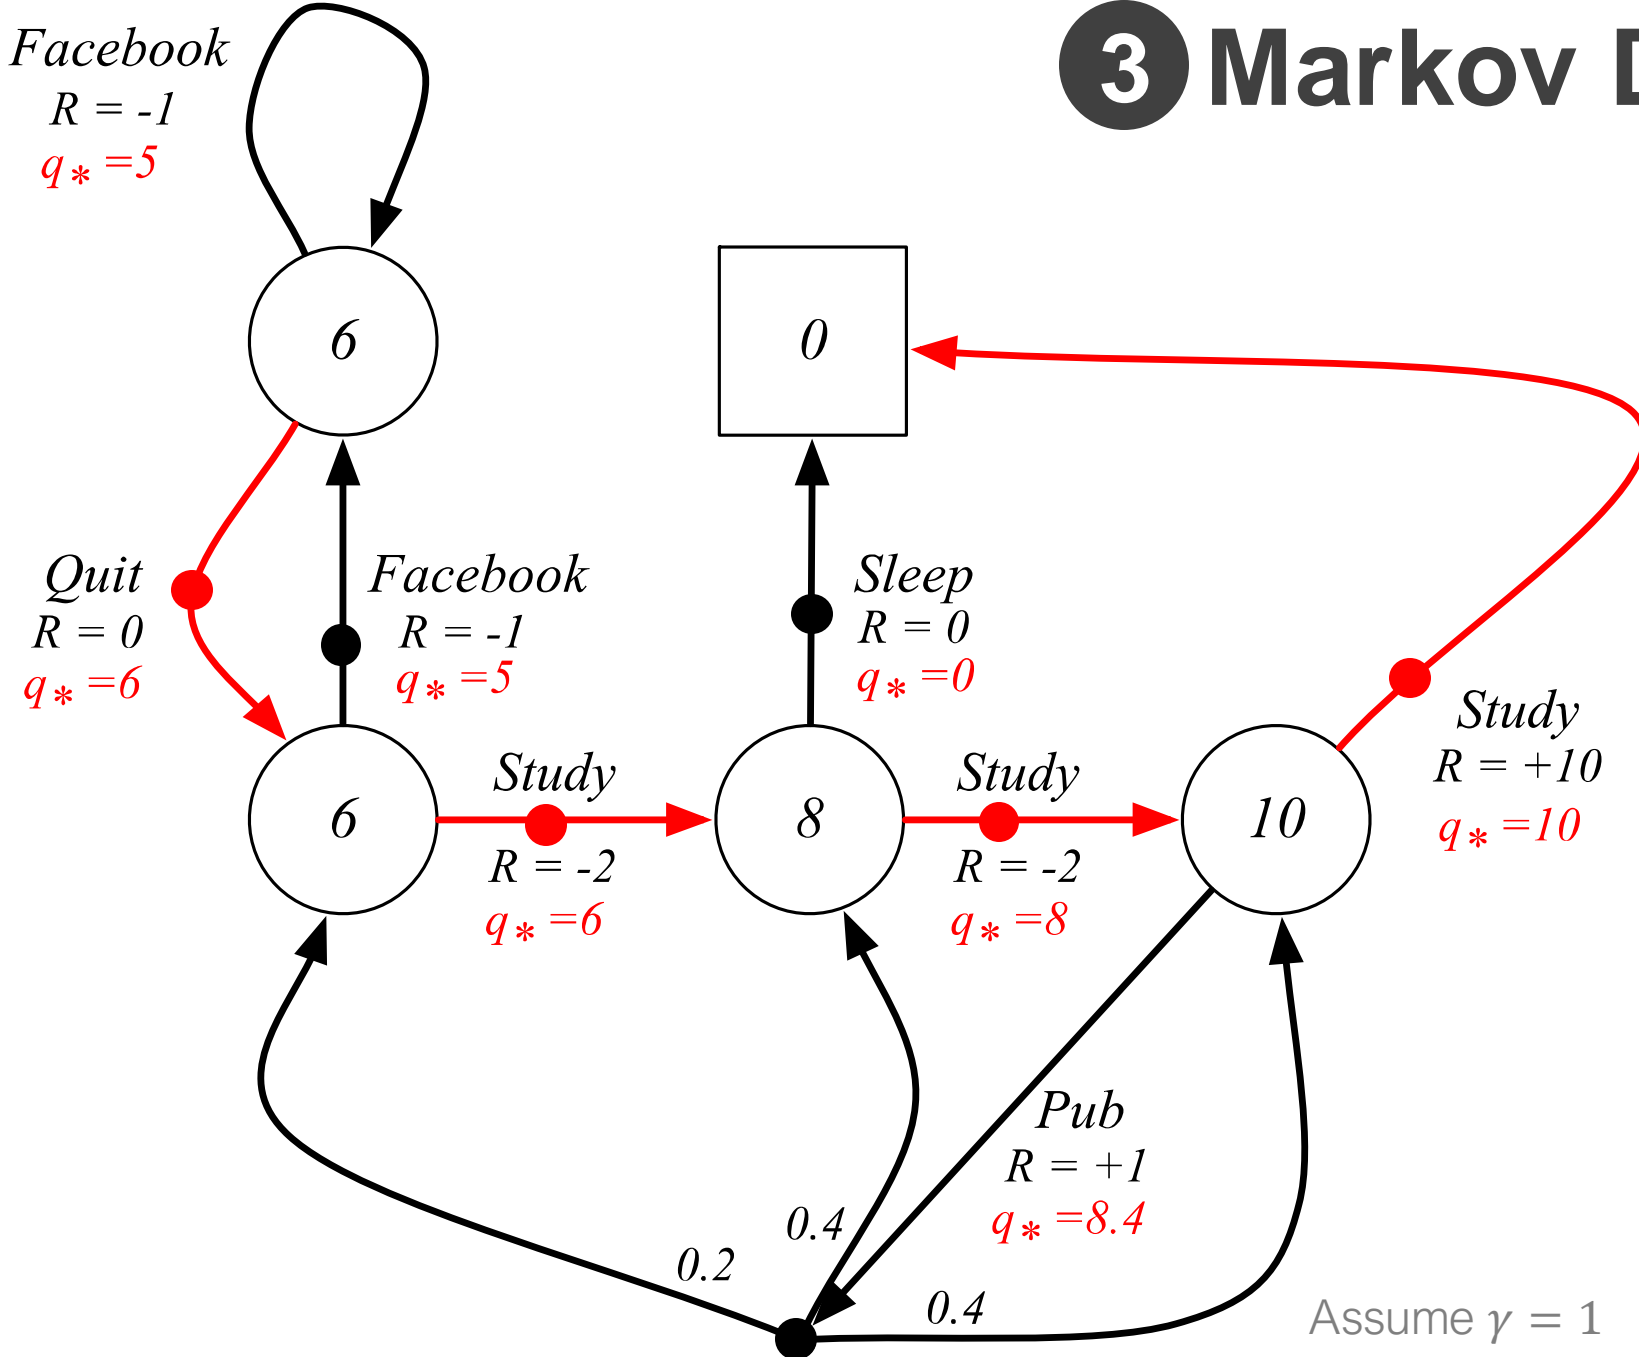
Optimal **action-value** function,  $q_*(s, a)$

Maximum value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

Example from David Silver, UCL, 2015

# 3 Markov Decision Process



Optimal **policy**,  $\pi_*(s)$   
Which action to take at each moment

$$\pi_*(s) = \arg \max_a q_*(s, a)$$

Once we have the optimal value functions, we've "solved" the MDP!

Example from David Silver, UCL, 2015

# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)

- RL methods **do NOT ASSUME** knowledge of  $P$  or  $R$  (while dynamic programming does)
- RL learns/approximates that knowledge

Adapted from David Silver, 2015

# Reinforcement Learning Roadmap

- 1 Core concepts in reinforcement learning  
Actions, Rewards, Value, Environments, and Policies

- 2 Markov decision processes  
...and Markov chains and Markov reward processes

- 3 Dynamic Programming  
How do we find optimal policies?  
(Bellman equations)

- 4 Monte Carlo Control  
How do we estimate our value functions?  
How do we use the value functions to choose actions?  
How do we learn optimal policies from experience?

Knowledge of **Environment**

## Perfect knowledge

Known Markov  
Decision Process



## No knowledge

Must learn from  
experience